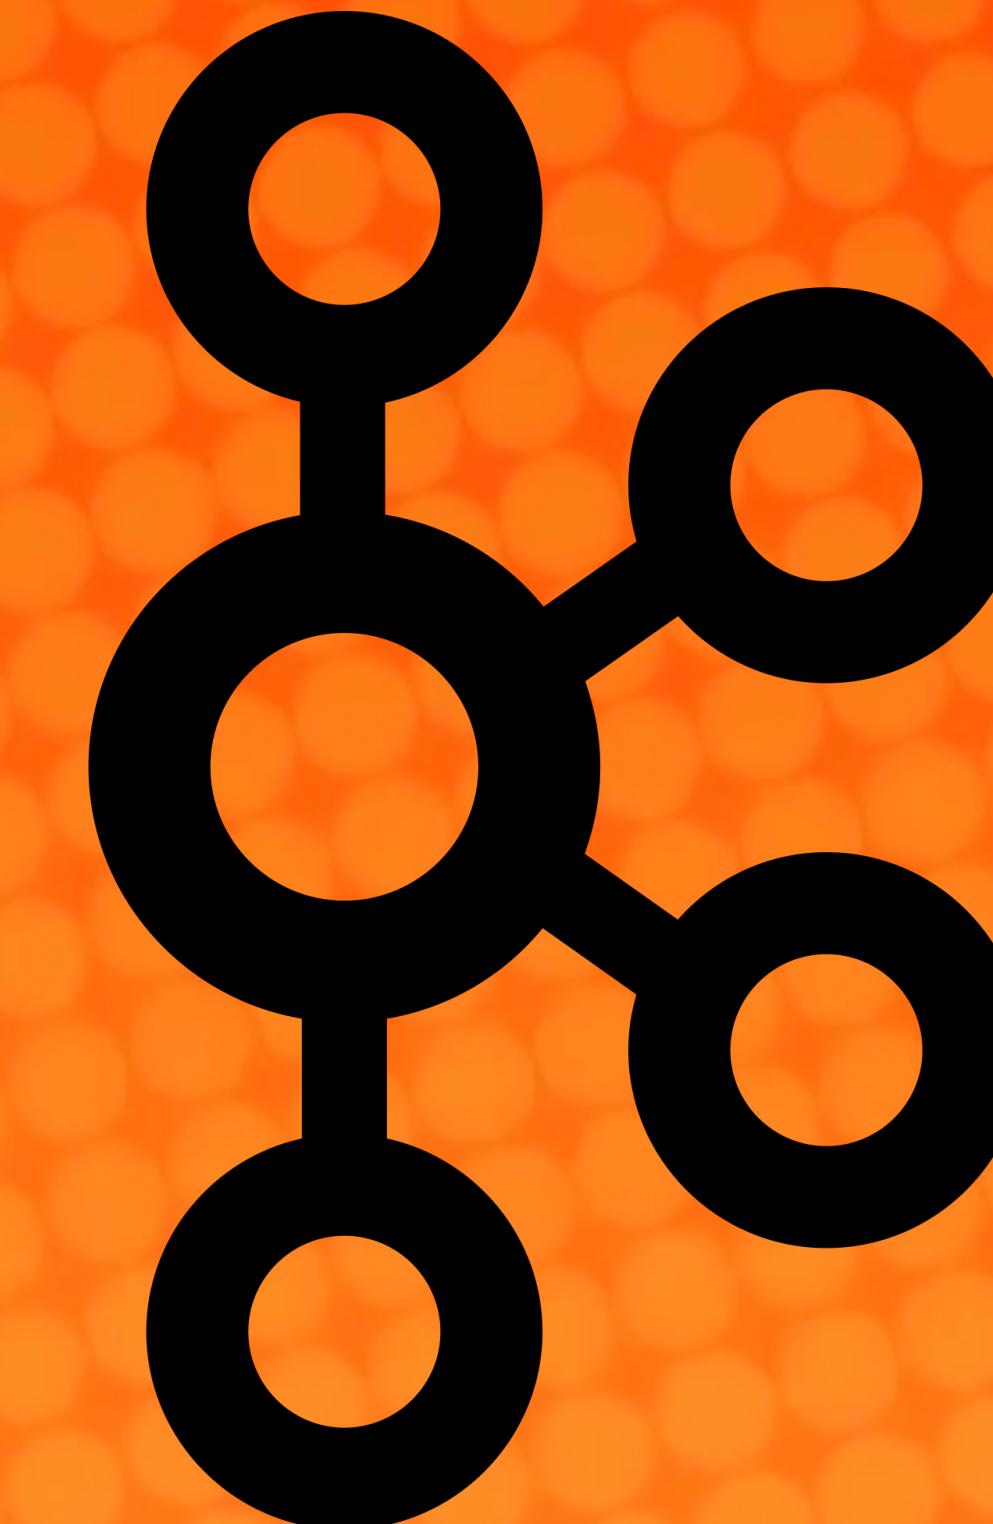


Apache Kafka

Mohammed Abdulrazeg



About Me...

Mohammed Abdulrazeg
Programmer, Consultant,
Trainer

Contact:

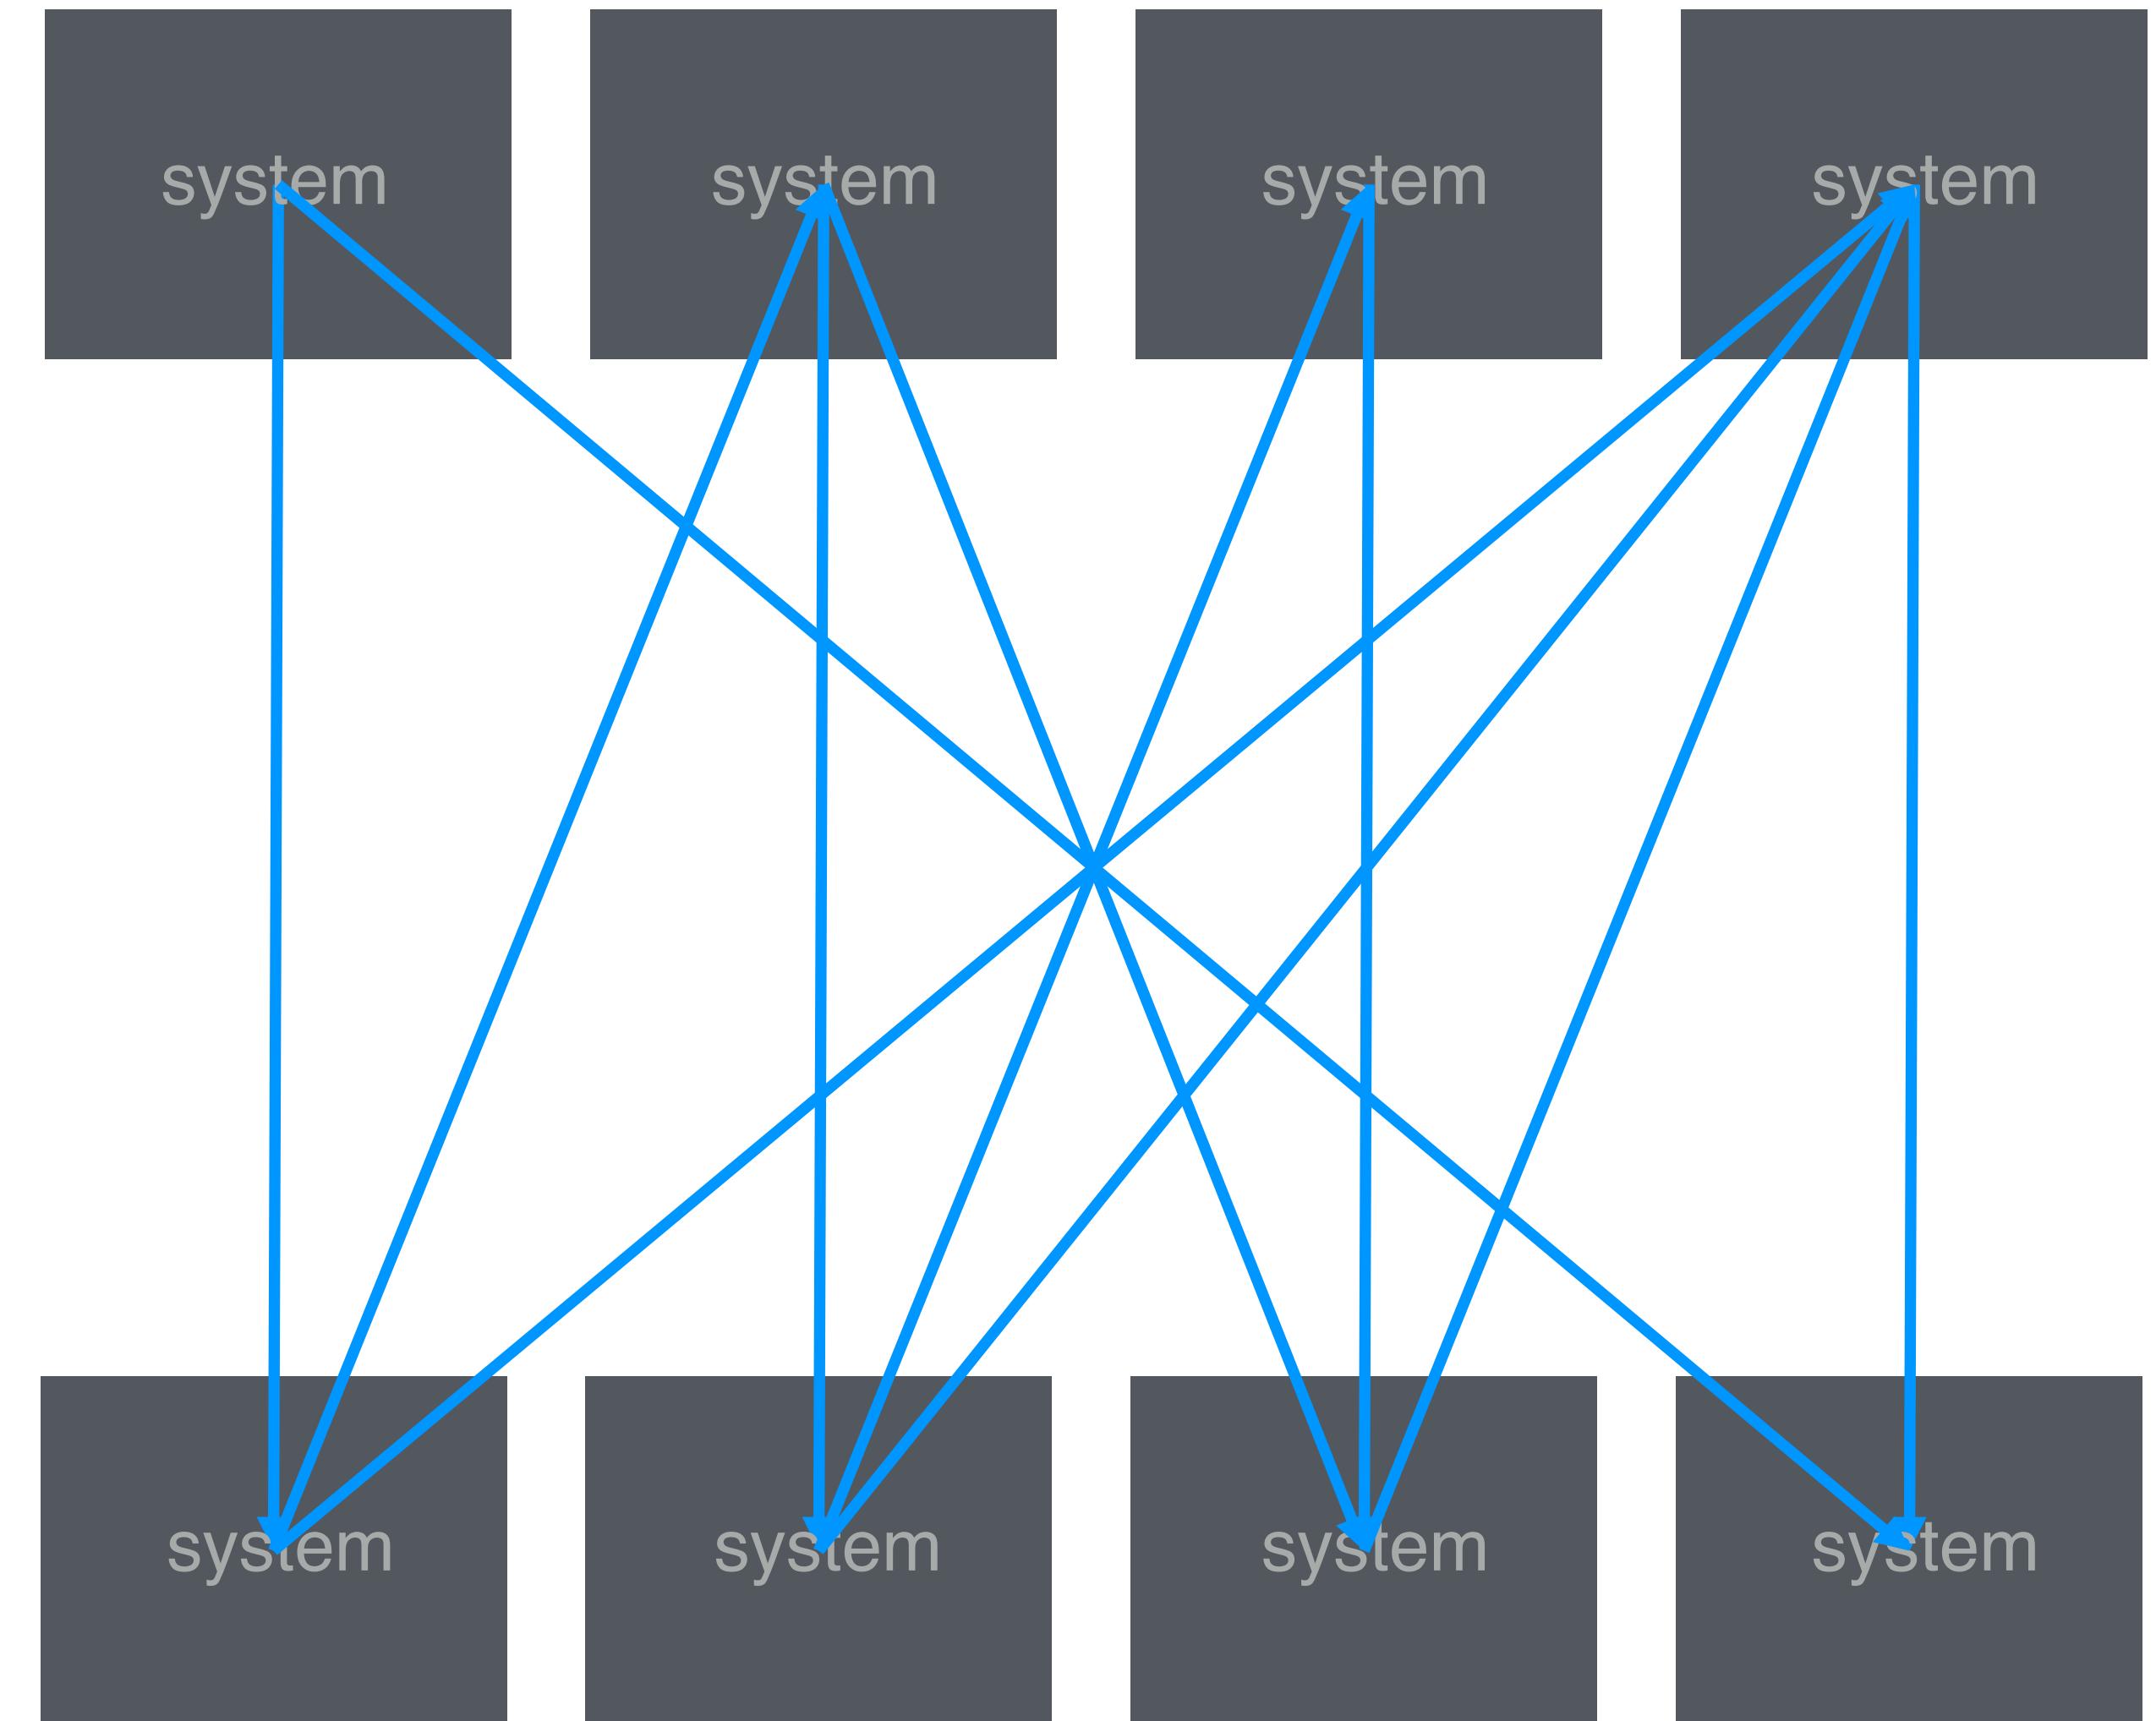
hello@mzawi.me

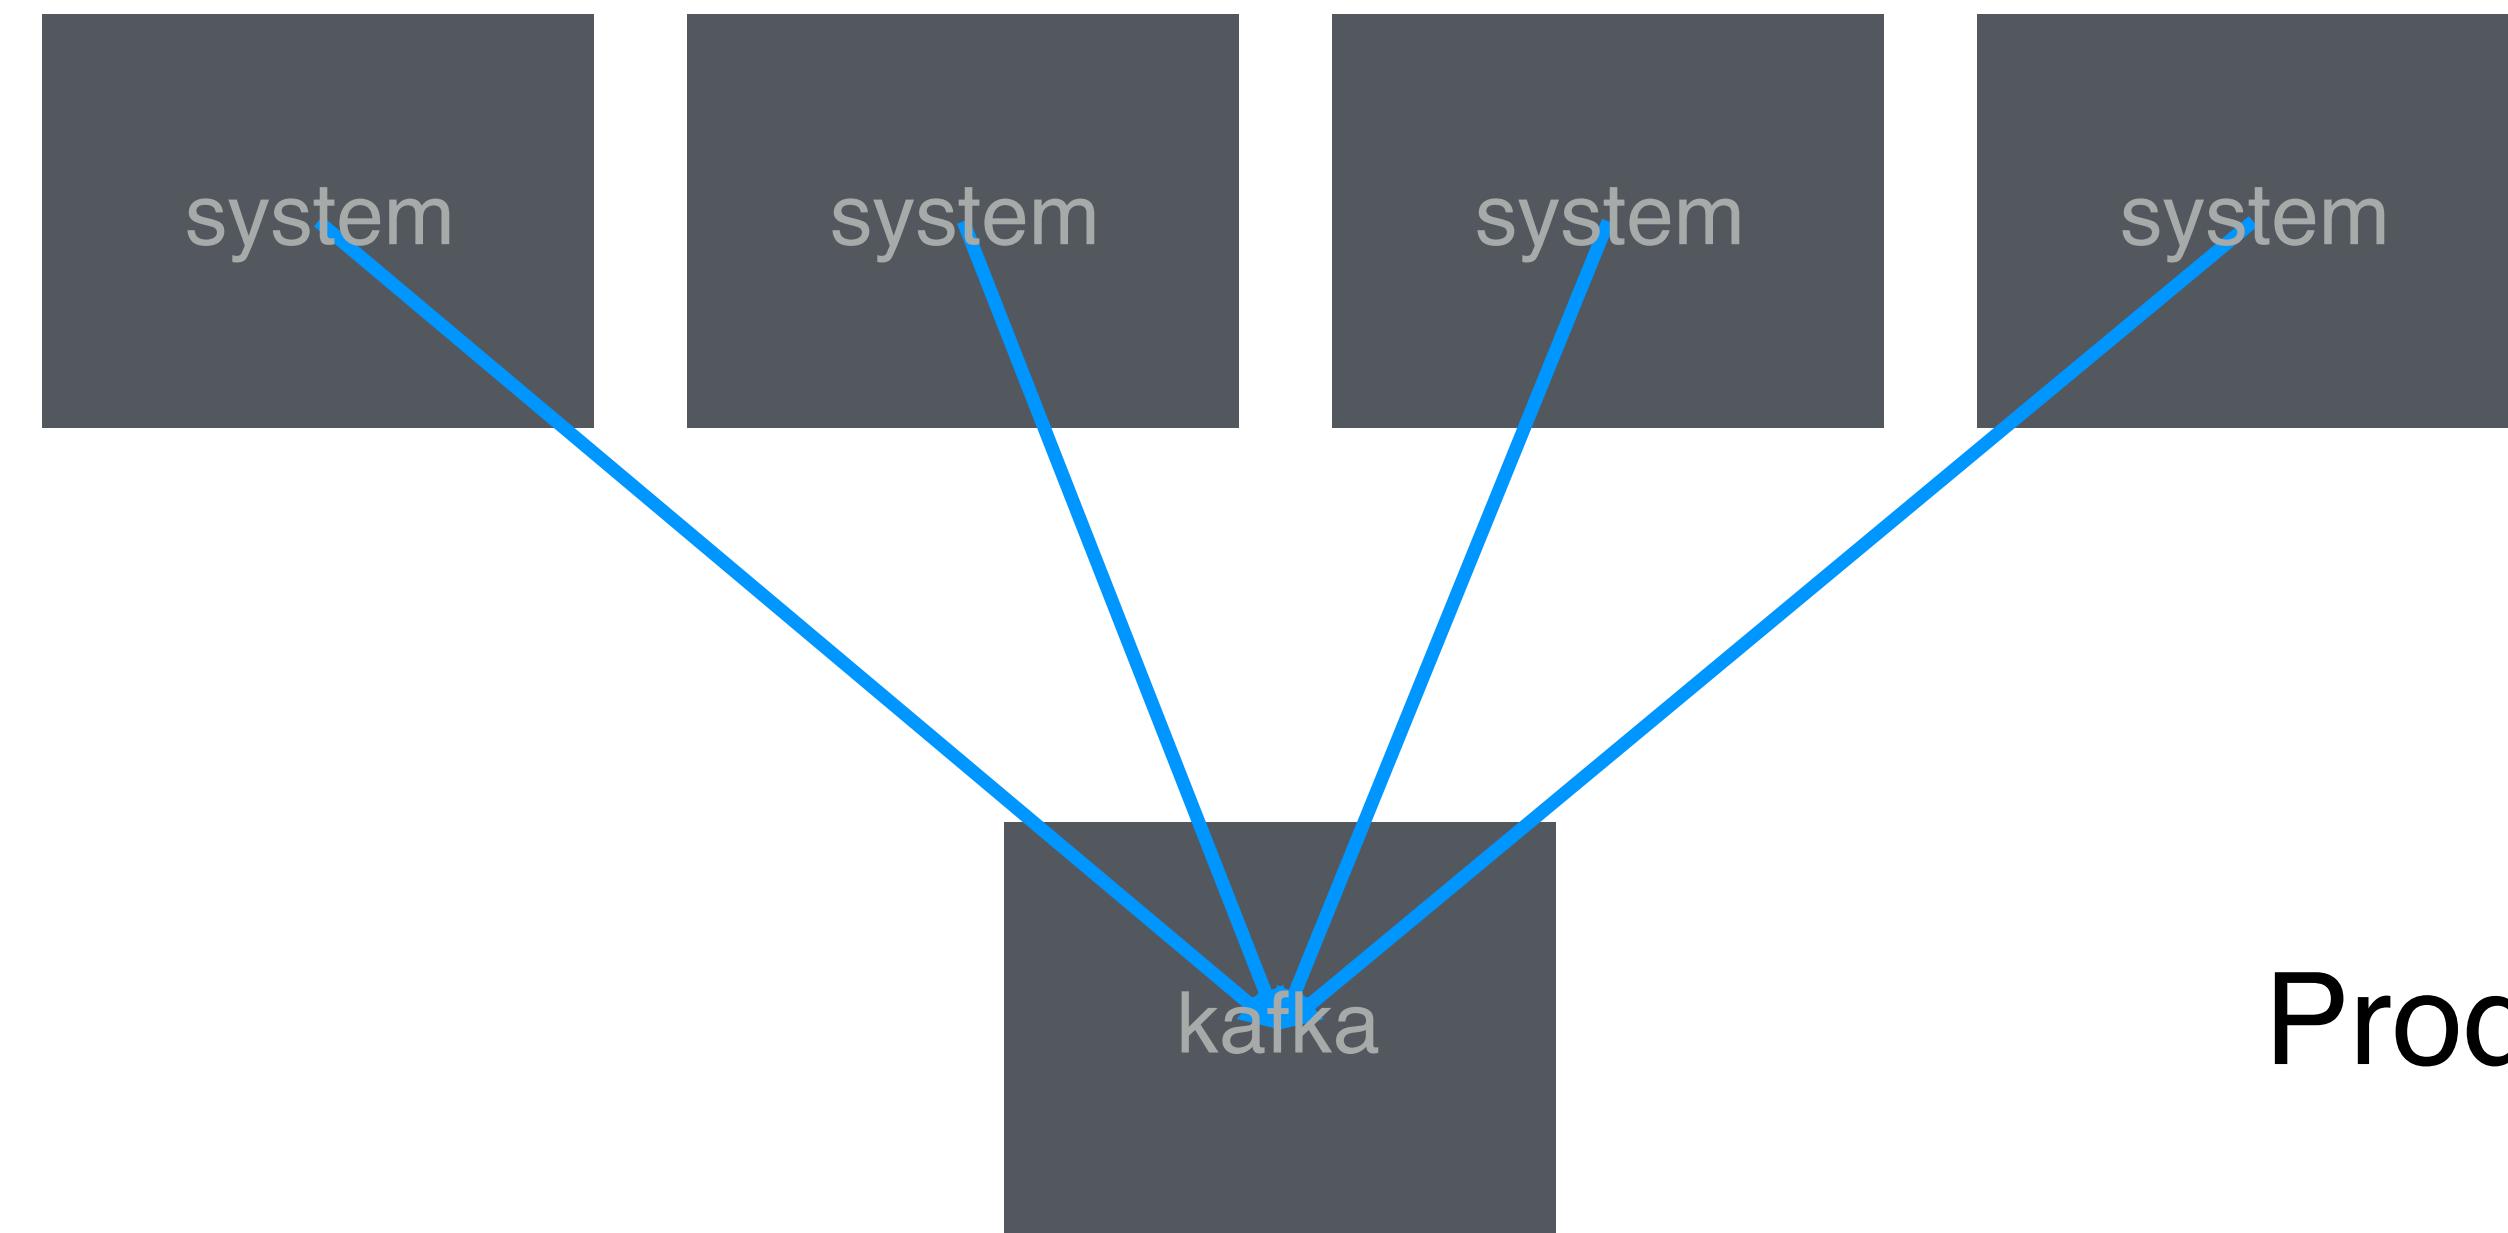
Agenda

Agenda

- ★ Understand Kafka
- ★ Understand Producer
- ★ Understand Consumer
- ★ Understand Rebalancing
- ★ Understand Streaming
- ★ Understand Tables
- ★ Understand KSQL

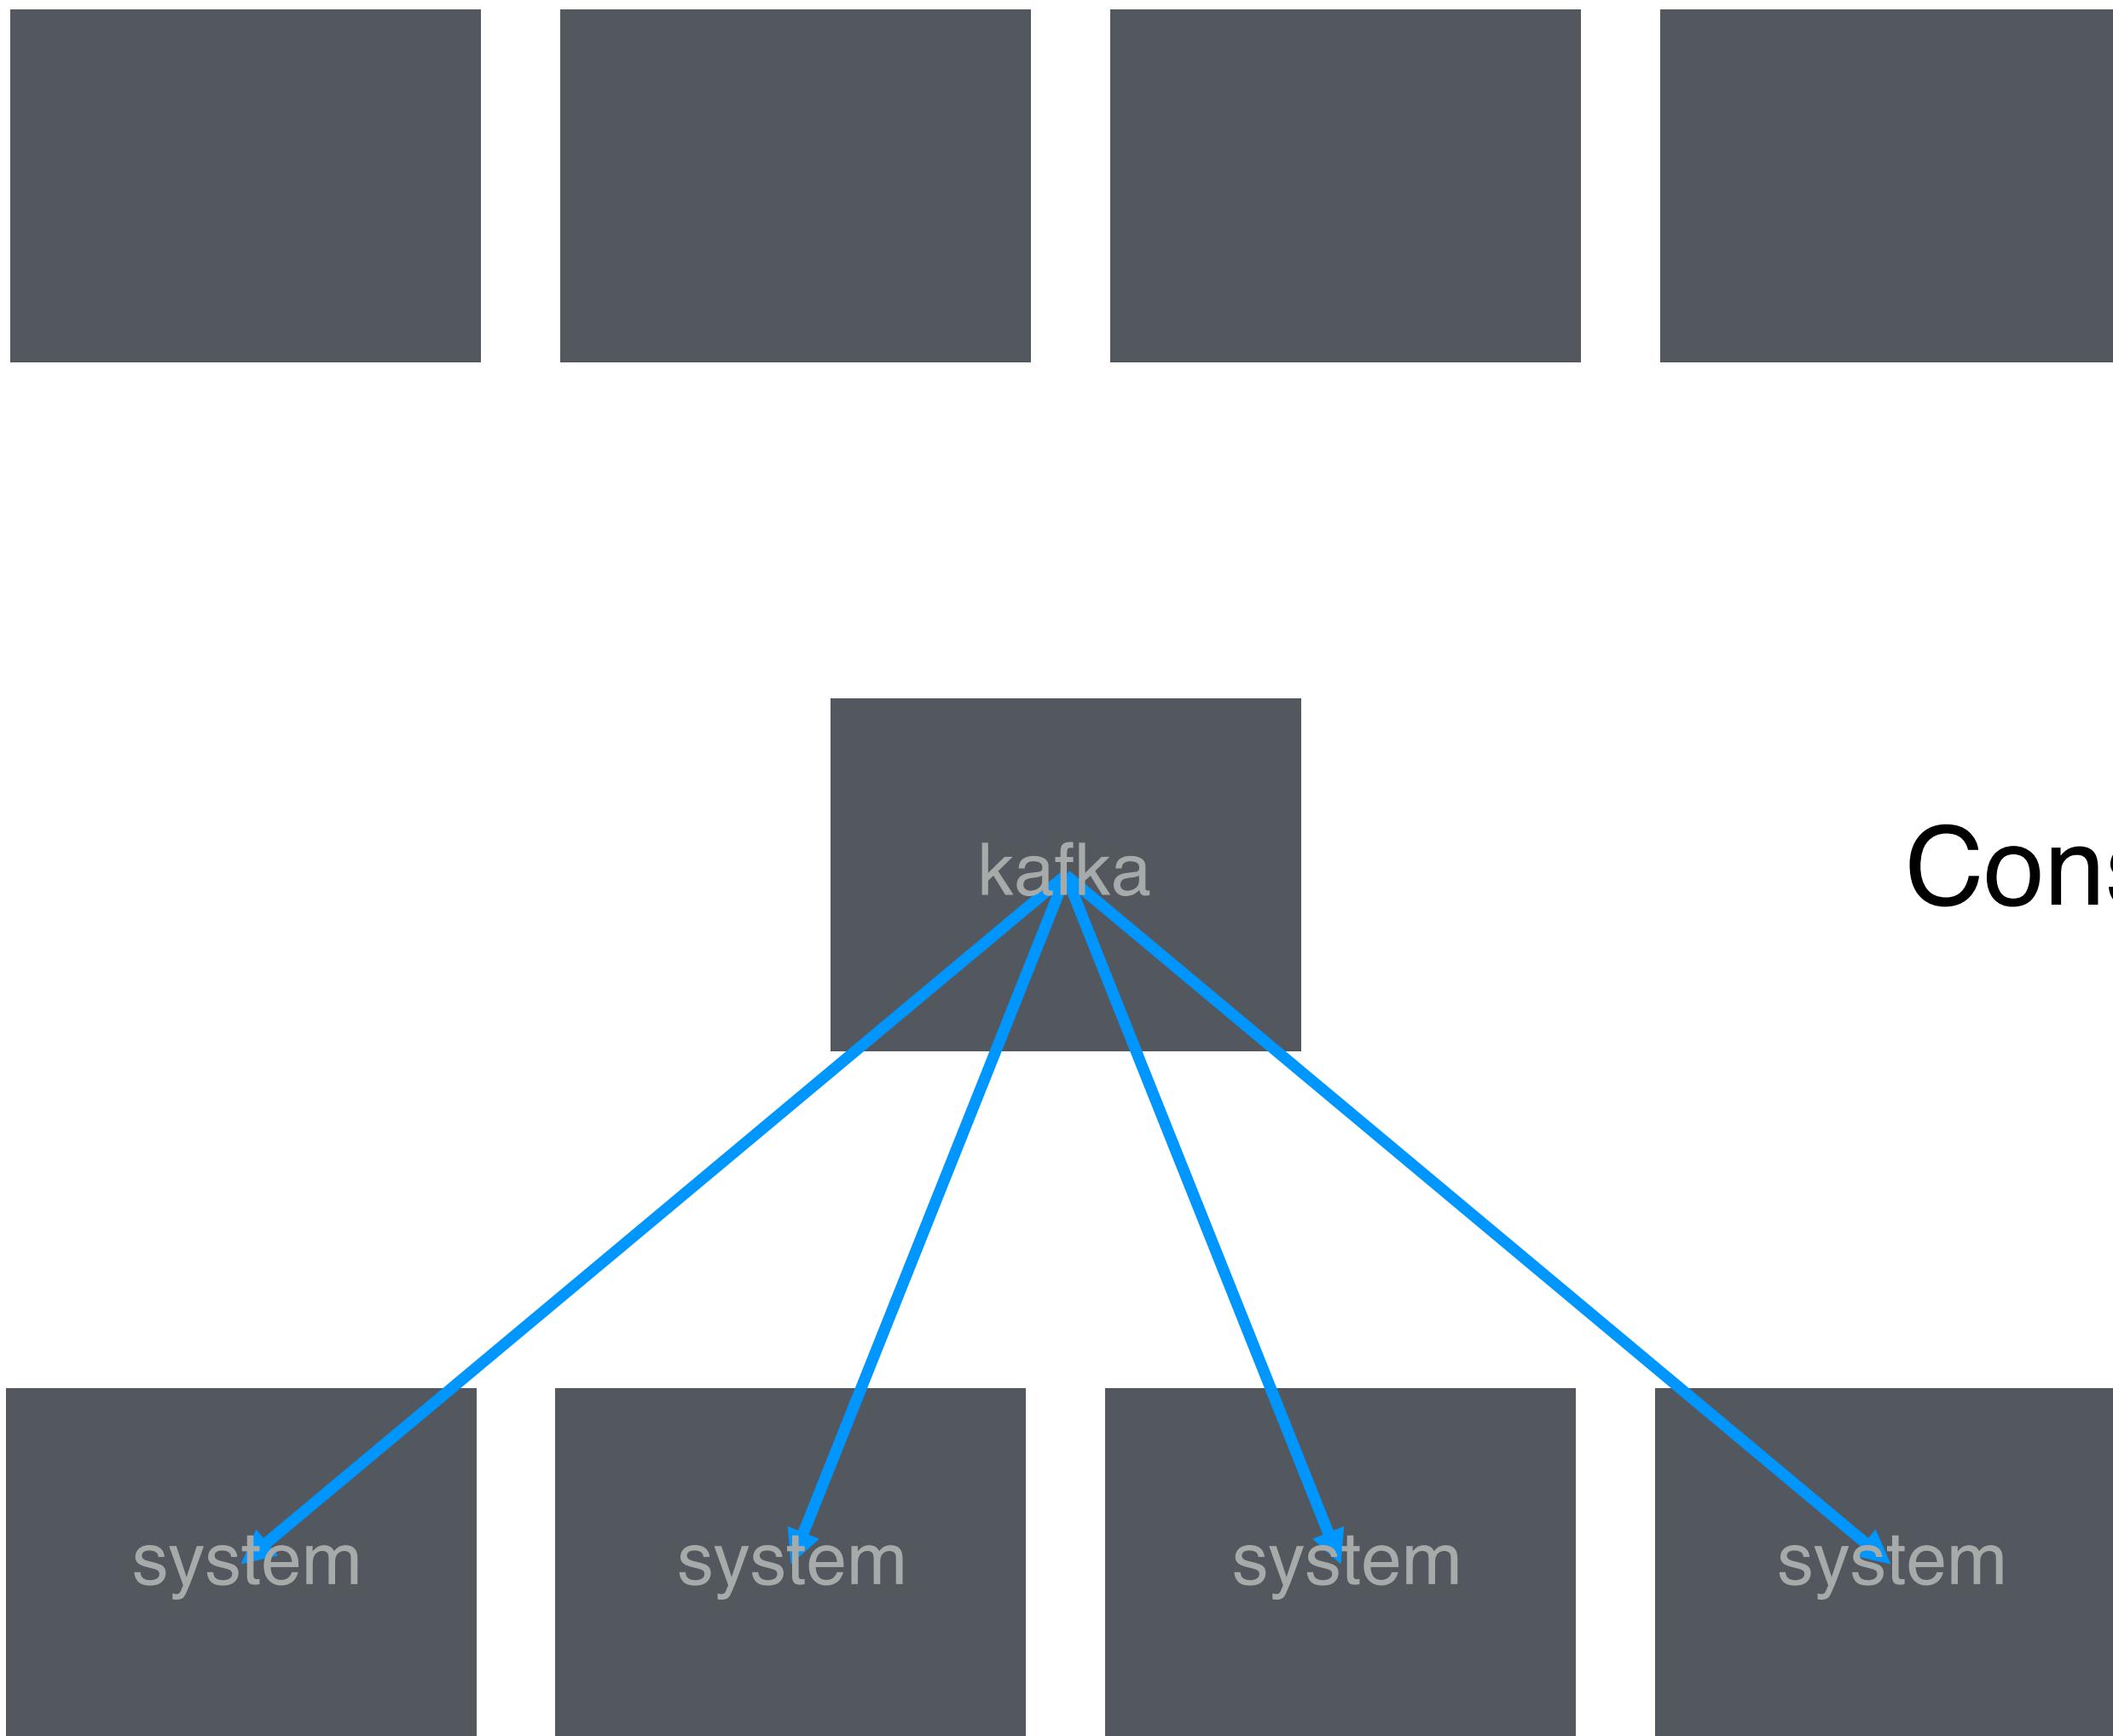
Kafka Introduction

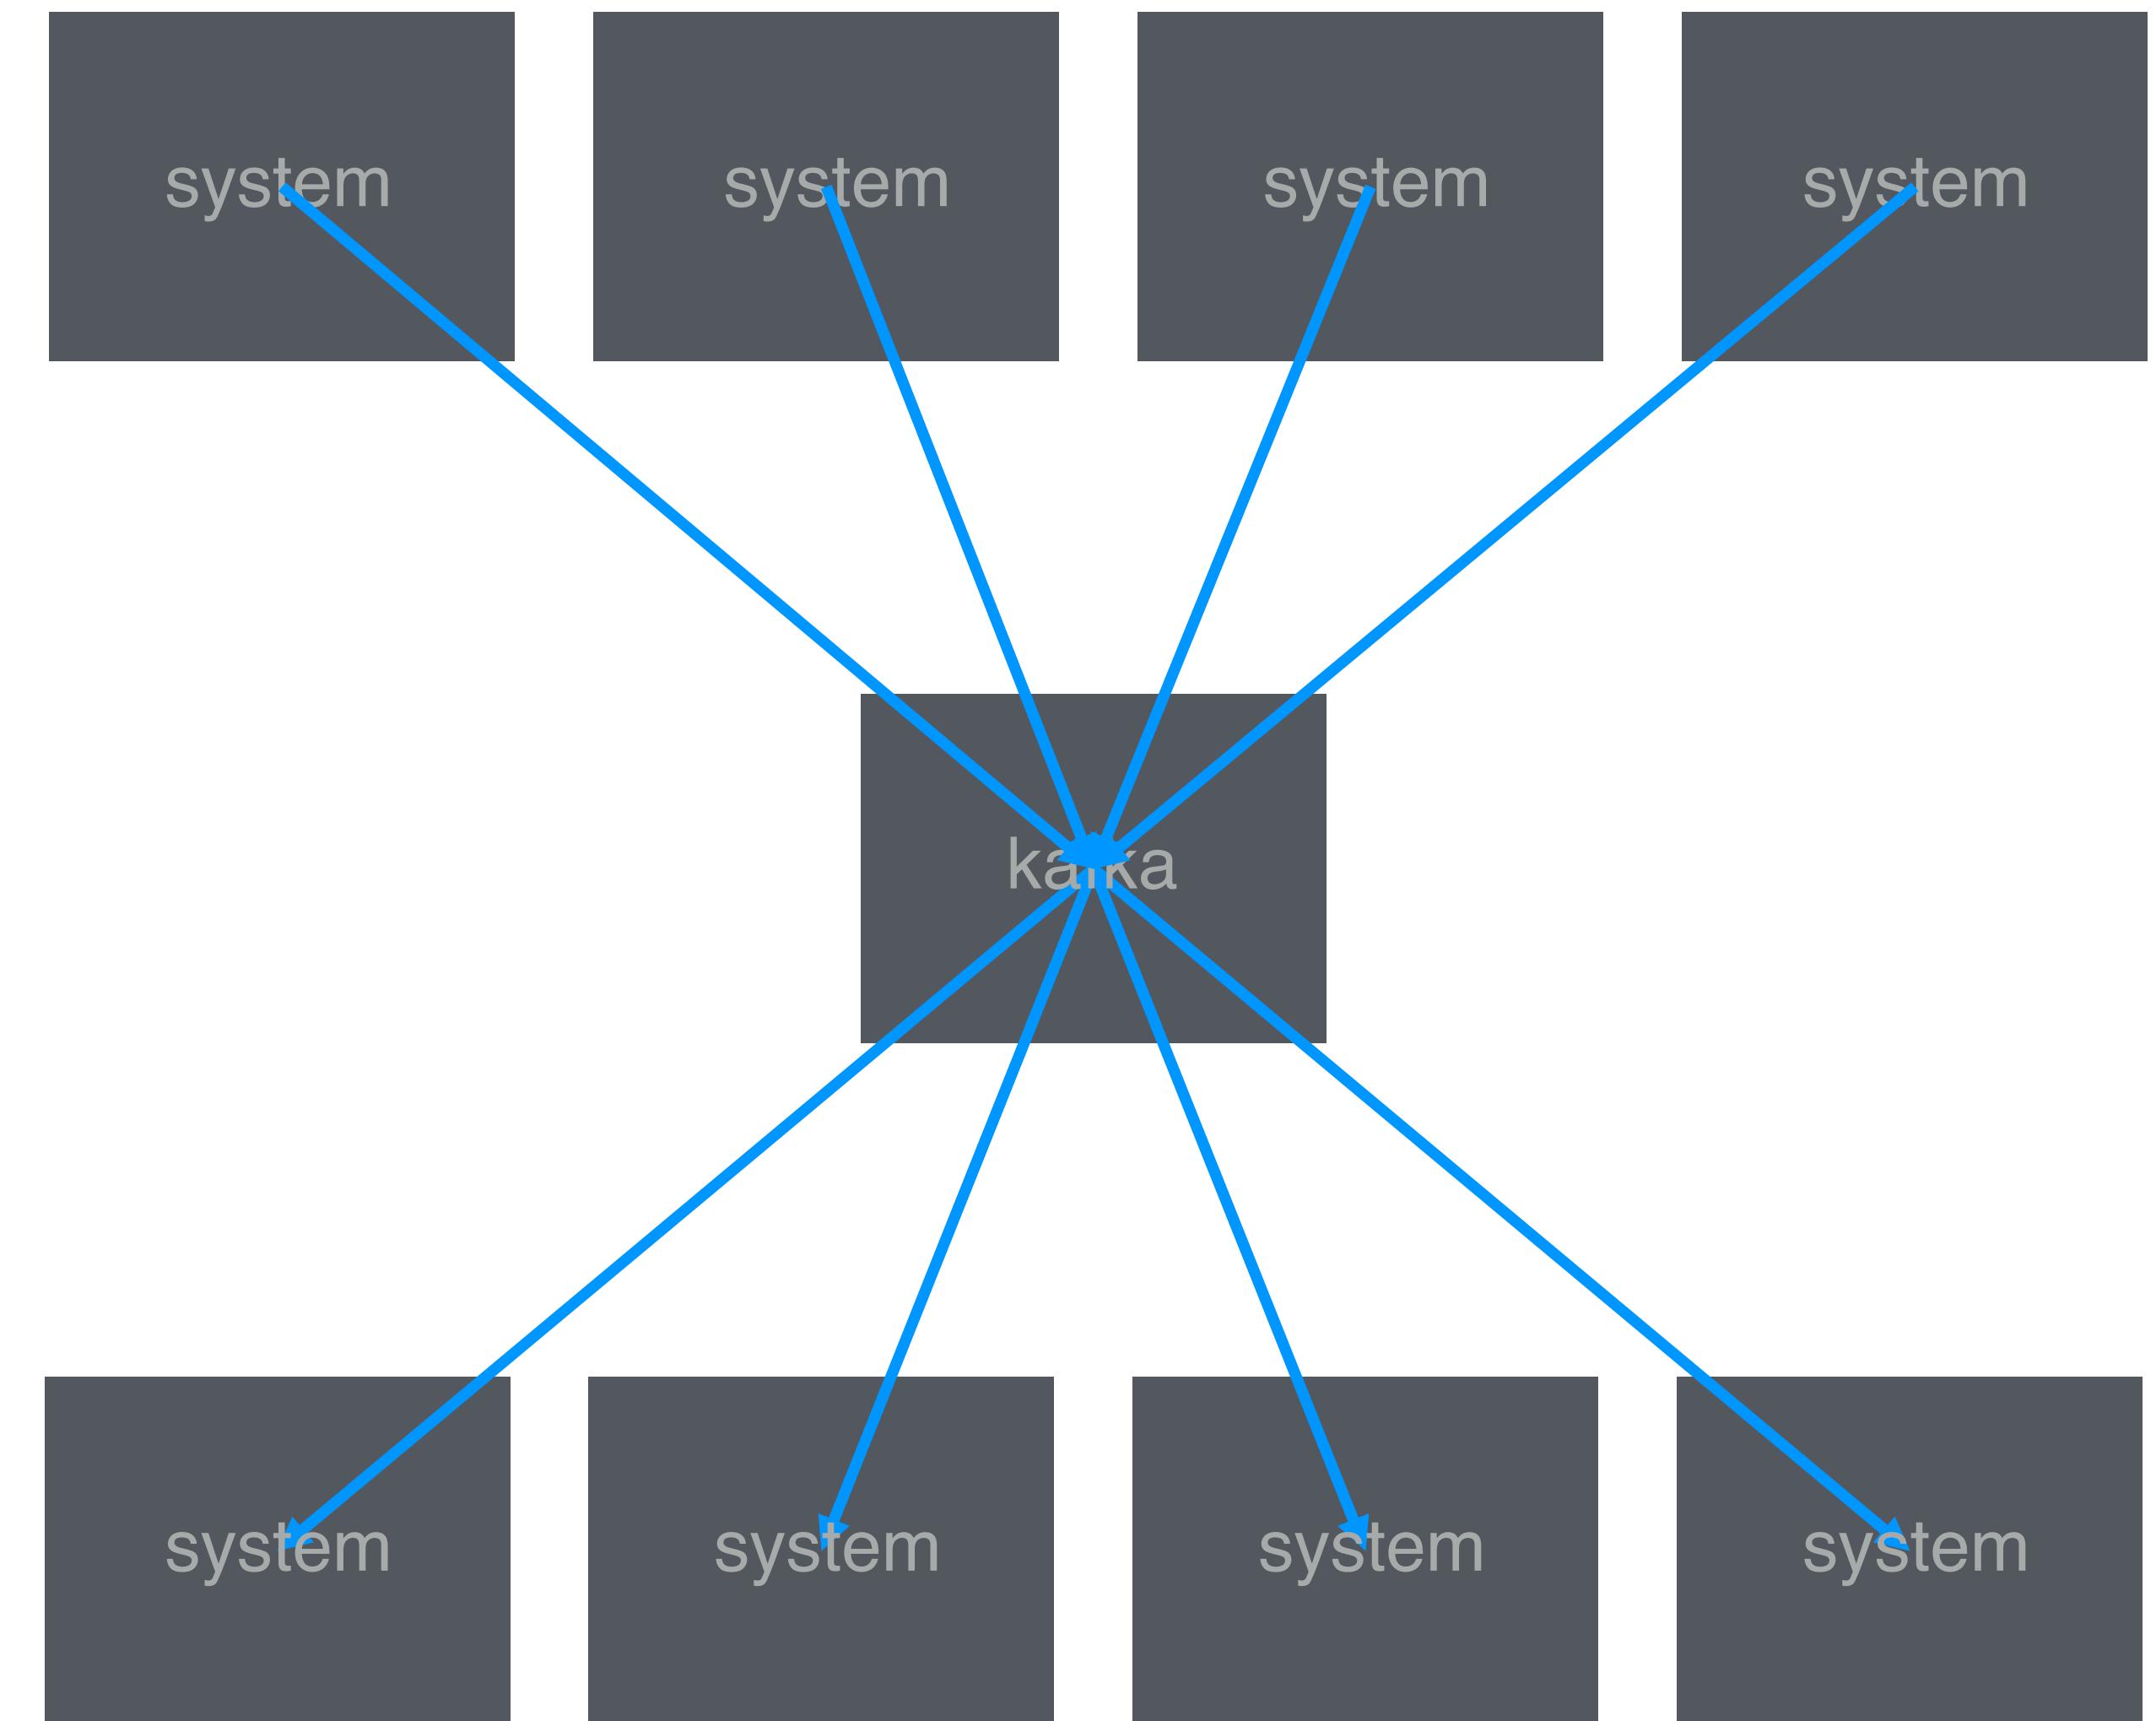


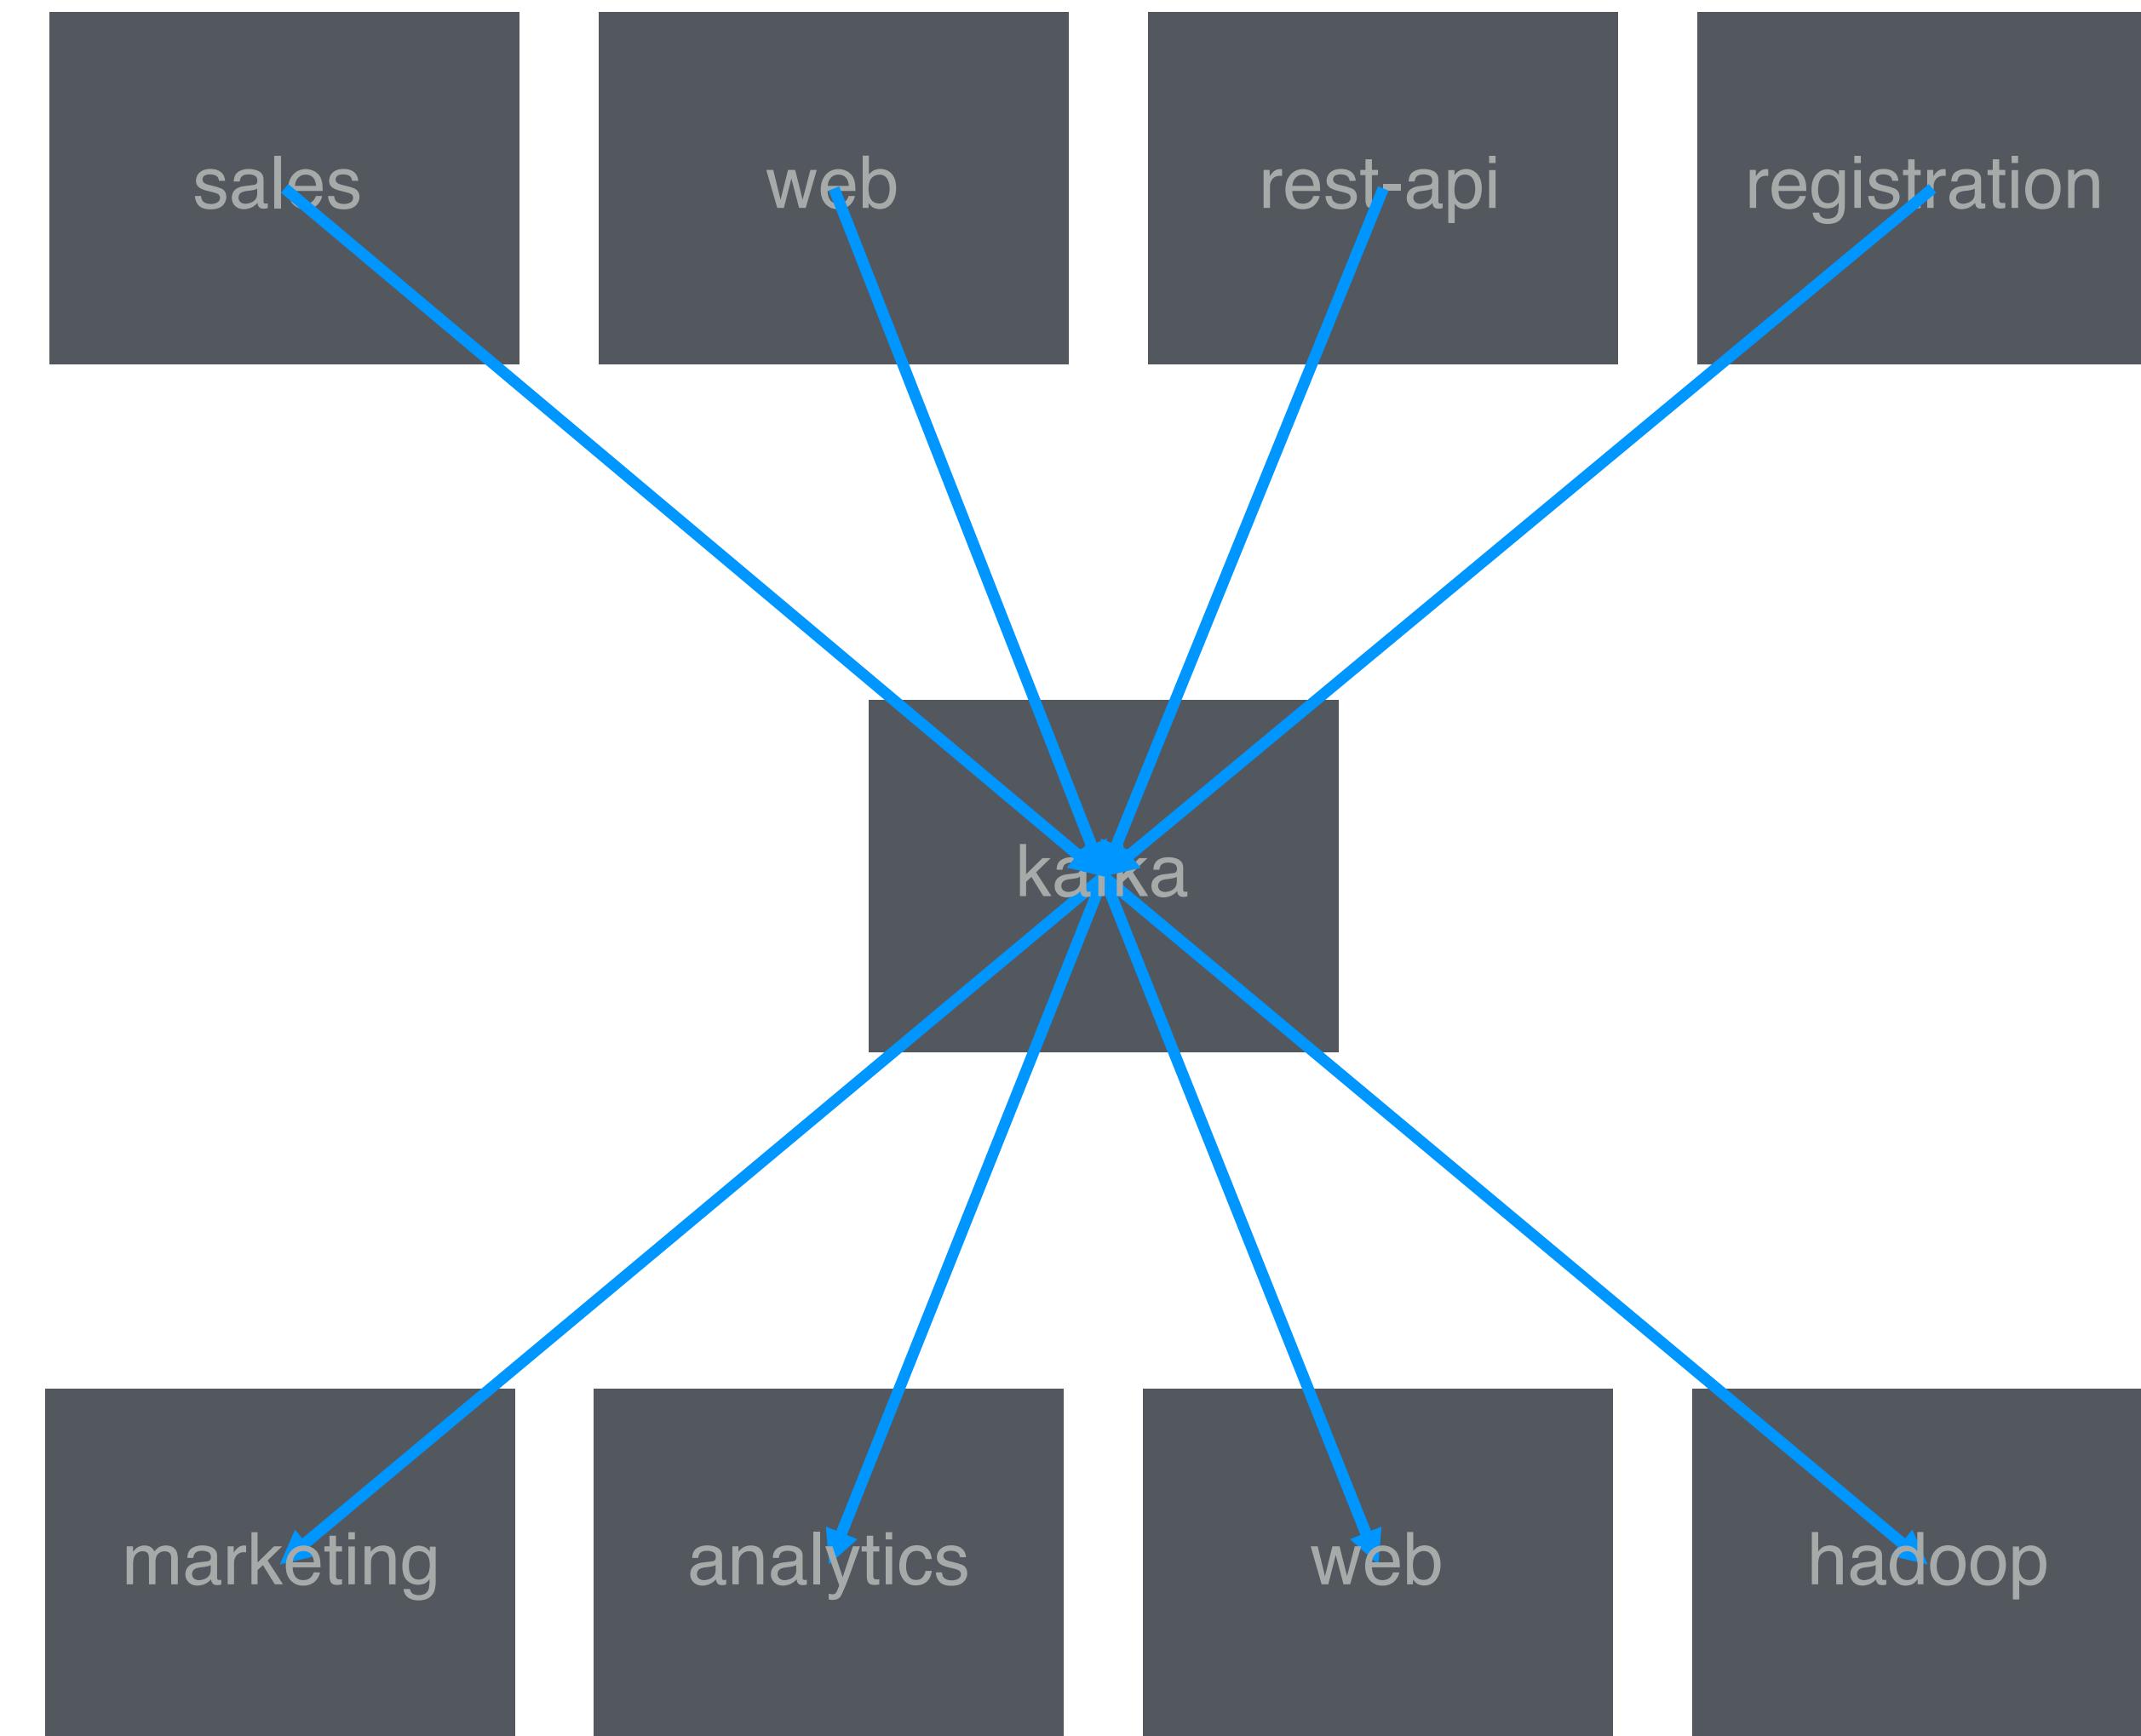


Producers









Note: A Producer can be a Consumer

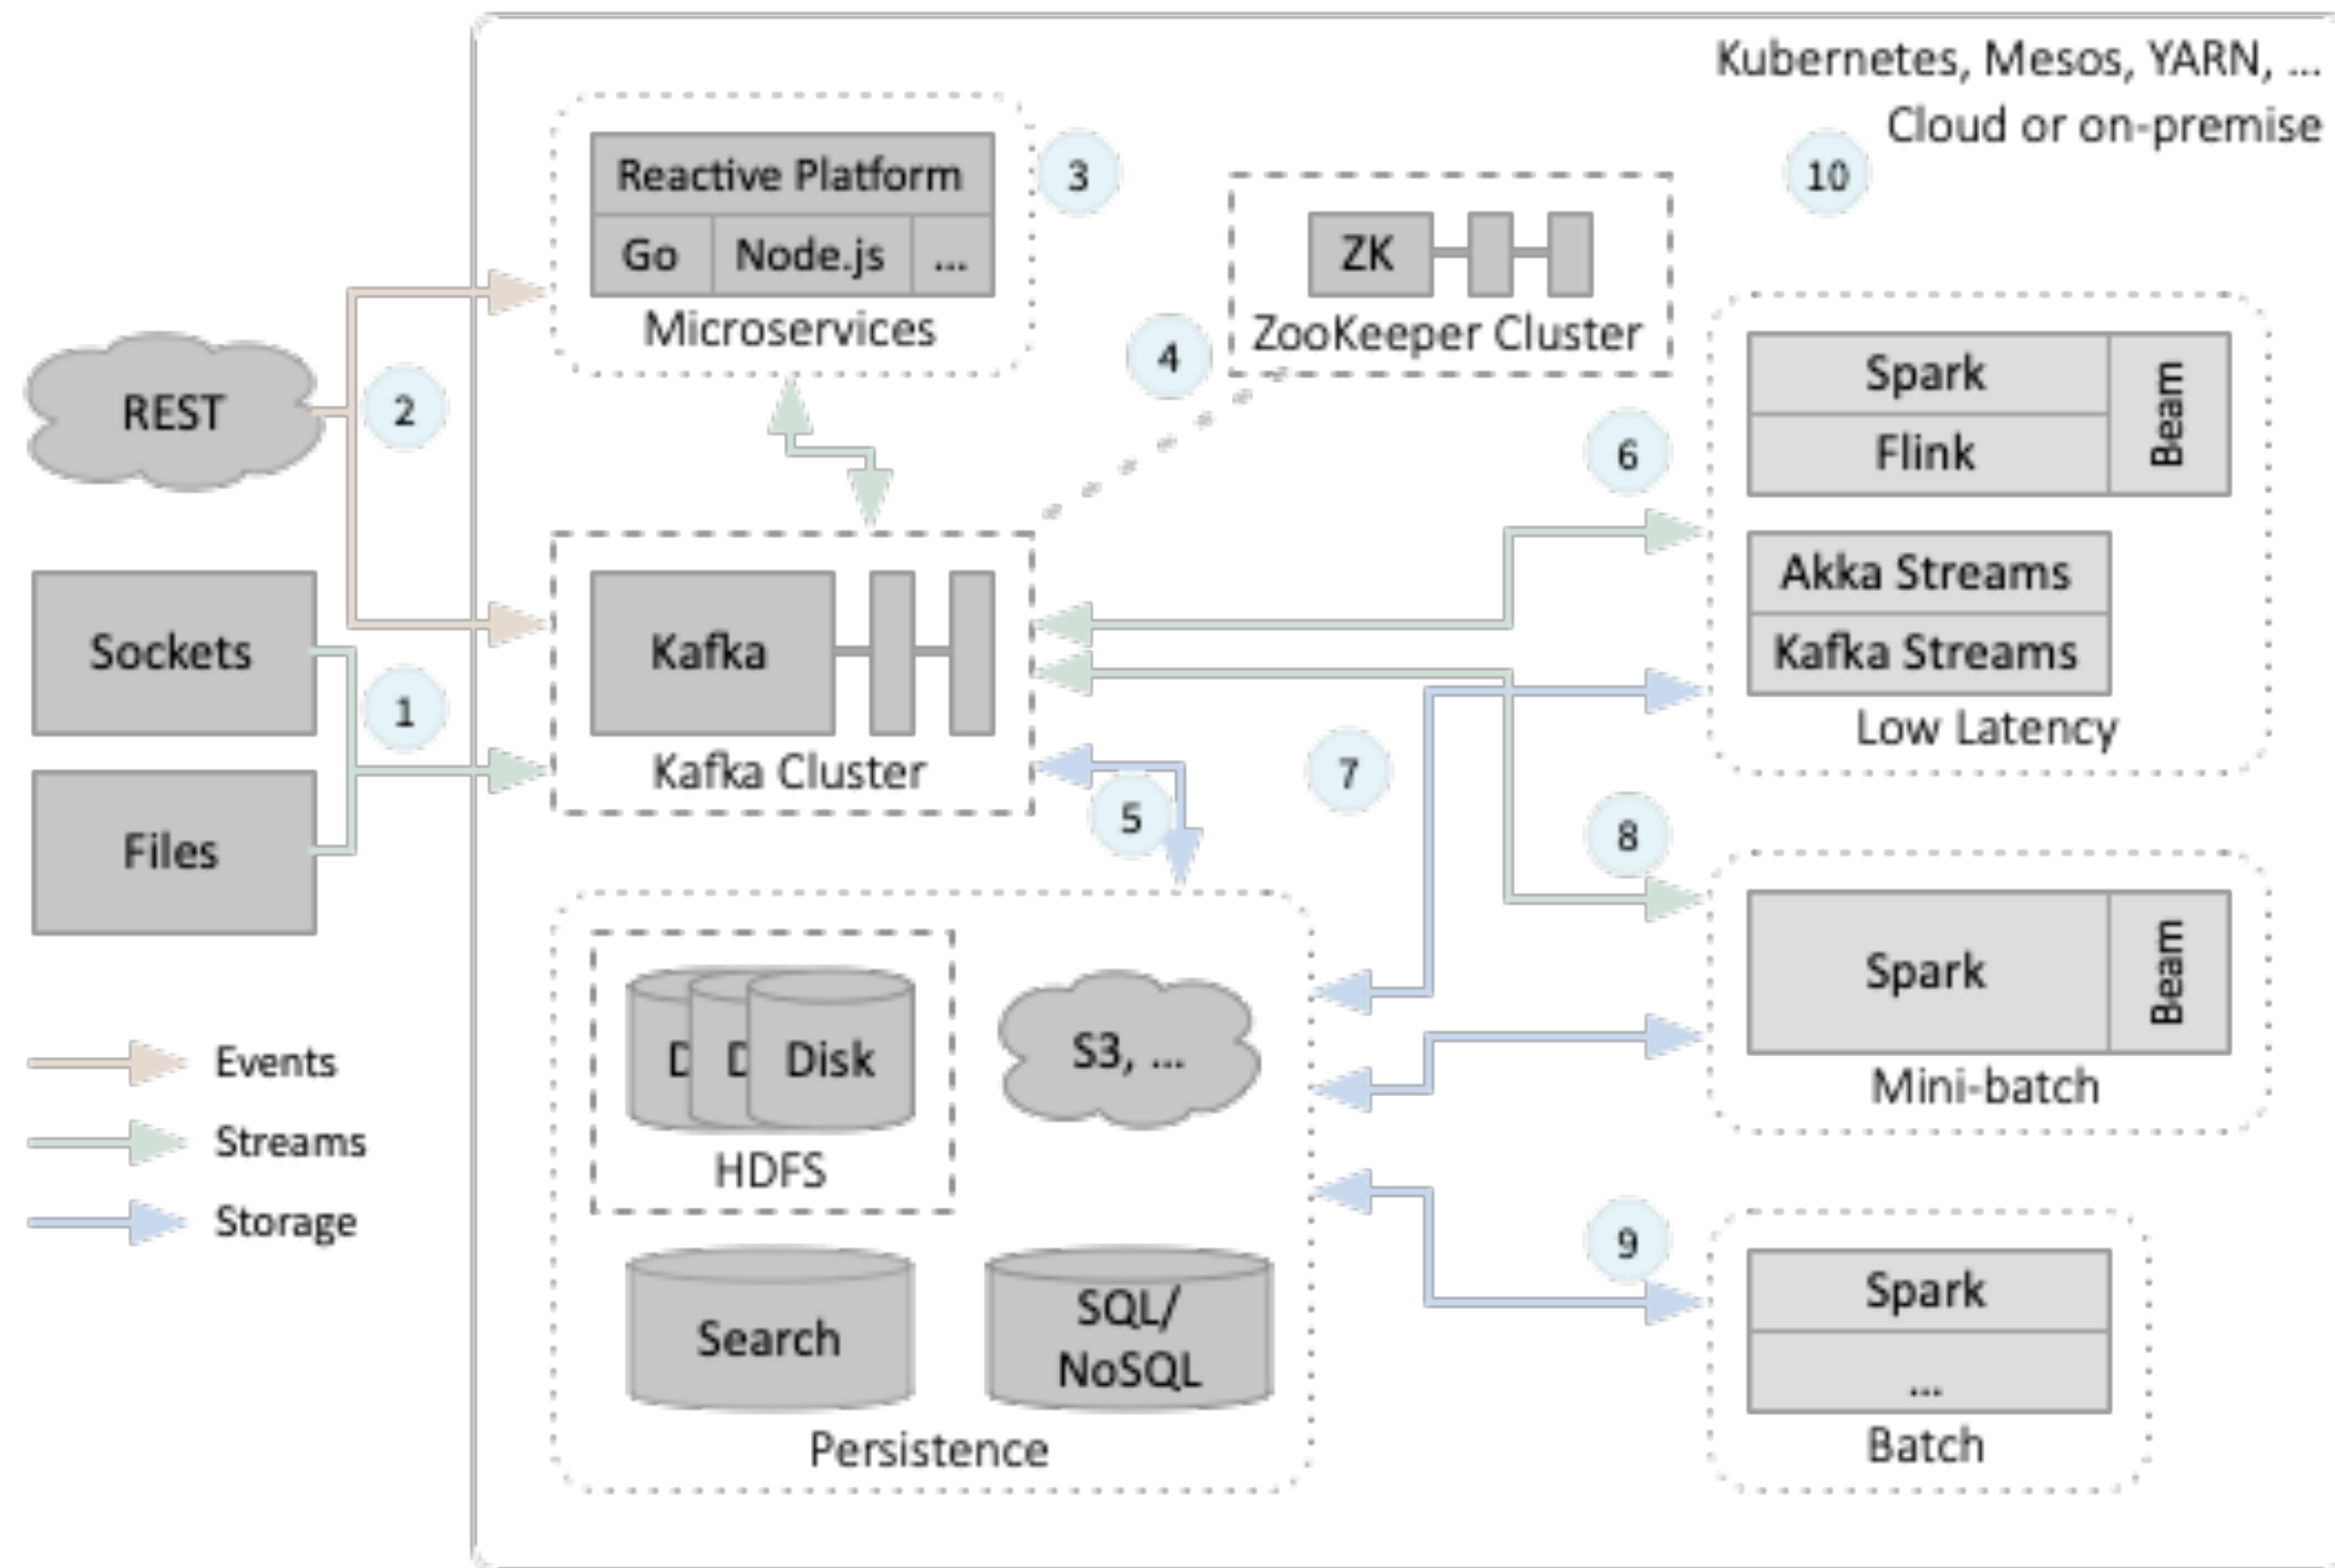
About Kafka

- Handles millions of messages per second, high throughput, high volume
- Distributed and Replicated Commit-log
- Real Time Data Processing
- Stream processing

Kafka Use-Cases

- Activity Tracking
- Page Views
- Click Tracking
- Advertising Patronage
- Feed Artificial Intelligence
- Internet of Things

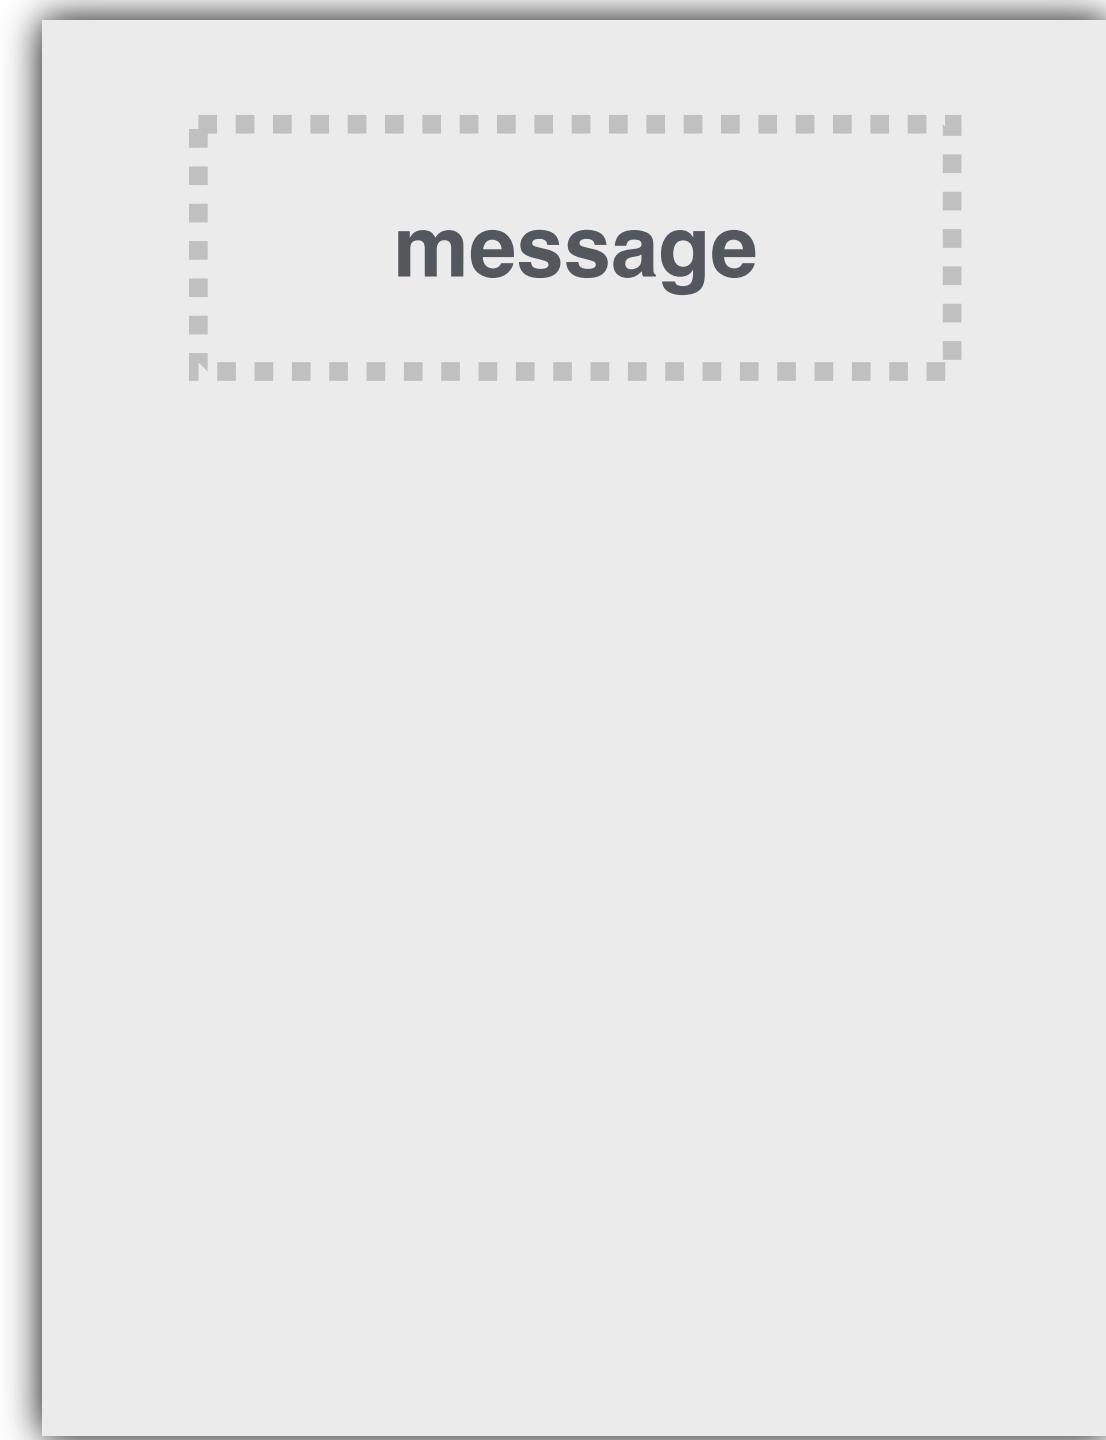
Dean Wampler Architecture



Kafka Messages

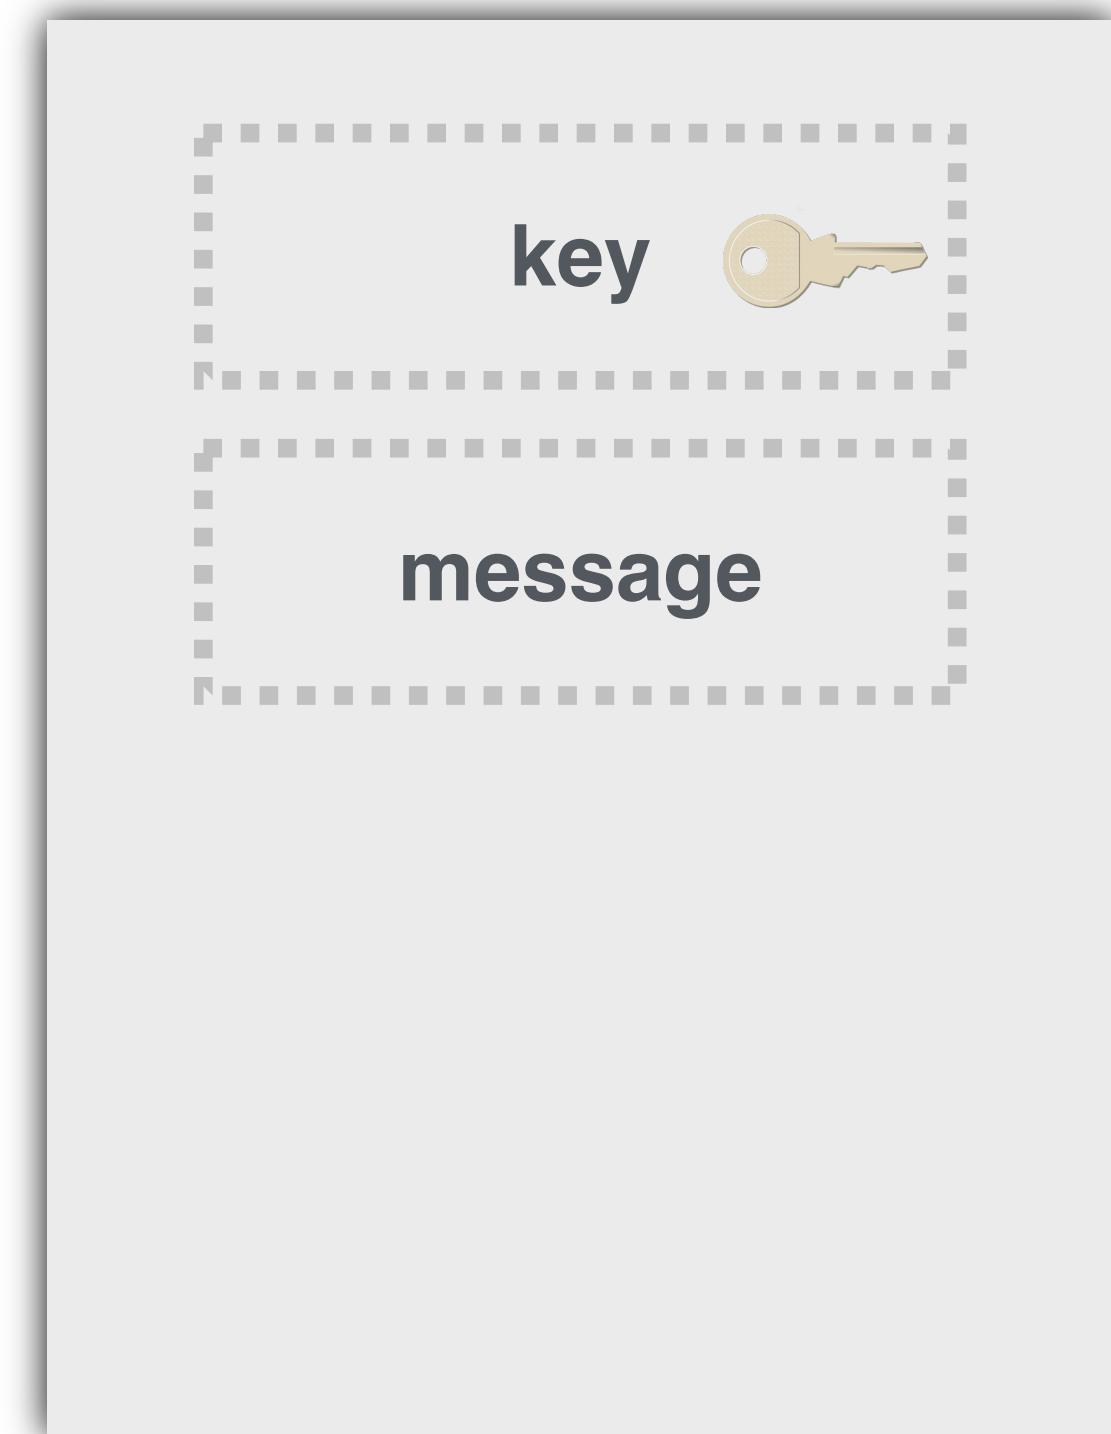
A Kafka Message

- Similar to a *row* or a *record*
- Message is an array of bytes
- No special serialization, that is done at the producer or consumer



A Kafka Message Key

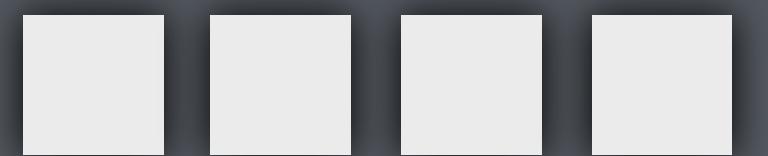
- Message may contain a *key* for better distribution to partitions
- The *key* is also an array of bytes
- If a *key* is provided, a partitioner will hash the key and map it to a single partition
- Therefore it is the only time that something is guaranteed to be in order



Kafka Producers

kafka broker: 0

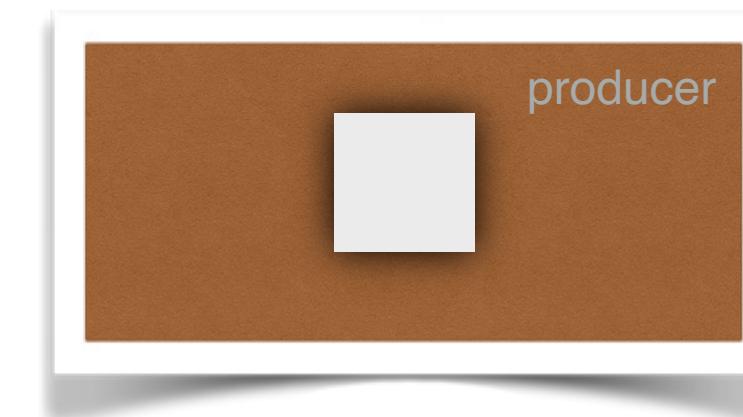
Partition 0:



[0] [1] [2] [3] [4]

topic-a

producer

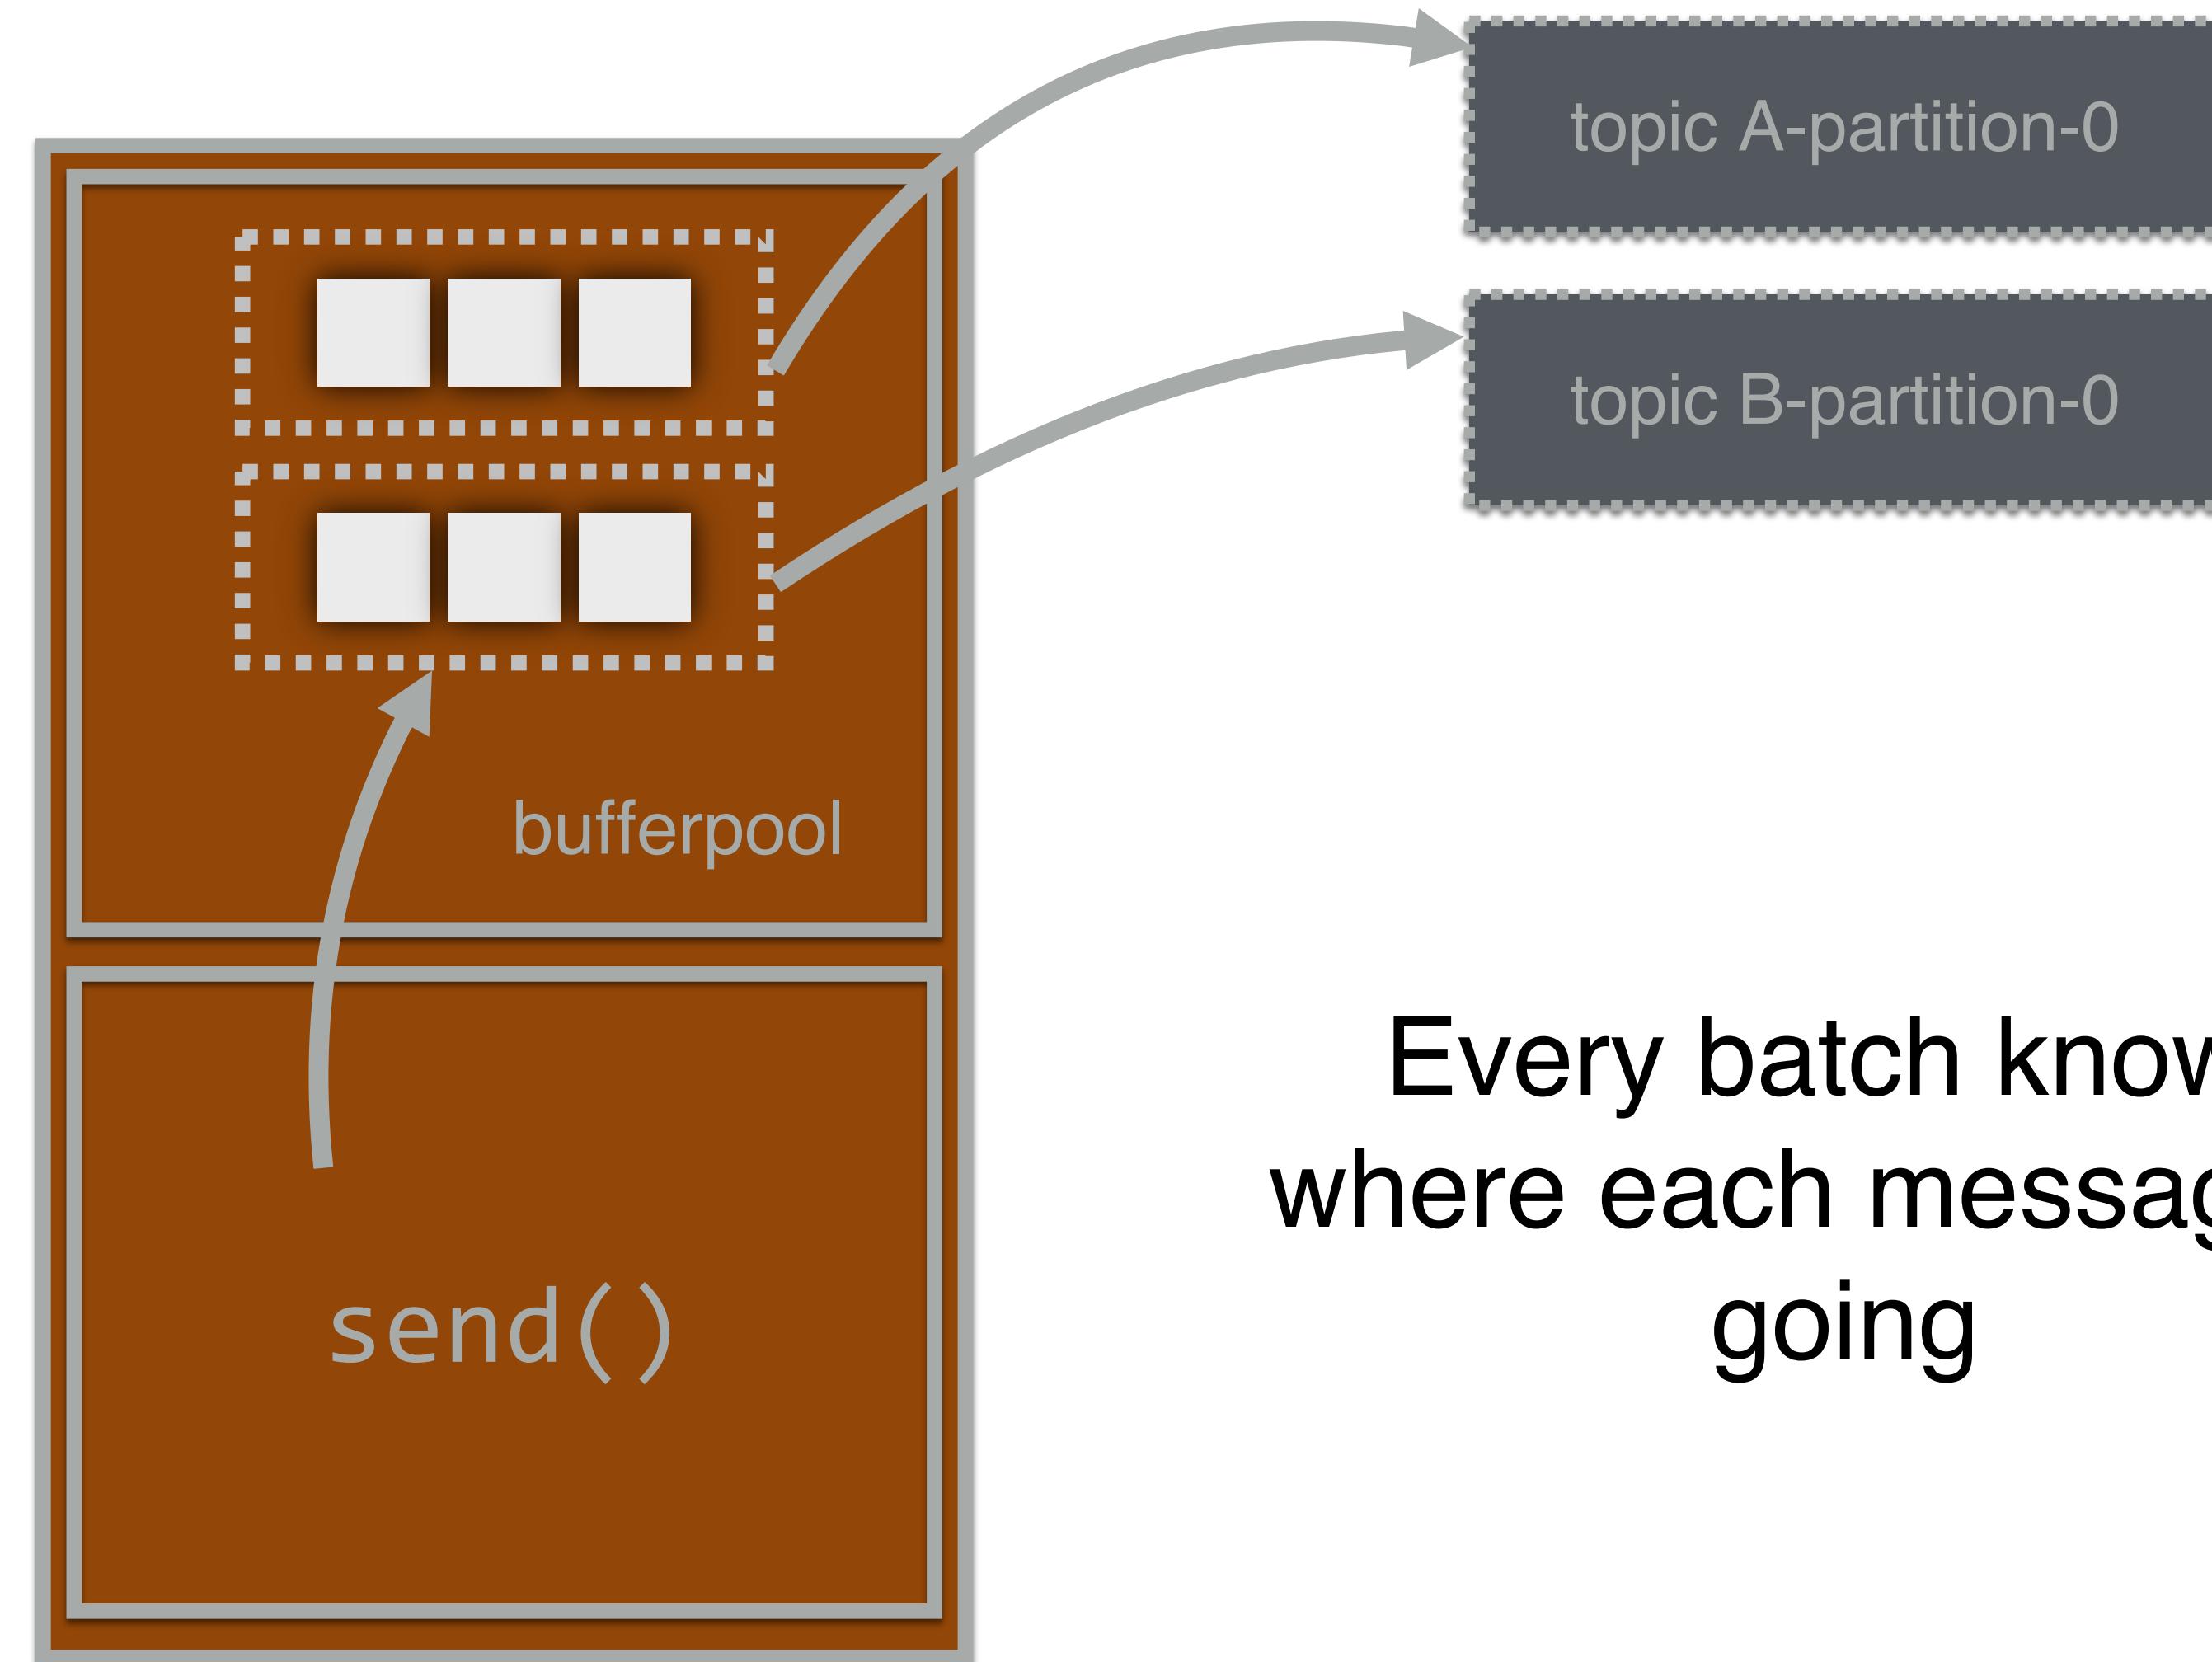


Retention: The data is temporary

Kafka Batch

- A collection of messages, that is sent, configured in *bytes*
- Sent to the same *topic* and the same *partition*
- *Avoids overhead of sending multiple message over the wire*





Every batch knows
where each message is
going

$\text{murmur2(bytes)} \% \text{ number partitions}$

This could be changed with custom partitioner

~~murmur2(bytes) % number_partitions~~

```
if (key == "CSCO") {  
    else return 0;  
} else {  
    (murmur2(bytes) % number_partitions - 1) + 1  
}
```

A-E



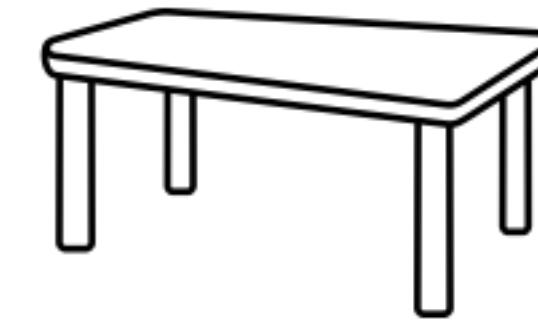
F-K



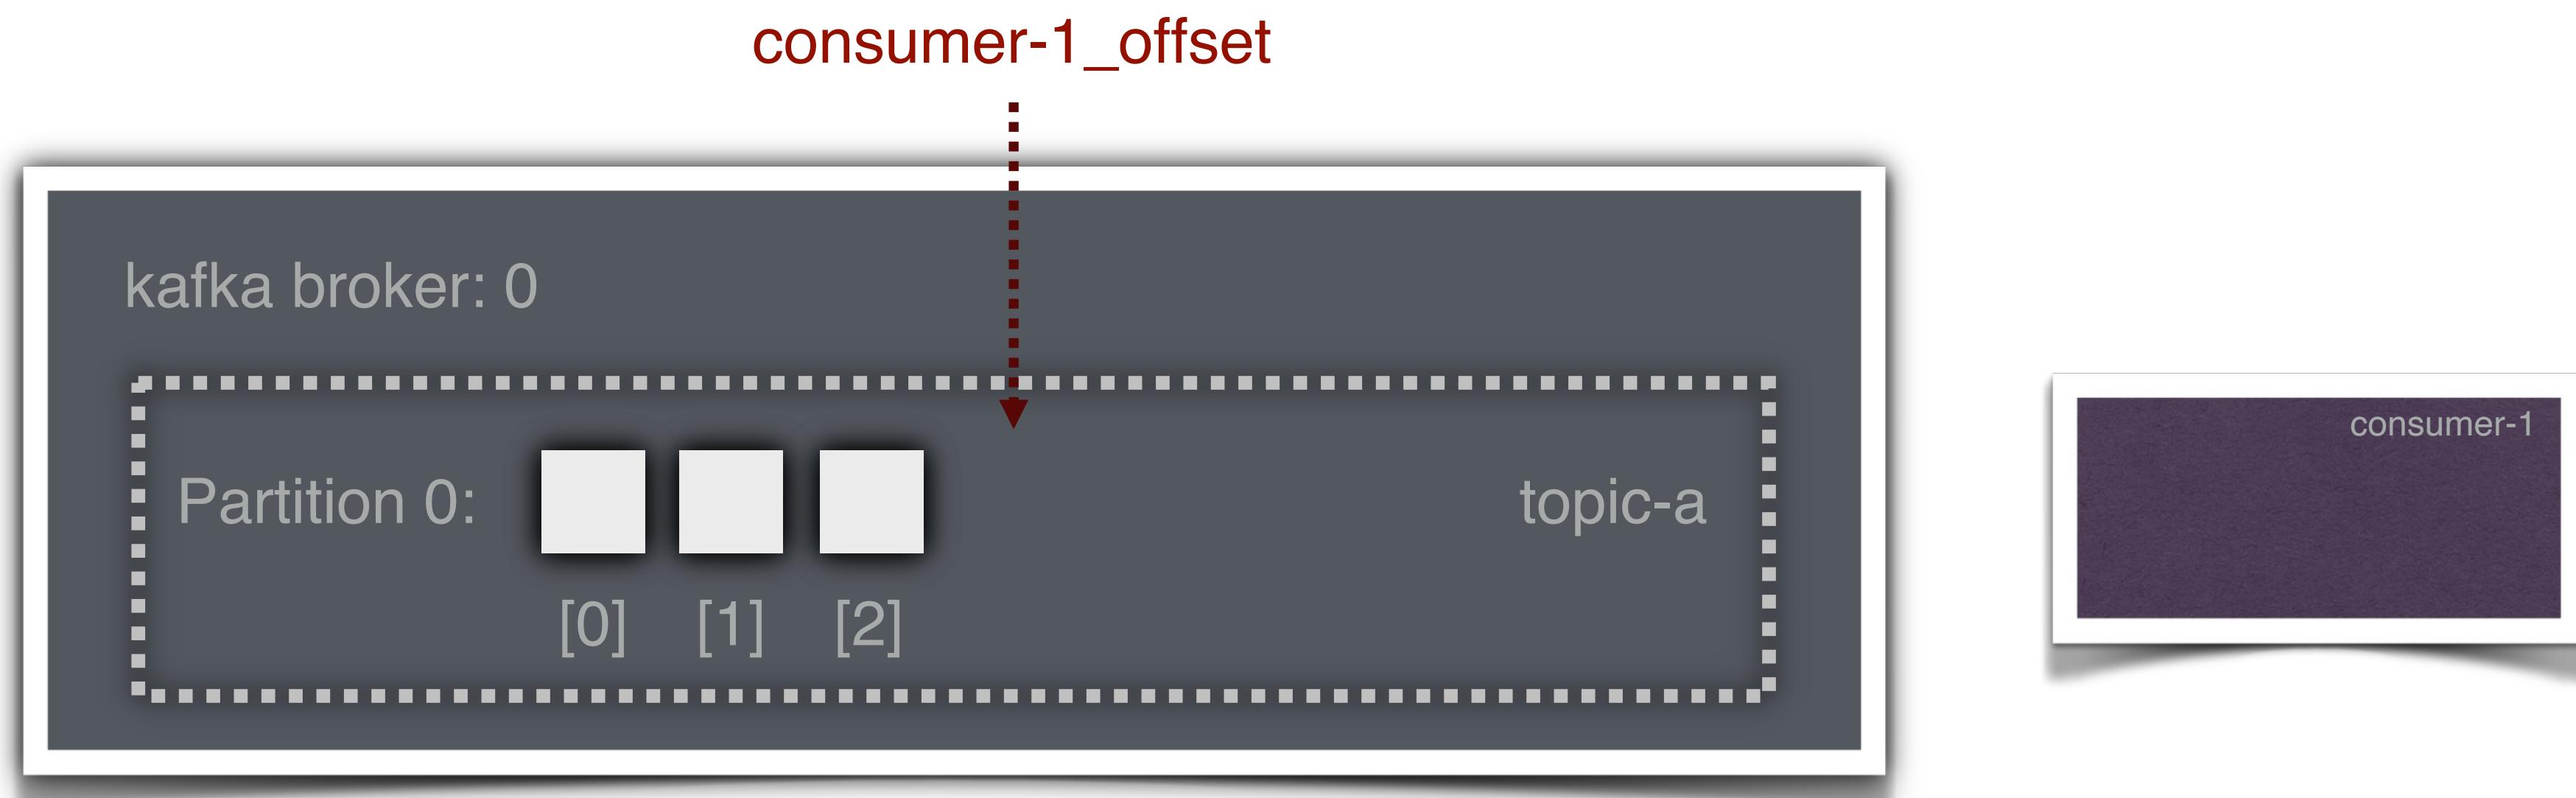
L-S

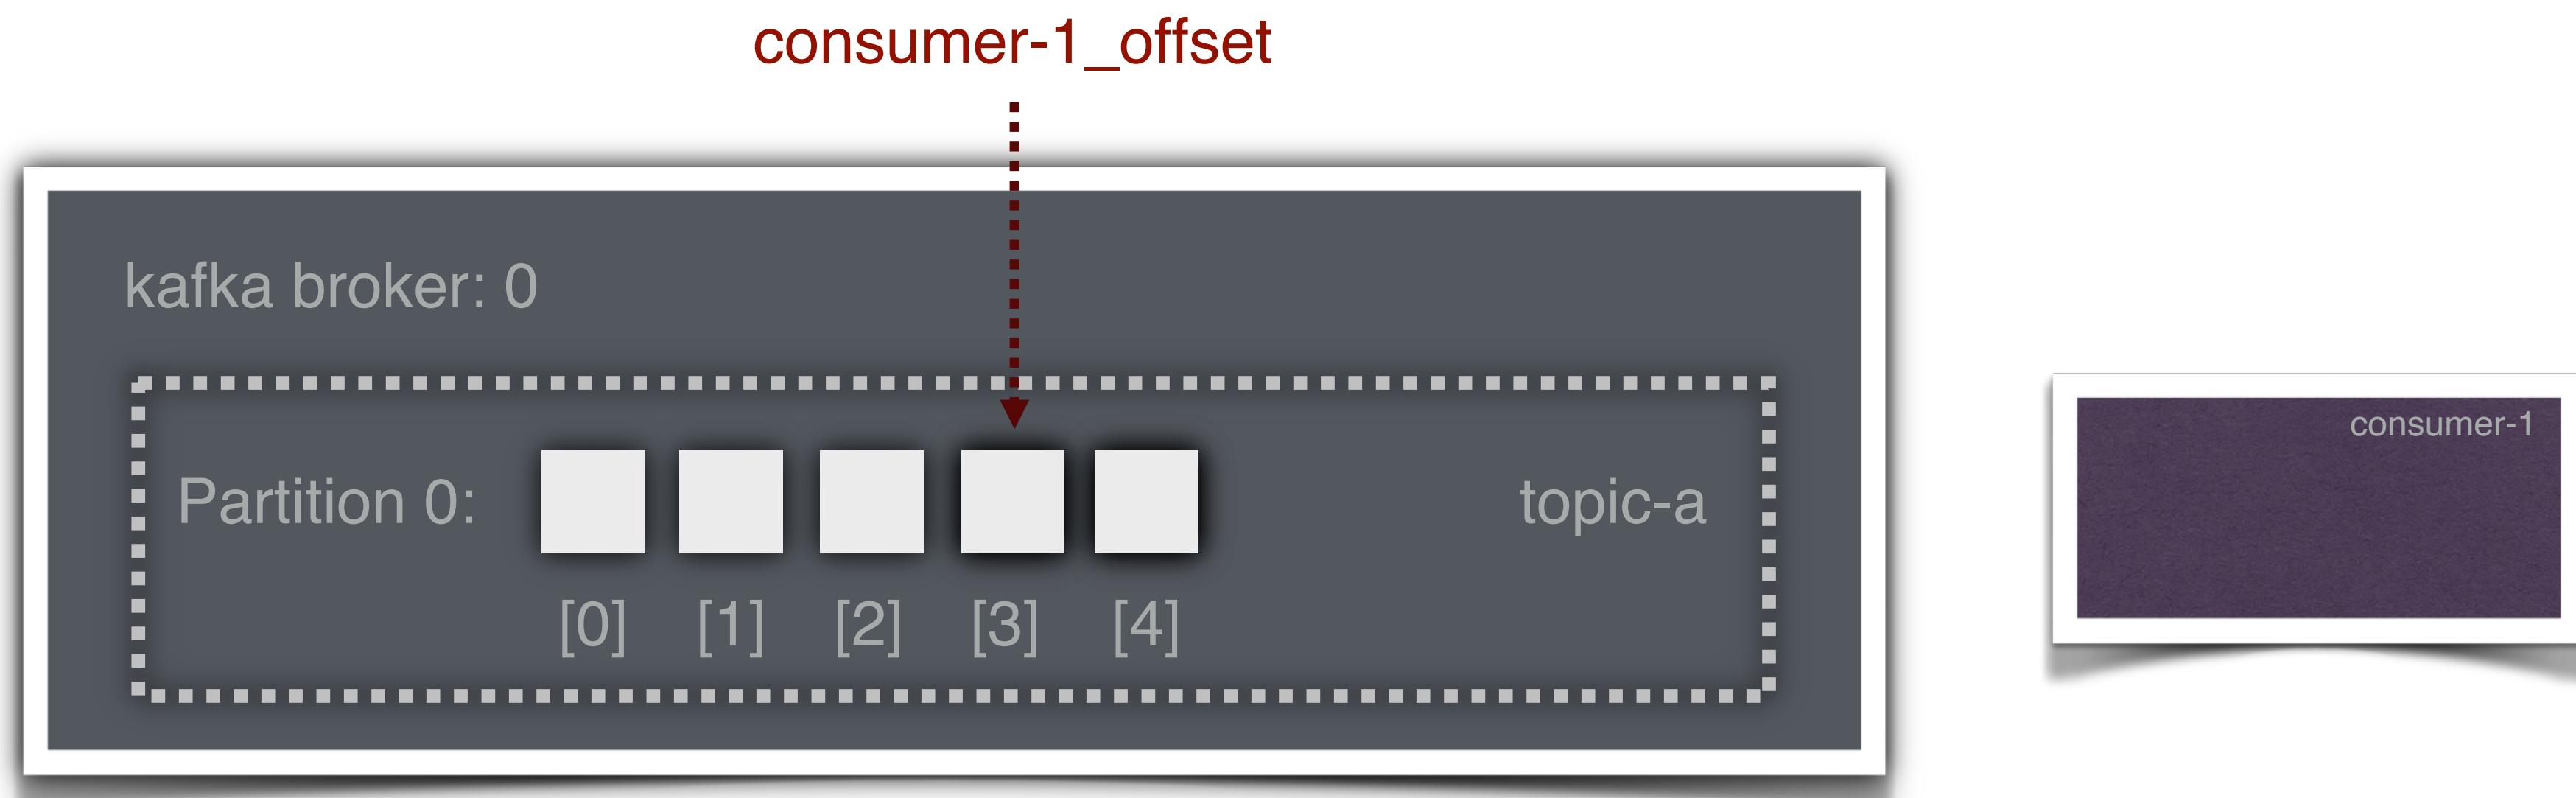


T-Z

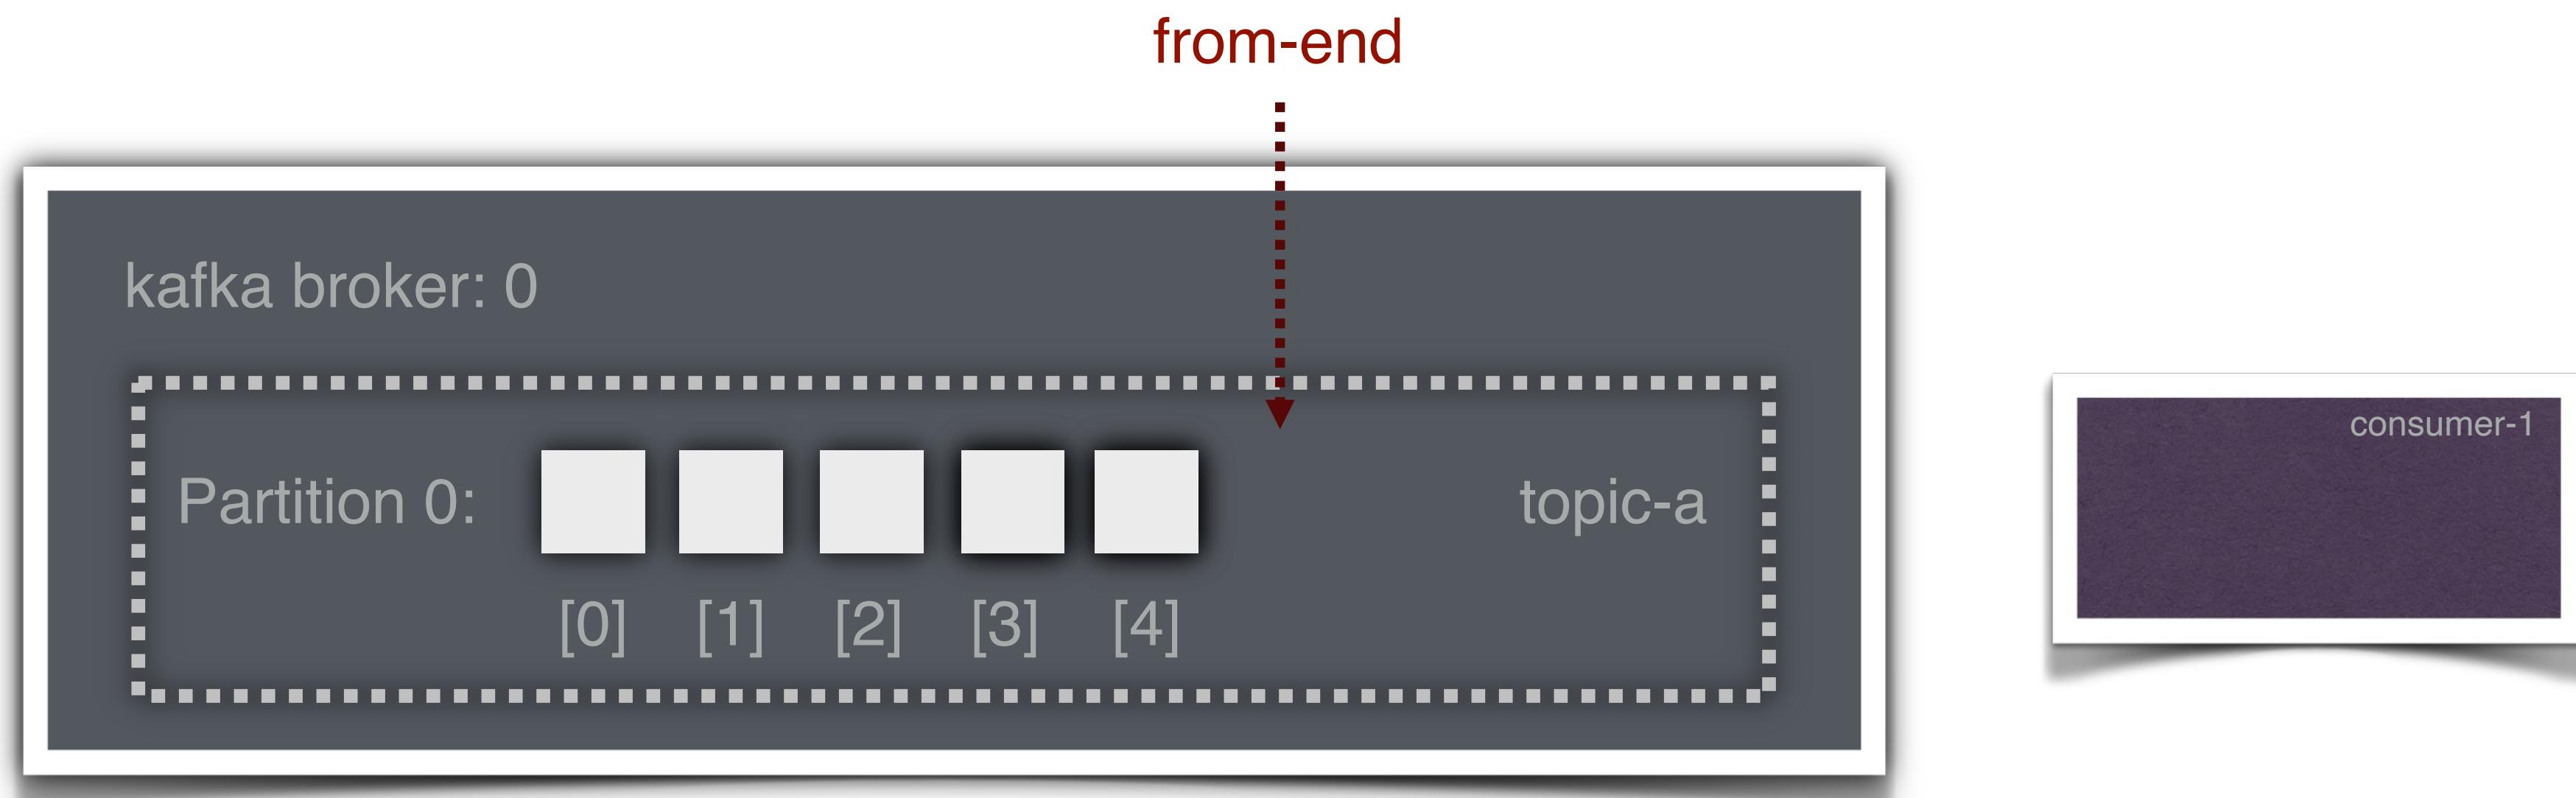


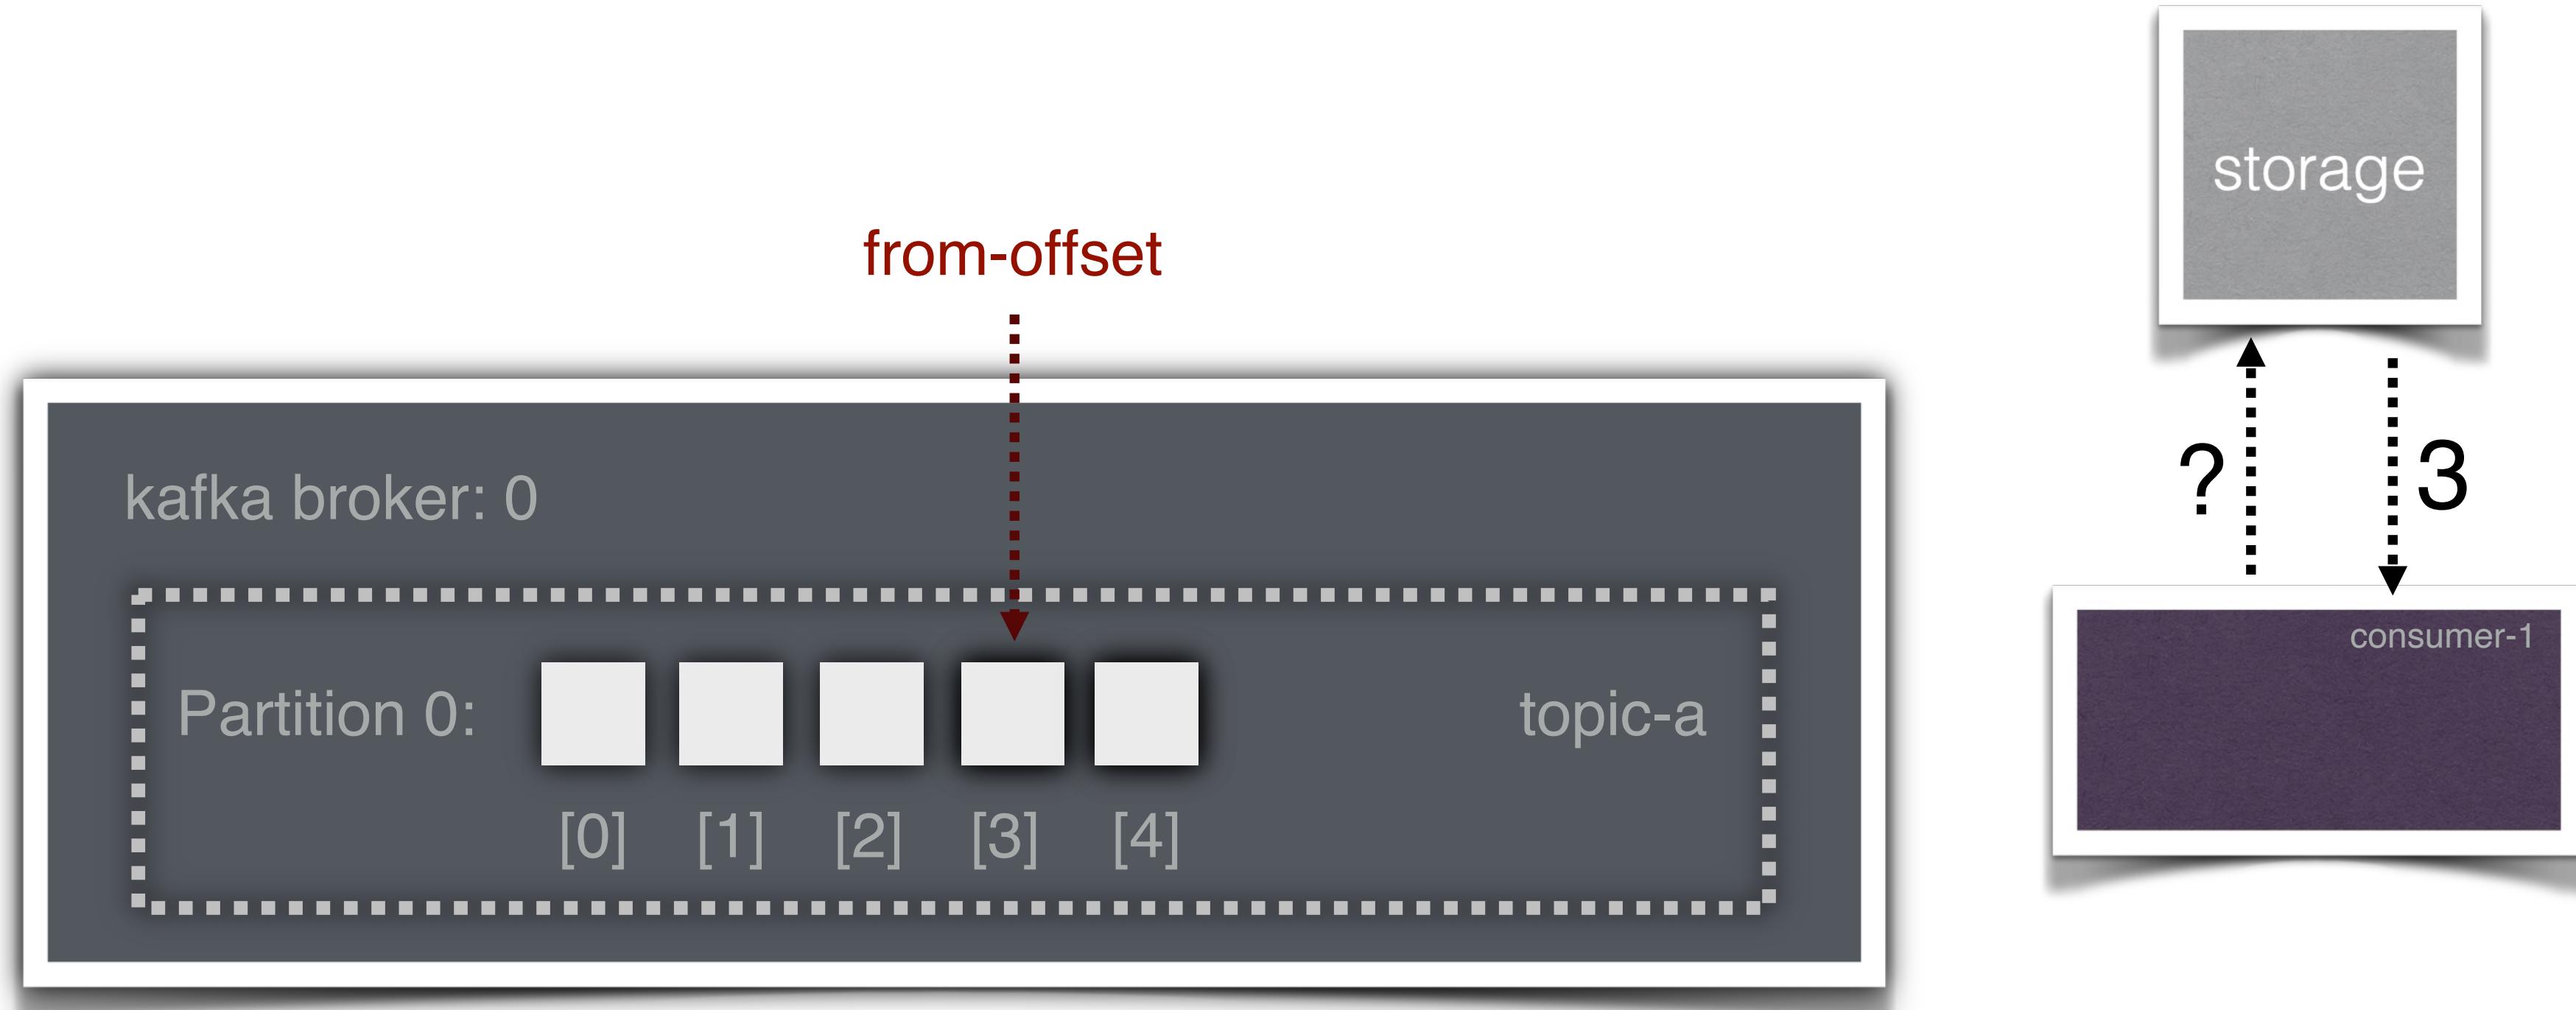
Kafka Consumers



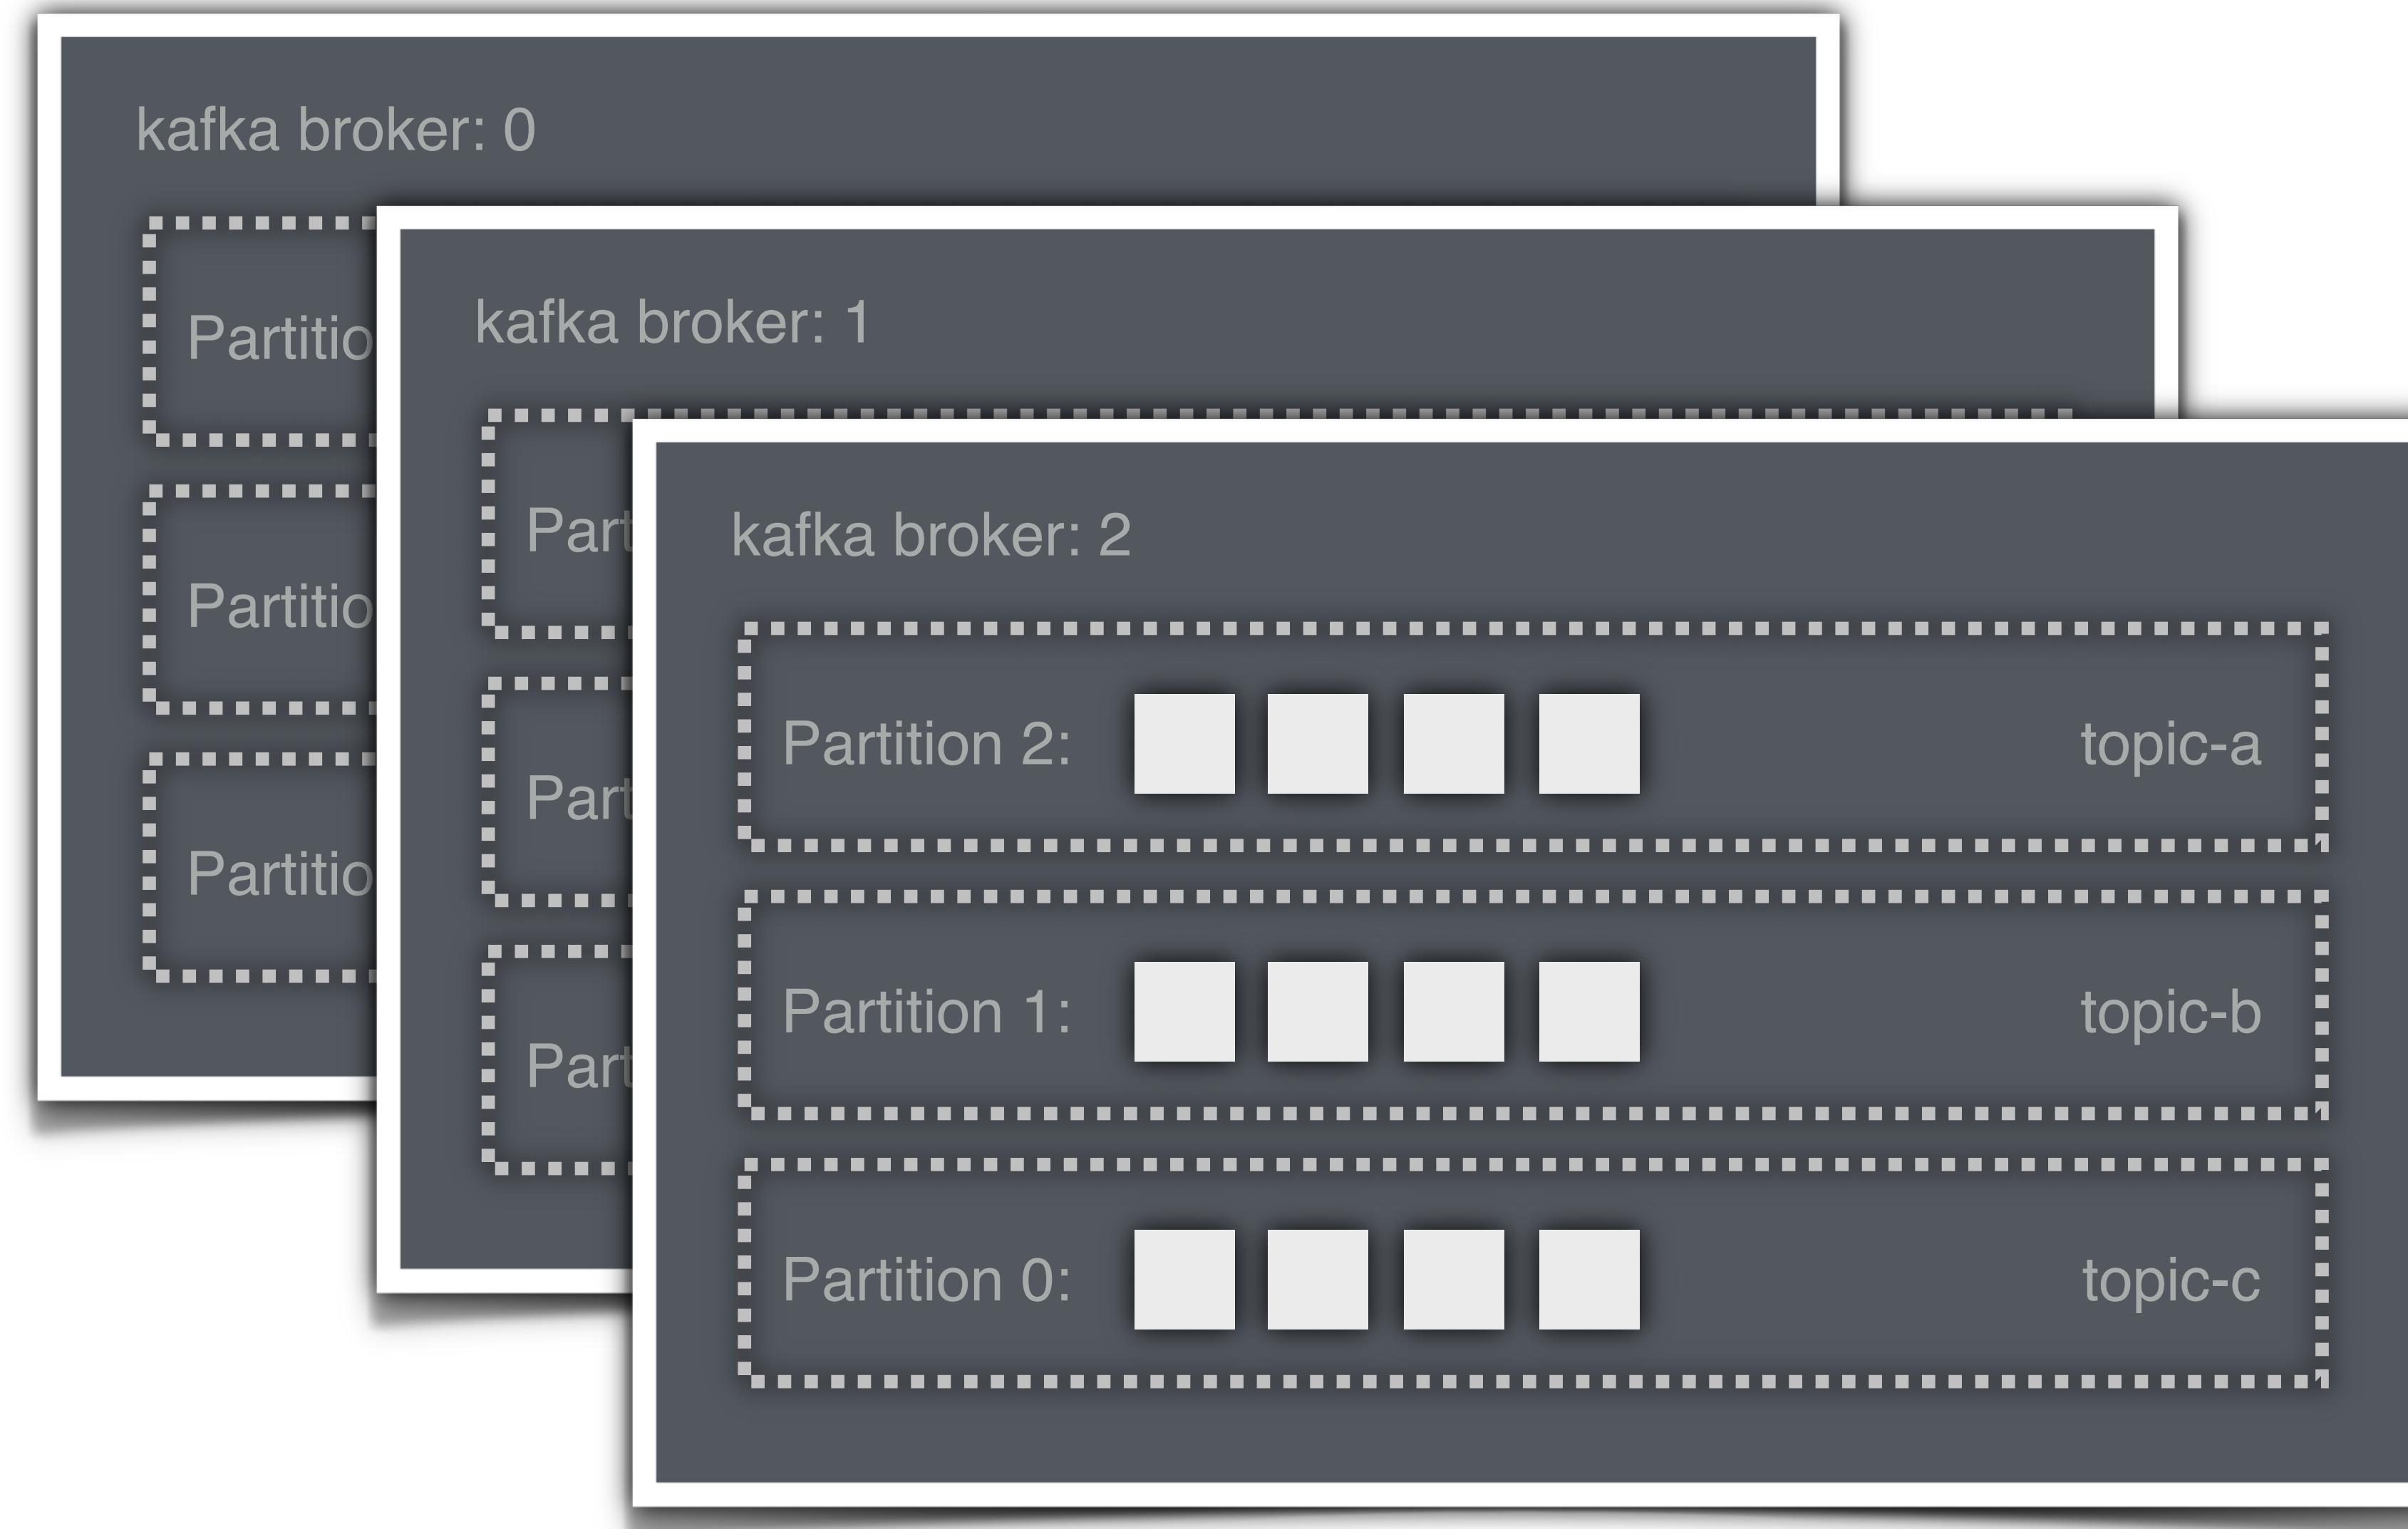








Kafka Partitions



Each partition is on a different broker,
therefore a single topic is scaled

kafka broker

Partition 0

Partition 1

Partition 2

Partition 3

Partition 4

kafka broker

Partition 0

Partition 1

Partition 2

Partition 3

Partition 4

kafka broker: 2

Partition 0:  topic-a

Partition 2:  topic-a

Partition 1:  topic-b

Partition 0:  topic-c

Partition 2:  topic-c

Partition 1:  topic-a

Partition 0:  topic-b

Partition 2:  topic-b

Partition 1:  topic-c

A-E



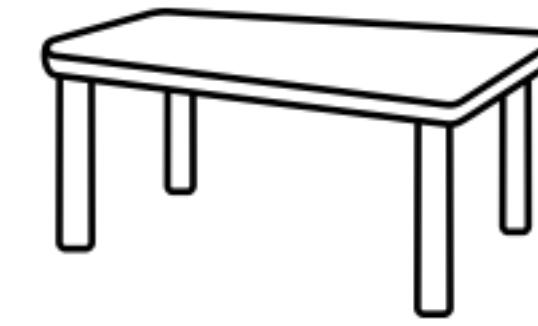
F-K



L-S



T-Z



Kafka Pub/Sub Comparisons

Comparison Pub/Sub

- Kafka as the distributed "git log"
- Ability to replayed in a consistent manner
- Kafka is also stored durable, in order, and deterministic, if key is available
- Data can also be distributed for resiliency
- Scalable and Elastic

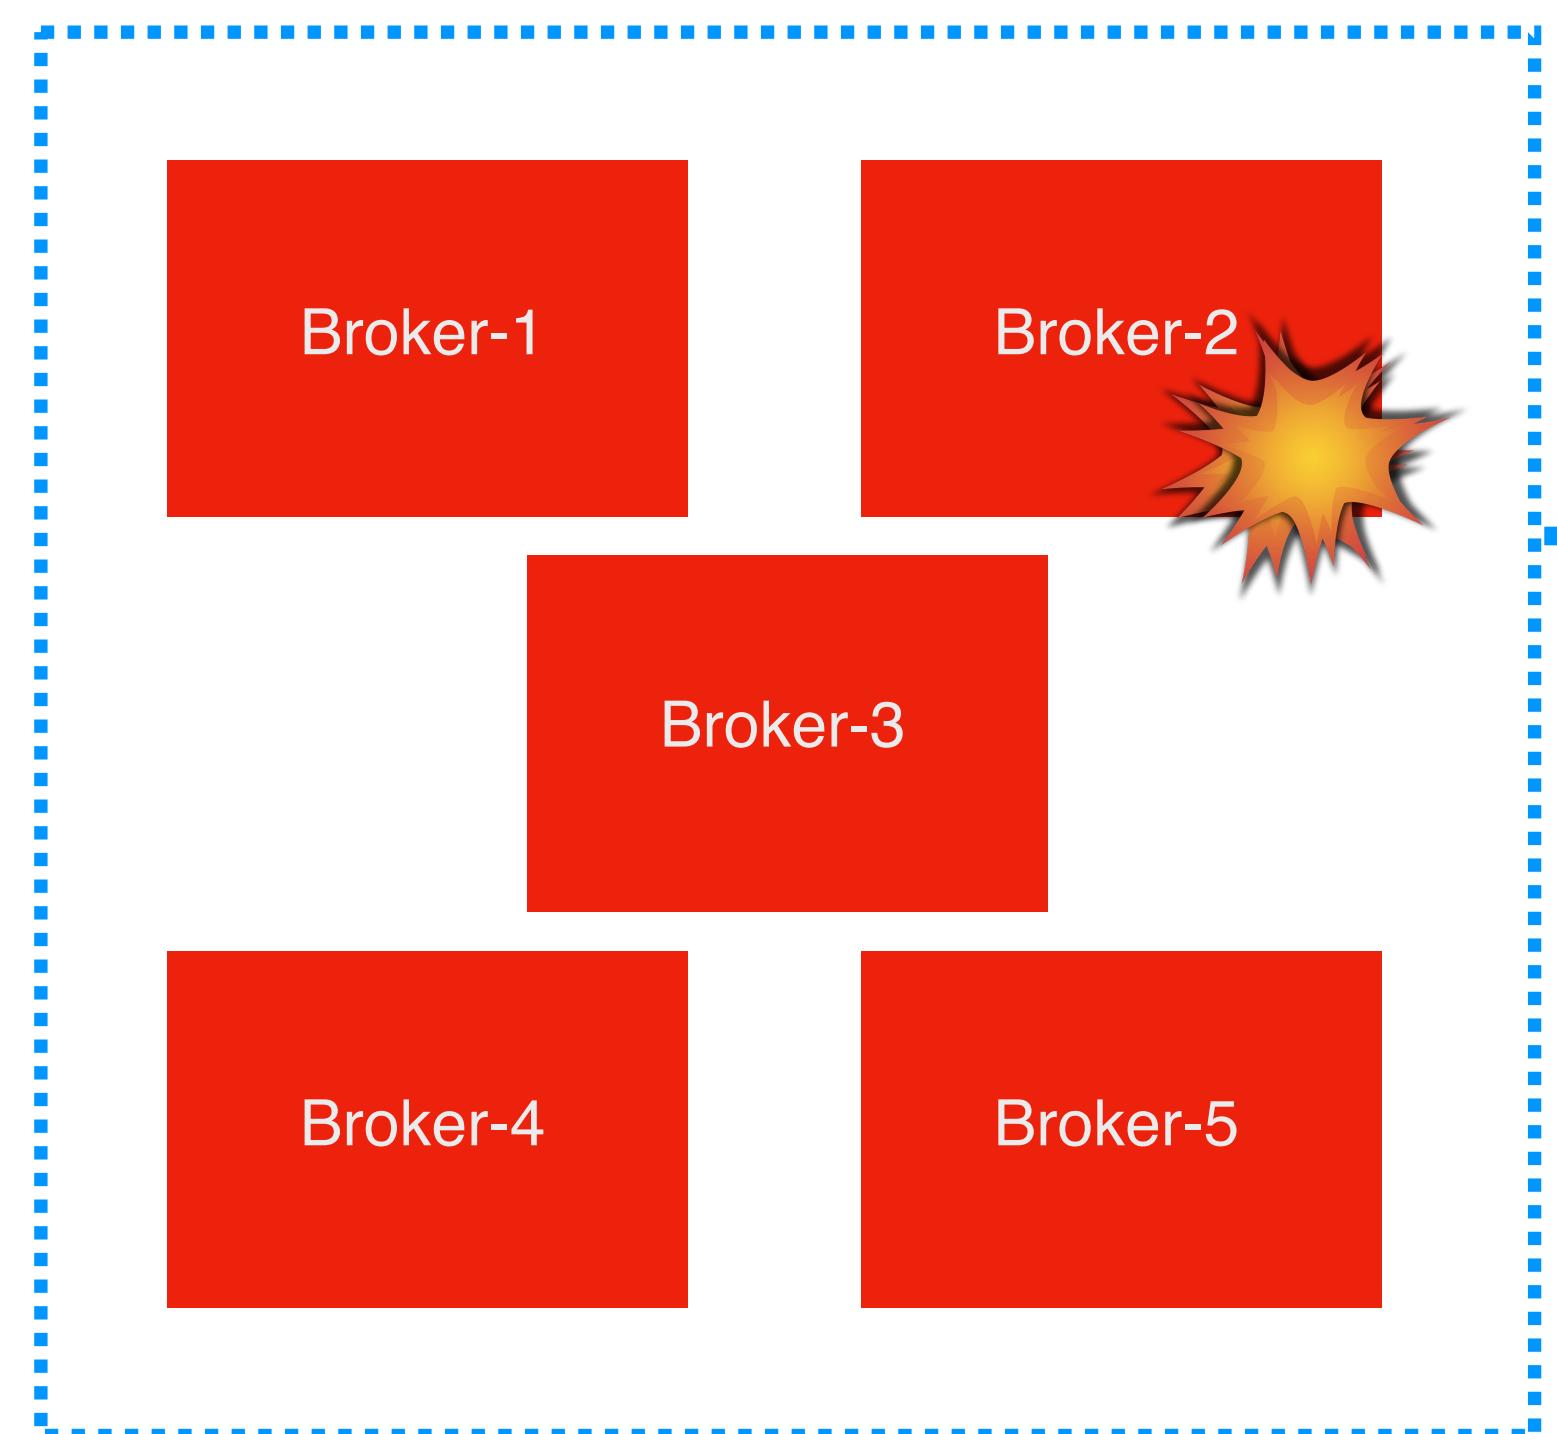
Zookeeper

Zookeeper Essence

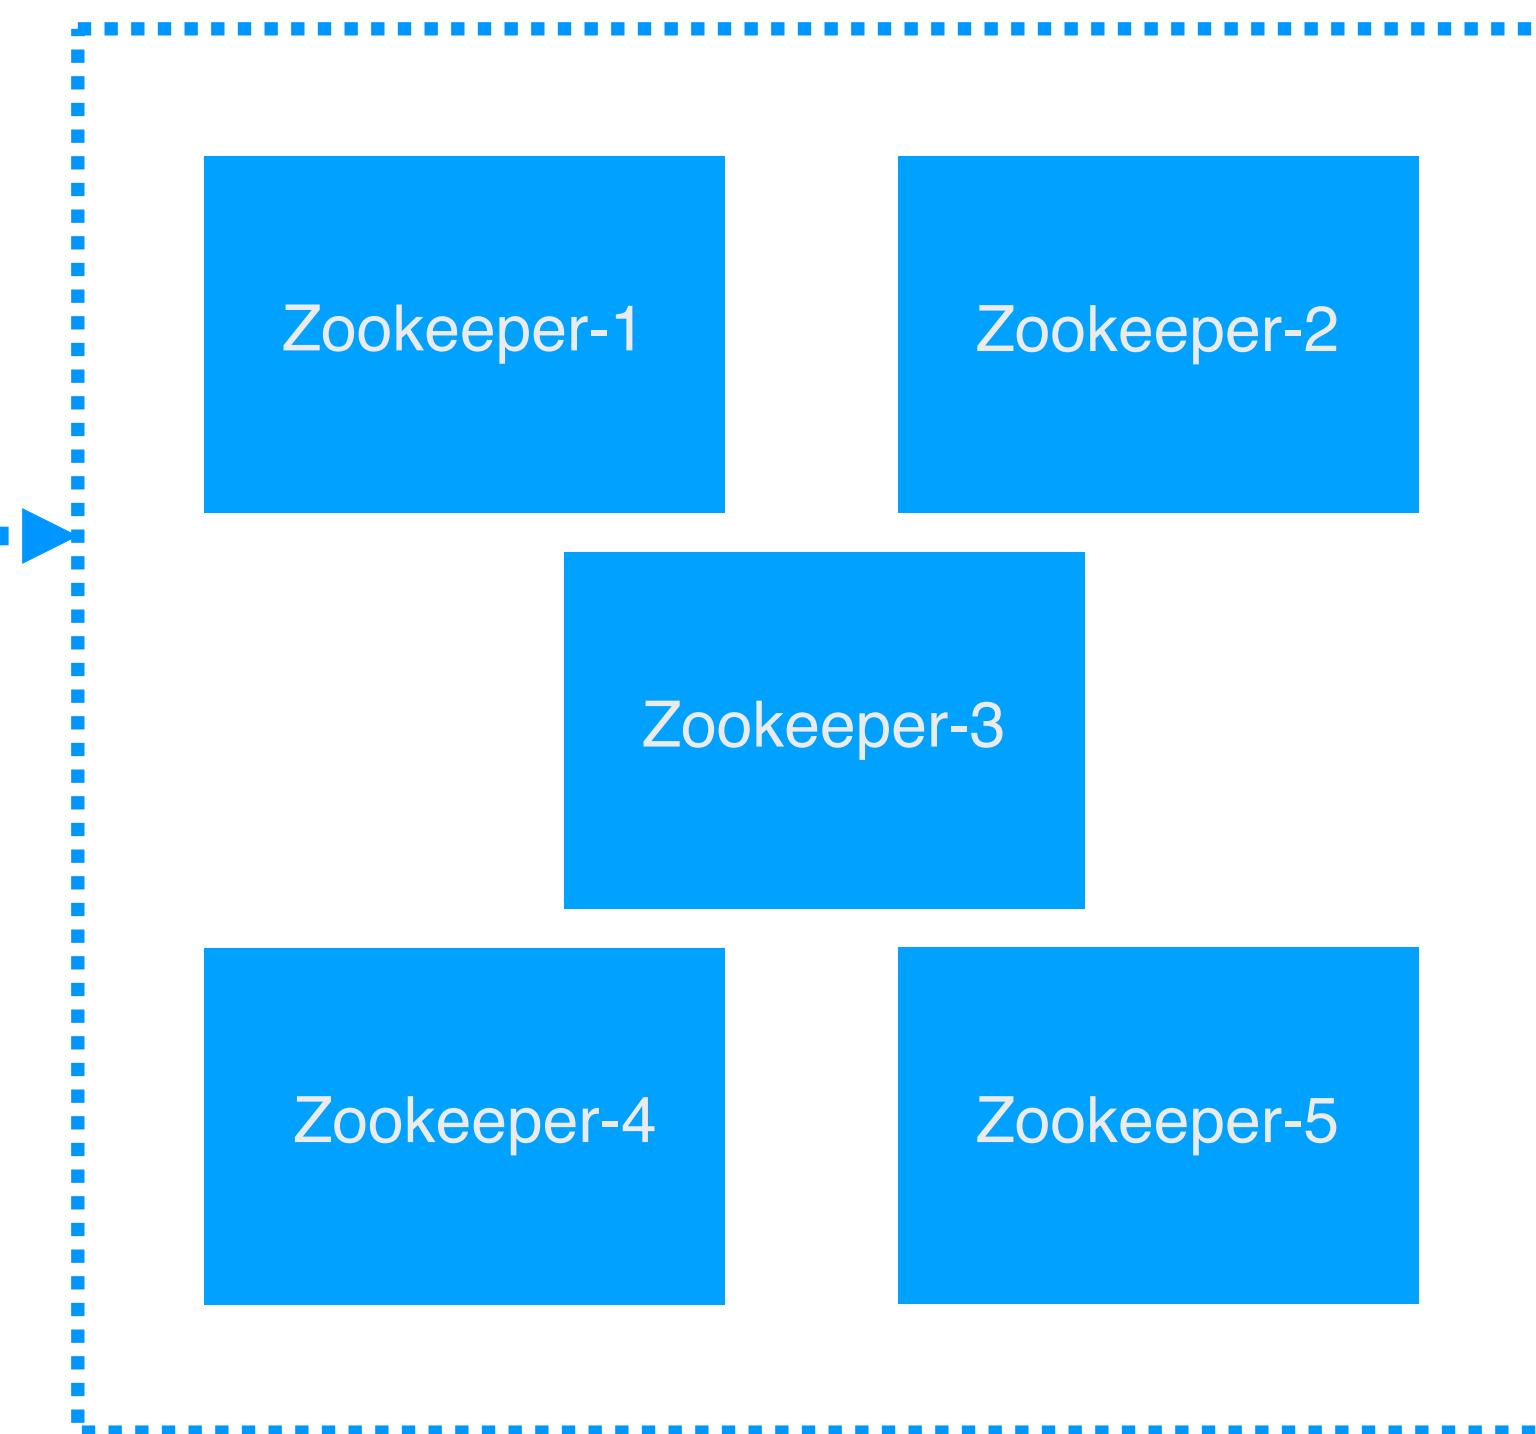
- Centralized Coordination Service
- Maintains:
 - Metadata
 - Naming
 - Configuration



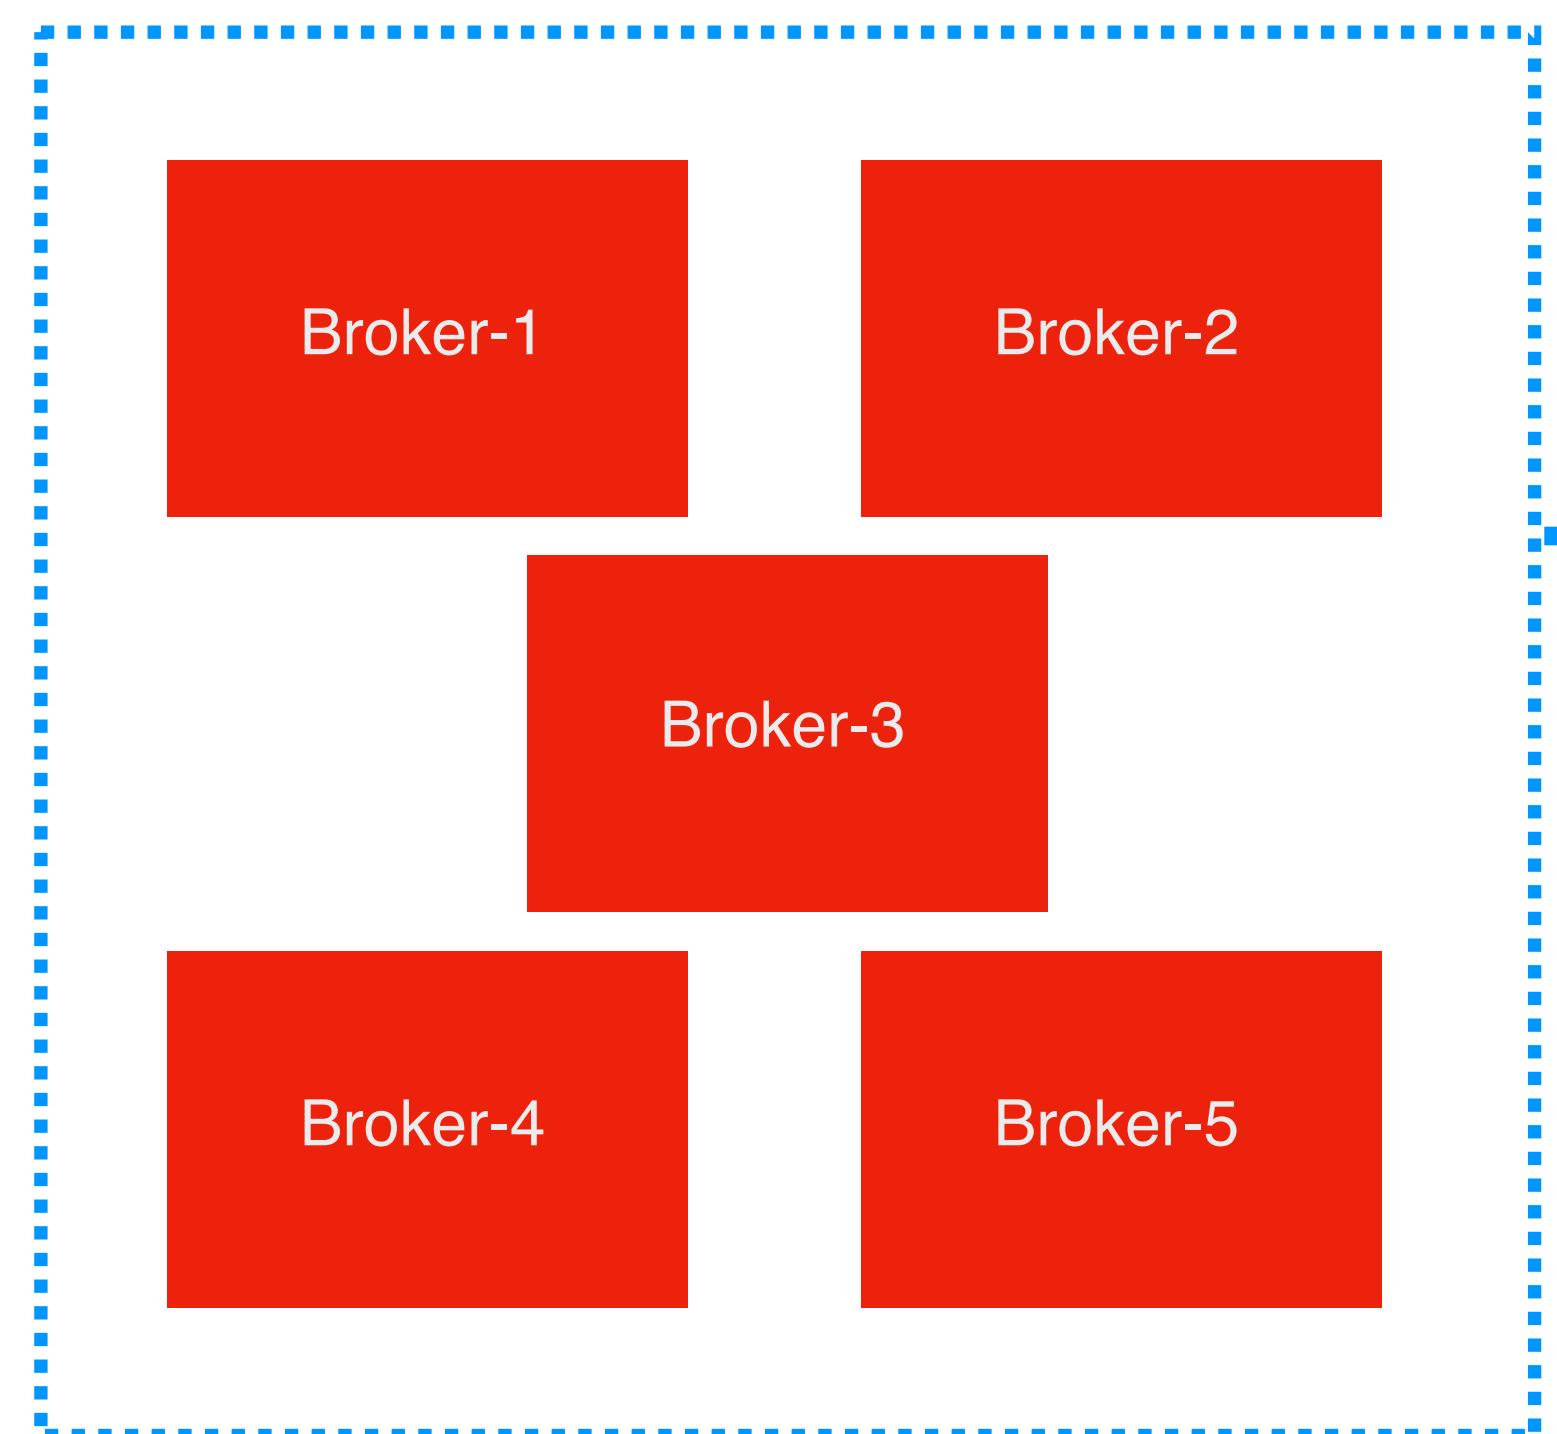
Kafka Cluster



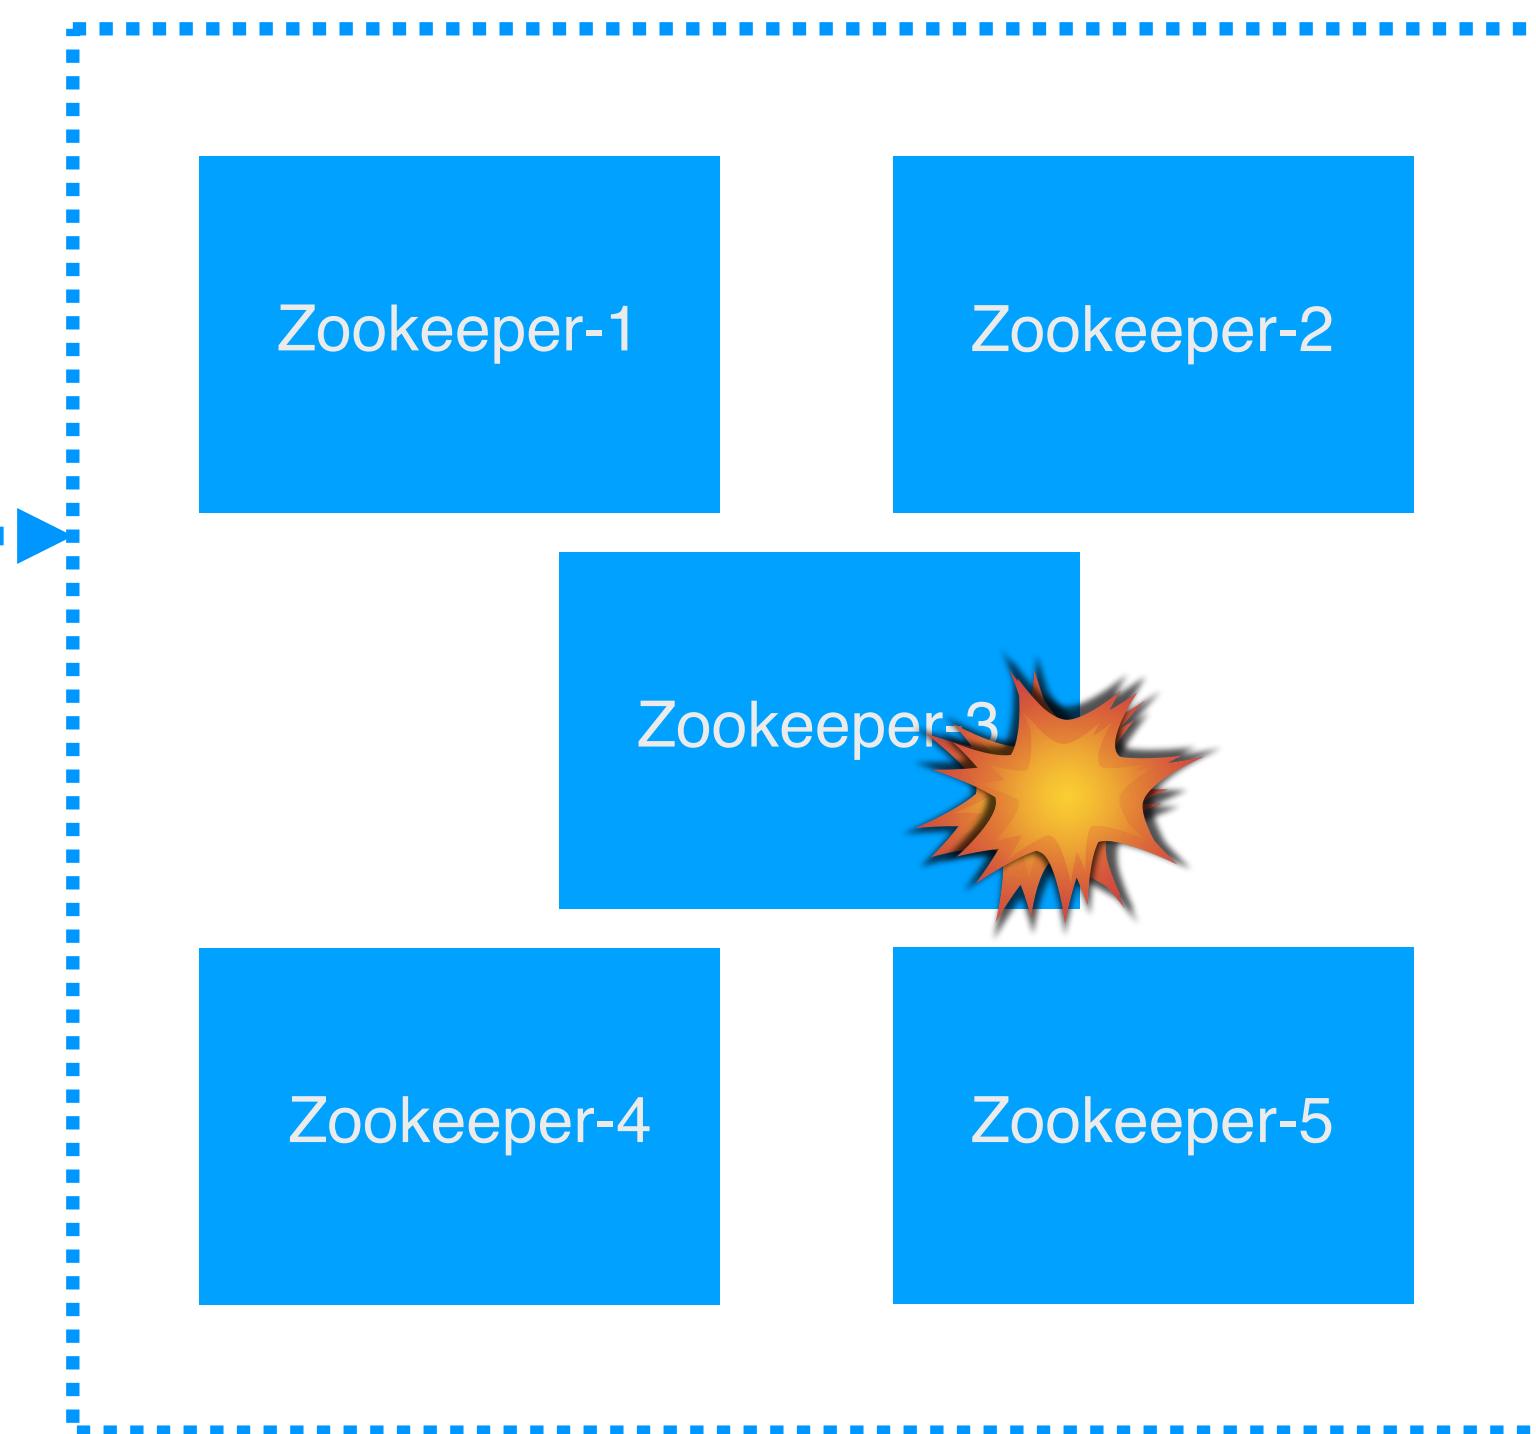
Zookeeper Ensemble



Kafka Cluster



Zookeeper Ensemble



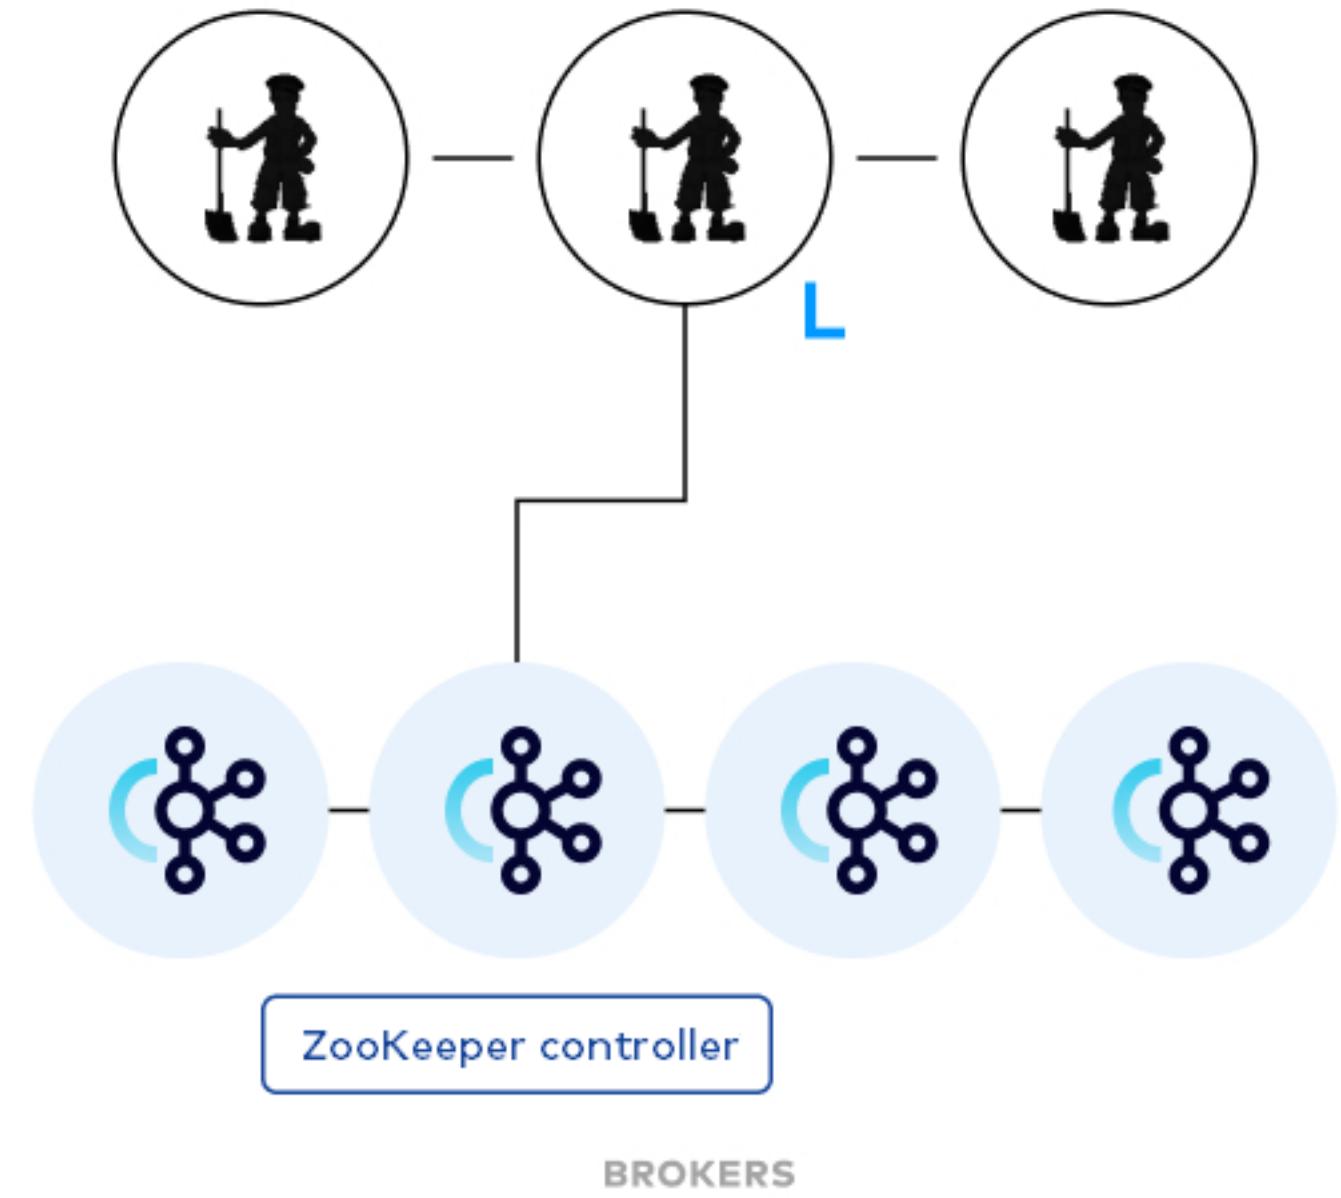
Kafka without ZooKeeper

- Kafka Improvement Proposal (KIP) 500 introduced idea to remove ZooKeeper
- Motivation:
 - Remove dependency on ZooKeeper, which is a project maintained by different people
 - Metadata, cluster membership and configuration management can be done by Kafka itself
 - Simpler deployment and configuration
 - <https://cwiki.apache.org/confluence/display/KAFKA/KIP-500%3A+Replace+ZooKeeper+with+a+Self-Managed+Metadata+Quorum#KIP500:ReplaceZooKeeperwithaSelfManagedMetadataQuorum-Motivation>

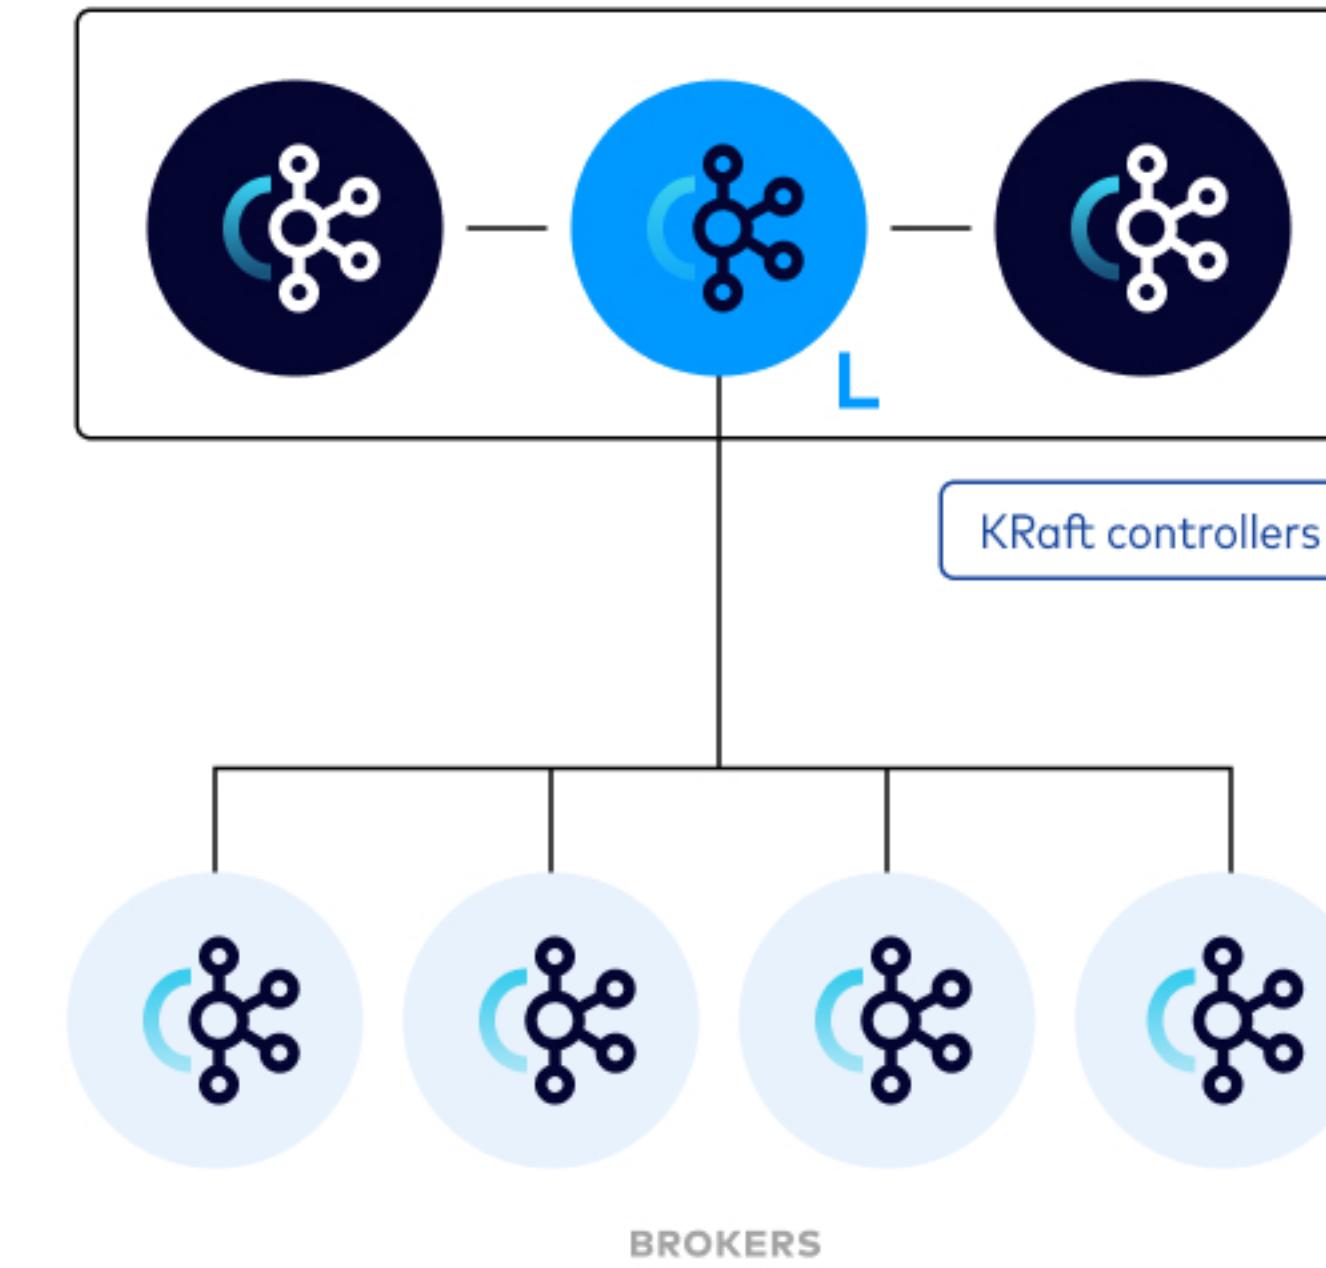
Kafka Raft (KRaft) mode

- Production ready since Kafka 3.3.0
- A new consensus protocol that replaces ZooKeeper
- KRaft mode is a new mode that replaces ZooKeeper
 - ZooKeeper is still needed for now
 - KRaft mode is still in development
- How it works:
 - Kafka brokers form a Raft quorum
 - Raft is a consensus algorithm that is simpler than ZooKeeper
 - Raft is used to manage metadata and configuration
 - ZooKeeper is still used for ACLs and other features

ZooKeeper



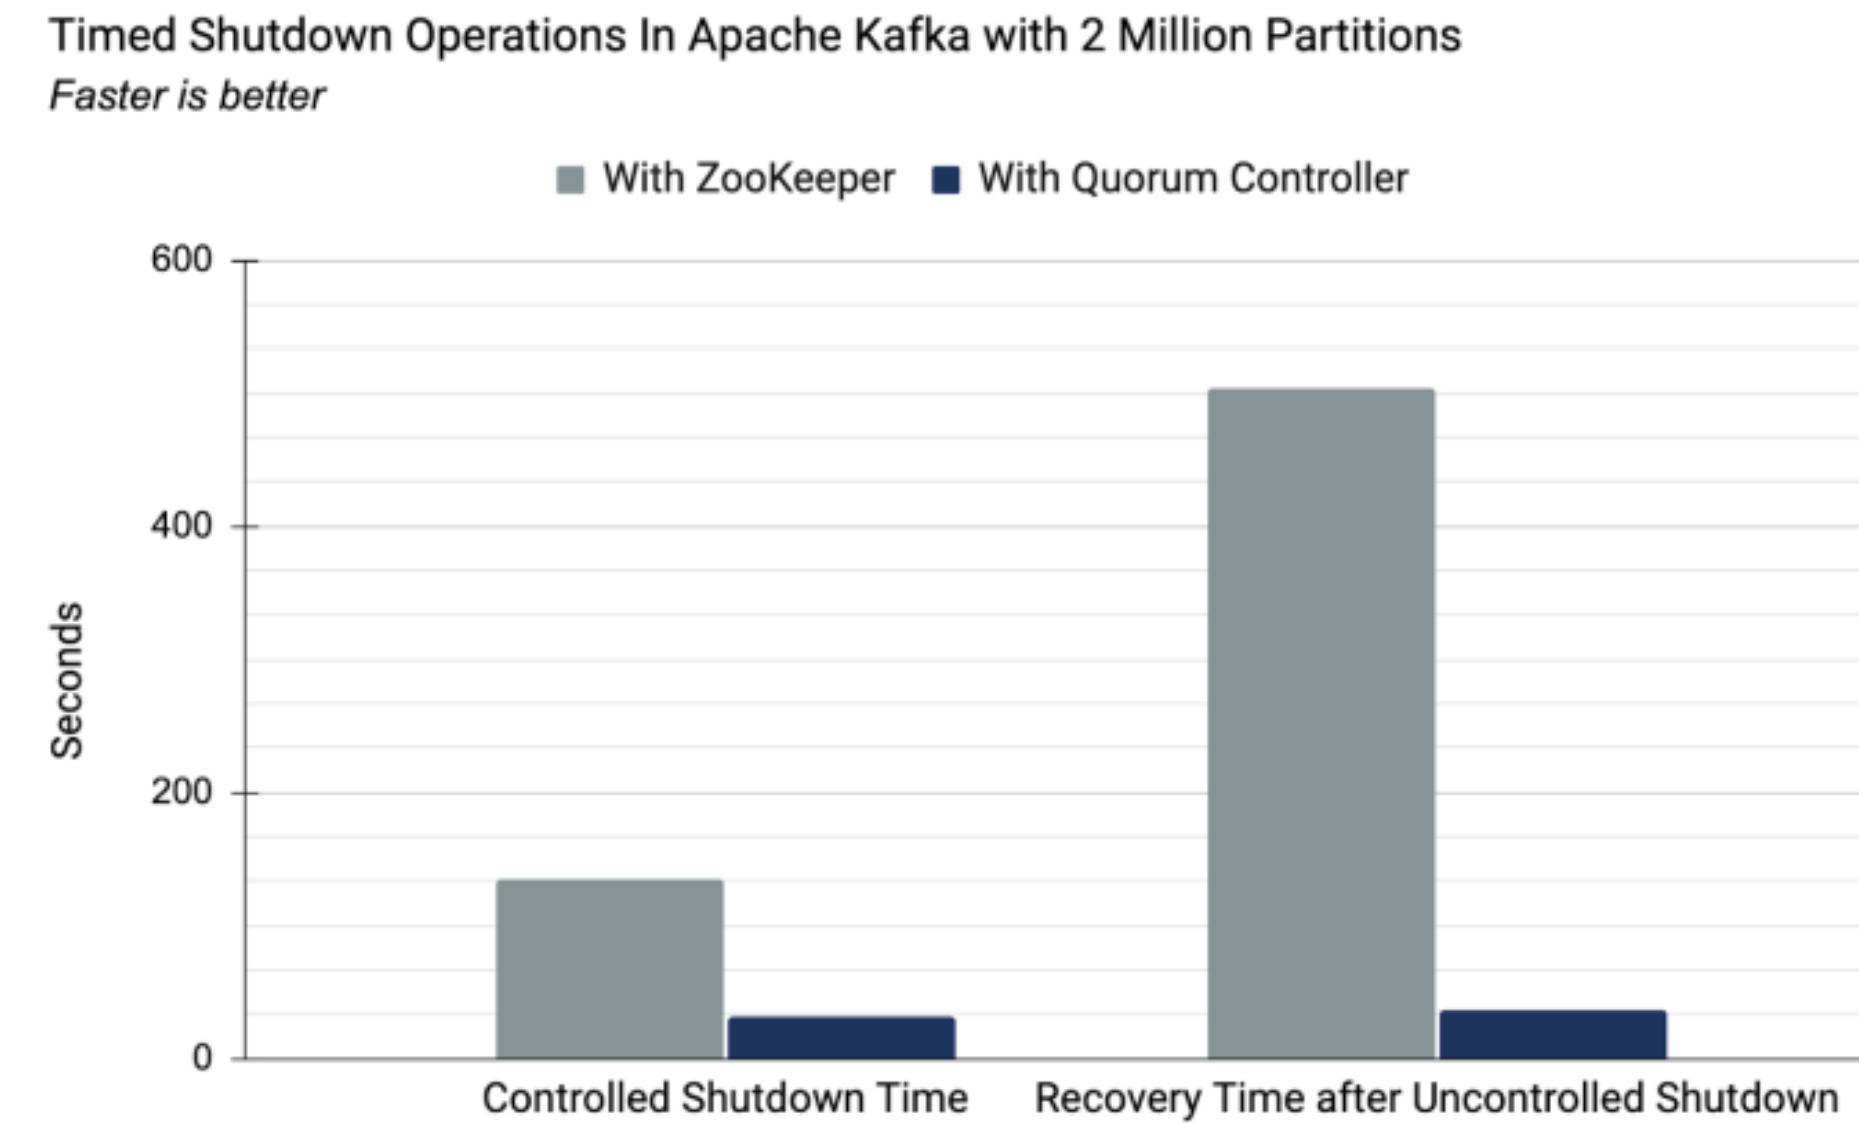
KRaft



L denotes cluster metadata leader

Kafka Raft (KRaft) mode - How it works

- In a cluster, one broker is the leader
- Metadata changes are captured by events
- A topic holds the metadata changes
- Other controllers update their metadata from the leader by consuming those events in the topic
 - This is called metadata failover and it is near instant
 - This event driven model is more efficient, simpler and reduces time to shutdown and recovery



Requirements



Hardware Requirements

- Do not co-locate other applications due to memory page cache pollution and will degrade performance
- Performant drive is required
- Storage capacity will need to be calculated by the expected messages per day and retention
- Slower networks can degrade the rate in which messages are produced

Cloud Requirements

- Analyze by data retention
- Analyze performance need by the producers
- If low latency is required, SSD should be considered
- Ephemeral Storage may be required (Elastic Block Storage)

Kafka Guarantees



Kafka Guarantees

- Messages sent by a producer to a particular topic partition will be appended in the order they are sent.
- A consumer instance sees records in the order they are stored in the log.
- For a topic with replication factor N, we will tolerate up to N-1 server failures without losing any records committed to the log.



Kafka CLI

Creating a Topic

Creating the topic `orders` with replication factor 2 and 4 partitions

```
$ kafka-topics.sh --create --bootstrap-server broker:9092 \
--replication-factor 2 --partitions 4 --topic orders
```

Listing Topics

List the topics using one of the zookeeper nodes

```
$ kafka-topics.sh --bootstrap-server broker:9092 --list
```

Sending a Message

Send a message to list of brokers for a particular topic

```
$ kafka-console-producer.sh \
  --broker-list <kafka-broker>:9092 --topic <topic>
> Hello
> I
> am
> sending
> six
> messages
```

Receiving a Message

Receiving the messages from the topic that was posted by the CLI producer

```
$ kafka-console-consumer.sh \
--bootstrap-server <kafka-broker>:9092 \
--topic <topic> \
--from-beginning

> Hello
> I
> am
> sending
> six
> messages
```

Showing Distributed Partitions

Showing how partitions are distributed

```
$ kafka-topics.sh --describe \  
  --topic <topic-name> \  
  --bootstrap-server broker:9092
```

Kafka Programming Producers

Establishing Properties

- Construct a `java.util.Properties` object
- Provide two or more locations where the bootstrap servers are located
- Provide a Serializer for the key
- Provide a Serializer for the value

```
Properties properties = new Properties();
properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
    StringSerializer.class);
properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
    IntegerSerializer.class);
```

Create a Producer Object

- Construct a `org.apache.kafka.clients.producer.KafkaProducer` object
- A Kafka Producer is thread-safe
- Inject the Properties into the Producer object

```
KafkaProducer<String, Integer> producer = new KafkaProducer<>(properties);
```

Creating a Record/Message

- Create a `org.apache.kafka.clients.producer.ProducerRecord`
- Accepts many parameters but the main ones are:
 - Topic
 - Key (if applicable)
 - Value

```
ProducerRecord<String, Integer> producerRecord =  
    new ProducerRecord<>("my_orders", state, amount);
```

Sending a message

- Send the record by calling send on the Producer
- Returns a Future object to process the results on another Thread.
- The Future will contain RecordMetadata, an object that has information about your send.

```
Future<RecordMetadata> send = producer.send(producerRecord);
```

RecordMetadata

Contains information about your send including the messages

```
if (metadata.hasOffset()) {  
    System.out.format("offset: %d\n",  
        metadata.offset());  
}  
System.out.format("partition: %d\n",  
    metadata.partition());  
System.out.format("timestamp: %d\n",  
    metadata.timestamp());  
System.out.format("topic: %s\n", metadata.topic());  
System.out.format("toString: %s\n",  
    metadata.toString());
```

Sending with a Callback

- Alternately, you can send with a Callback (lambda)
 - Callback is an interface, can be used as a lambda
 - If RecordMetadata is null there was an error, if Exception is null the send was successful

```
producer.send(producerRecord, new Callback() {  
    @Override  
    public void onCompletion(RecordMetadata metadata,  
                            Exception e) {  
        . . .  
    }  
}
```

Using a Closure to capture Key and Value

- RecordMetadata does not have information on key and value
- Using a closure you can obtain that in the block of your lambda

```
producer.send(producerRecord, (metadata, e) -> {
    if (metadata != null) {
        System.out.println(producerRecord.key());
        System.out.println(producerRecord.value());
    }
})
```

Be a good citizen, close your resources

- When you need to terminate, flush messages from the bufferpool
- Close the Producer

```
producer.flush();
producer.close();
```

Closing Resources in a Shutdown Hook

- `Runtime.getRuntime().addShutdownHook(...)` will listen for SIGTERM (CTRL+C)
- This would make an excellent place to flush and close, and close any loops that you may have created

```
Runtime.getRuntime().addShutdownHook(new Thread(() -> {
    done.set(true);
    producer.flush();
    producer.close();
}));
```

Acknowledgements & Retries

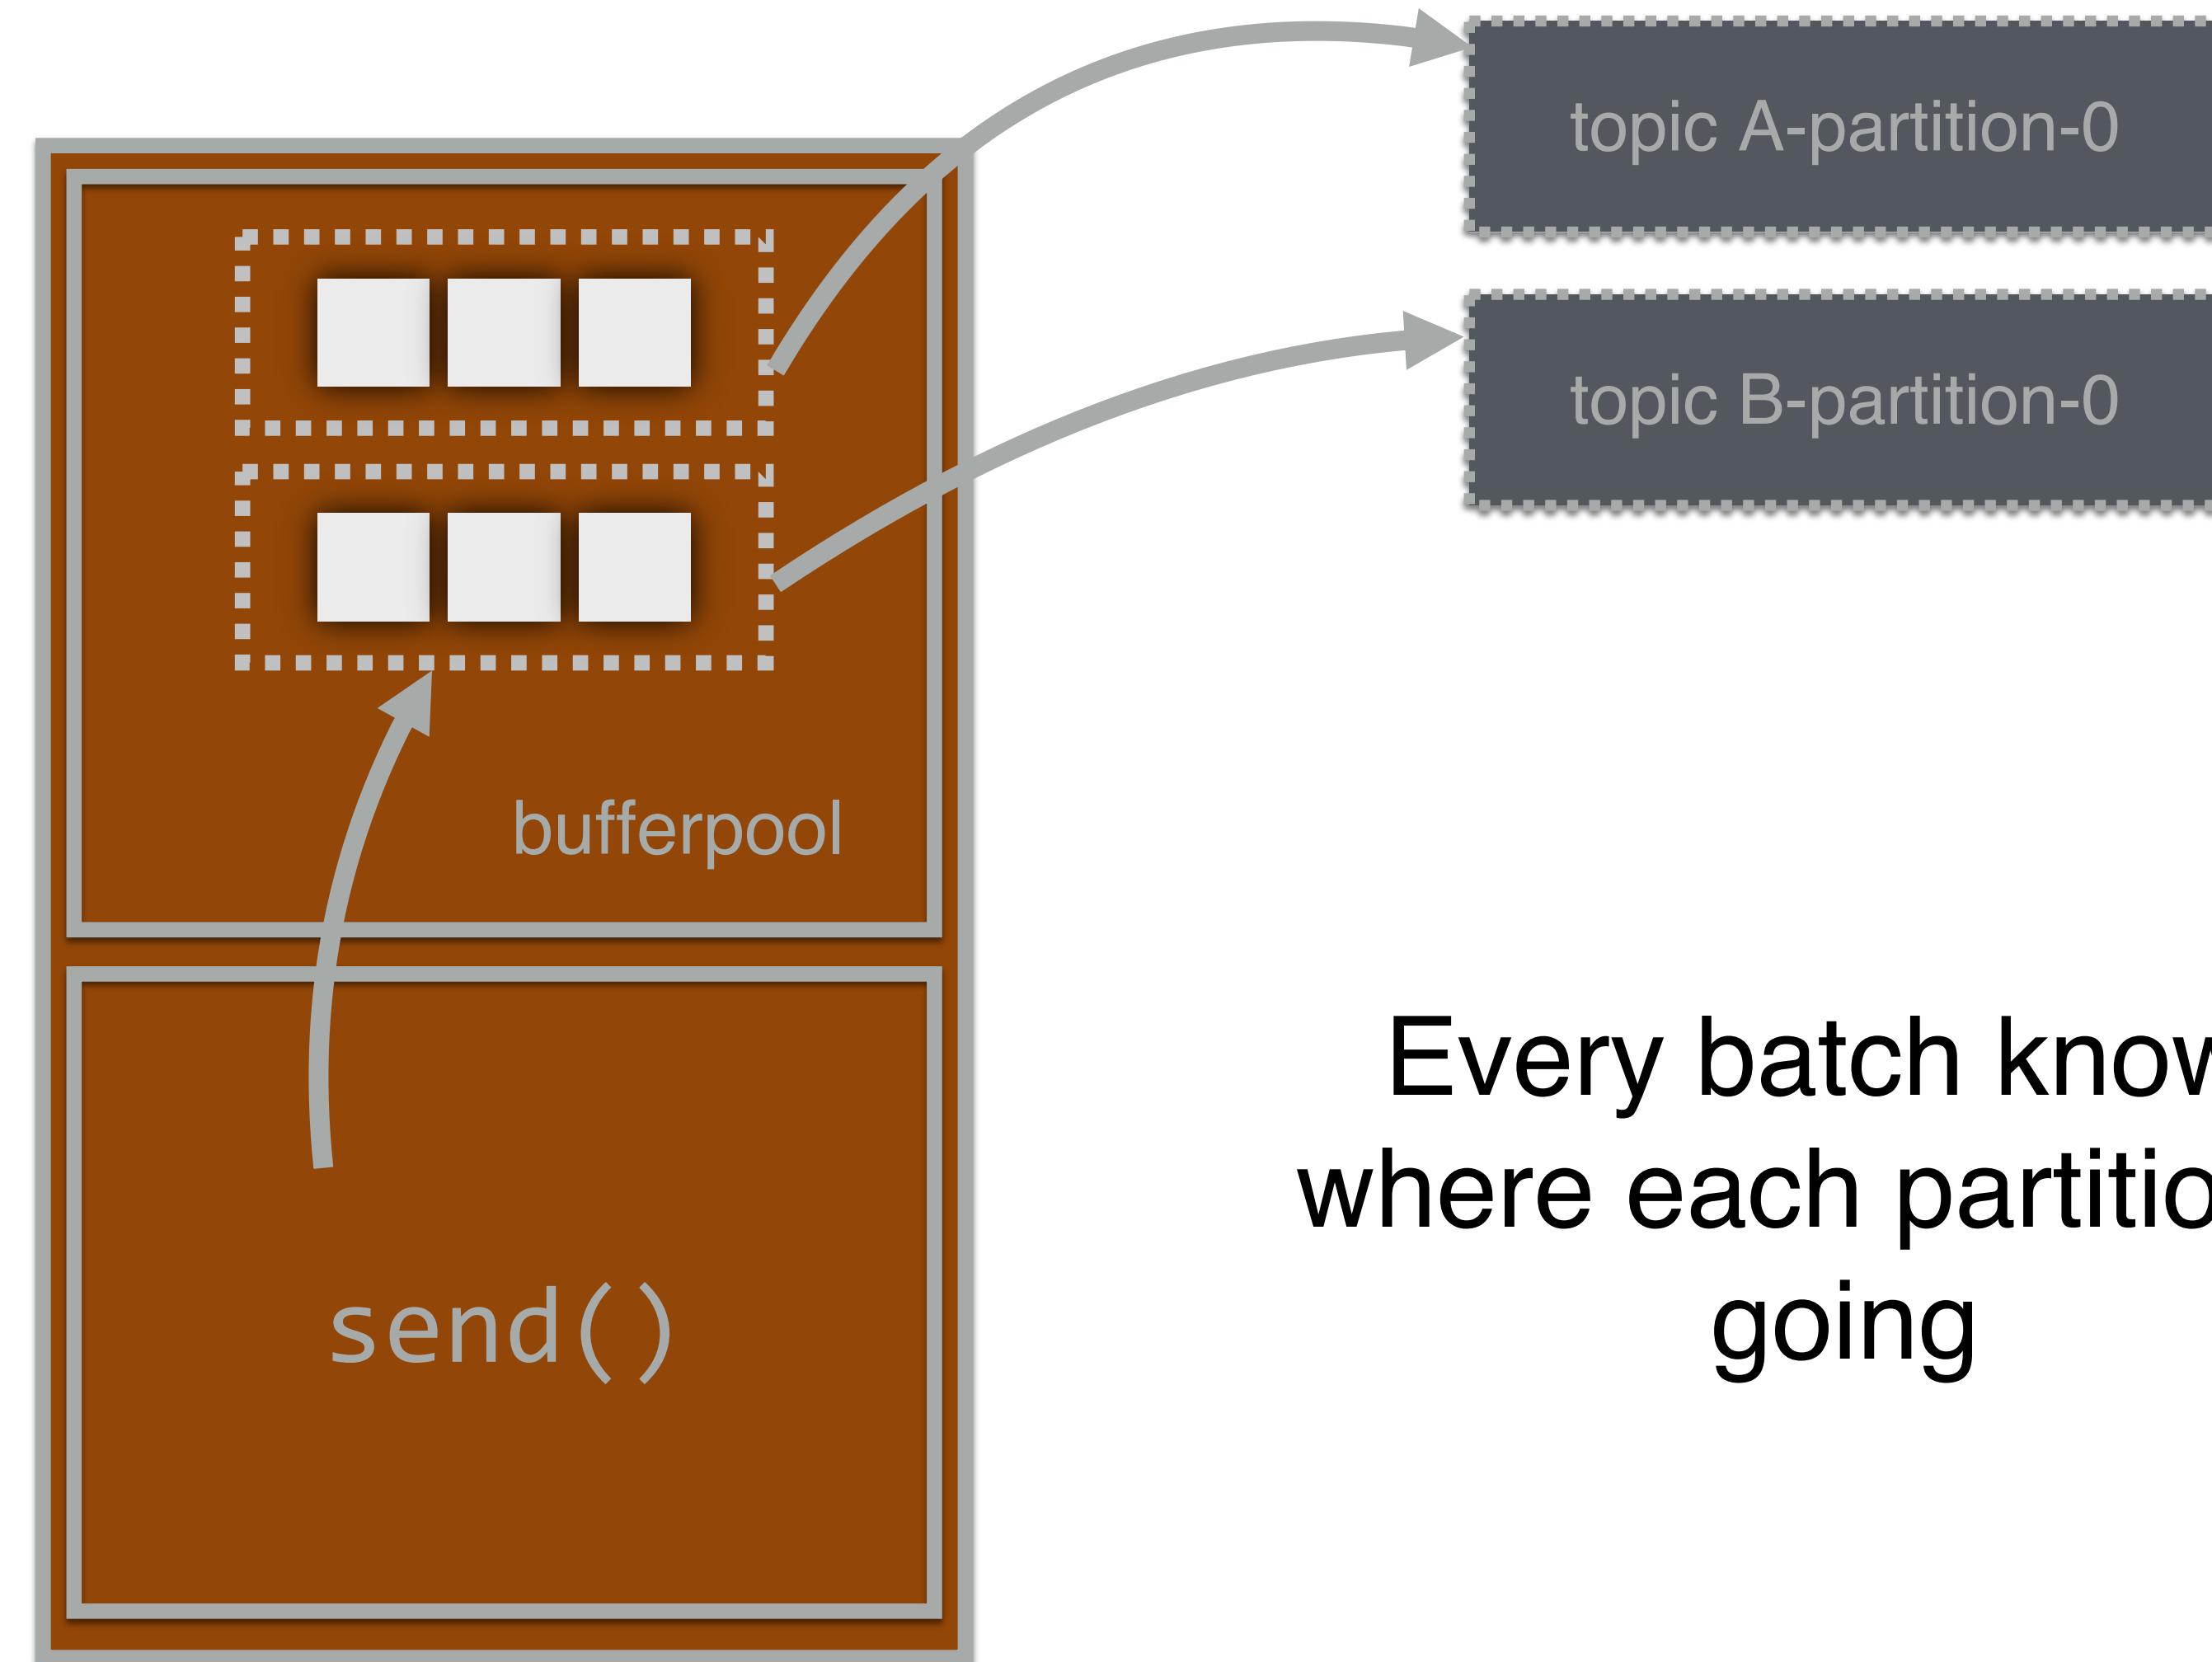
Acks

acks controls how many partition replicas must receive the record before the write is considered a success.

```
Properties properties = new Properties();
properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
    StringSerializer.class);
properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
    IntegerSerializer.class);
properties.put(ProducerConfig.ACKS_CONFIG, "all");
properties.put(ProducerConfig.RETRIES_CONFIG, 20);
properties.put(ProducerConfig.RETRY_BACKOFF_MS_CONFIG, 3000);
```

Acks

acks	description
0	No acknowledgment, assume all is well
1	At least one replica will Producer will receive a success response, error if unsuccessful, up to client to hand
all	All replicas must acknowledge. Higher latency, safest



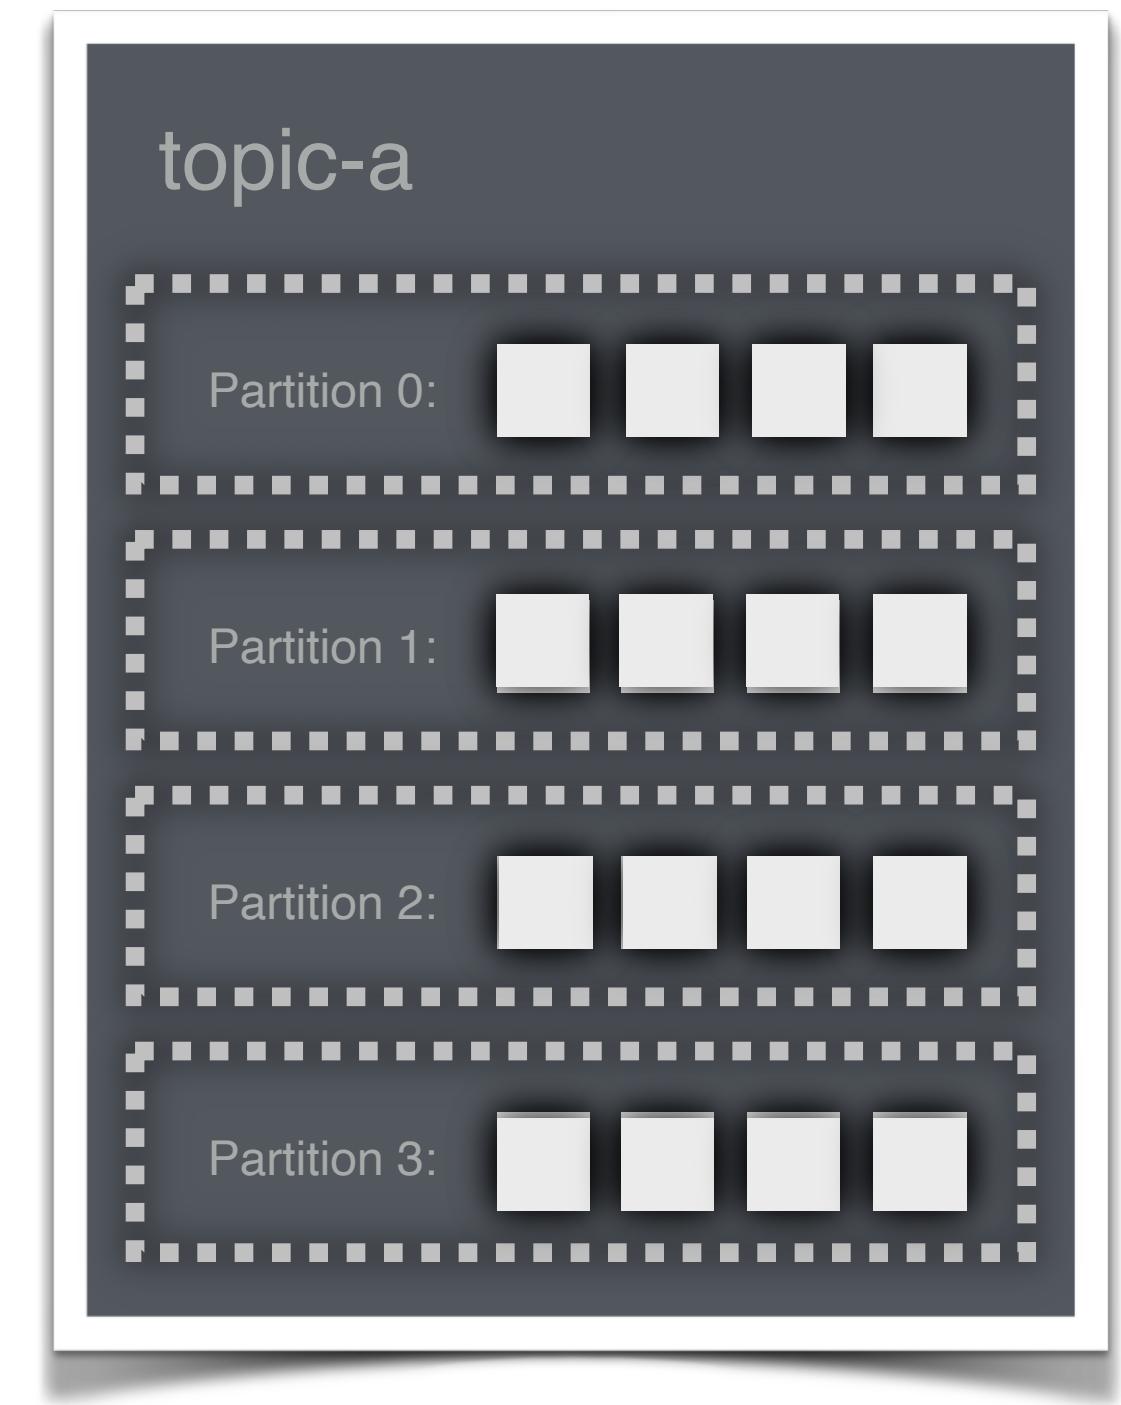
Every batch knows
where each partition is
going

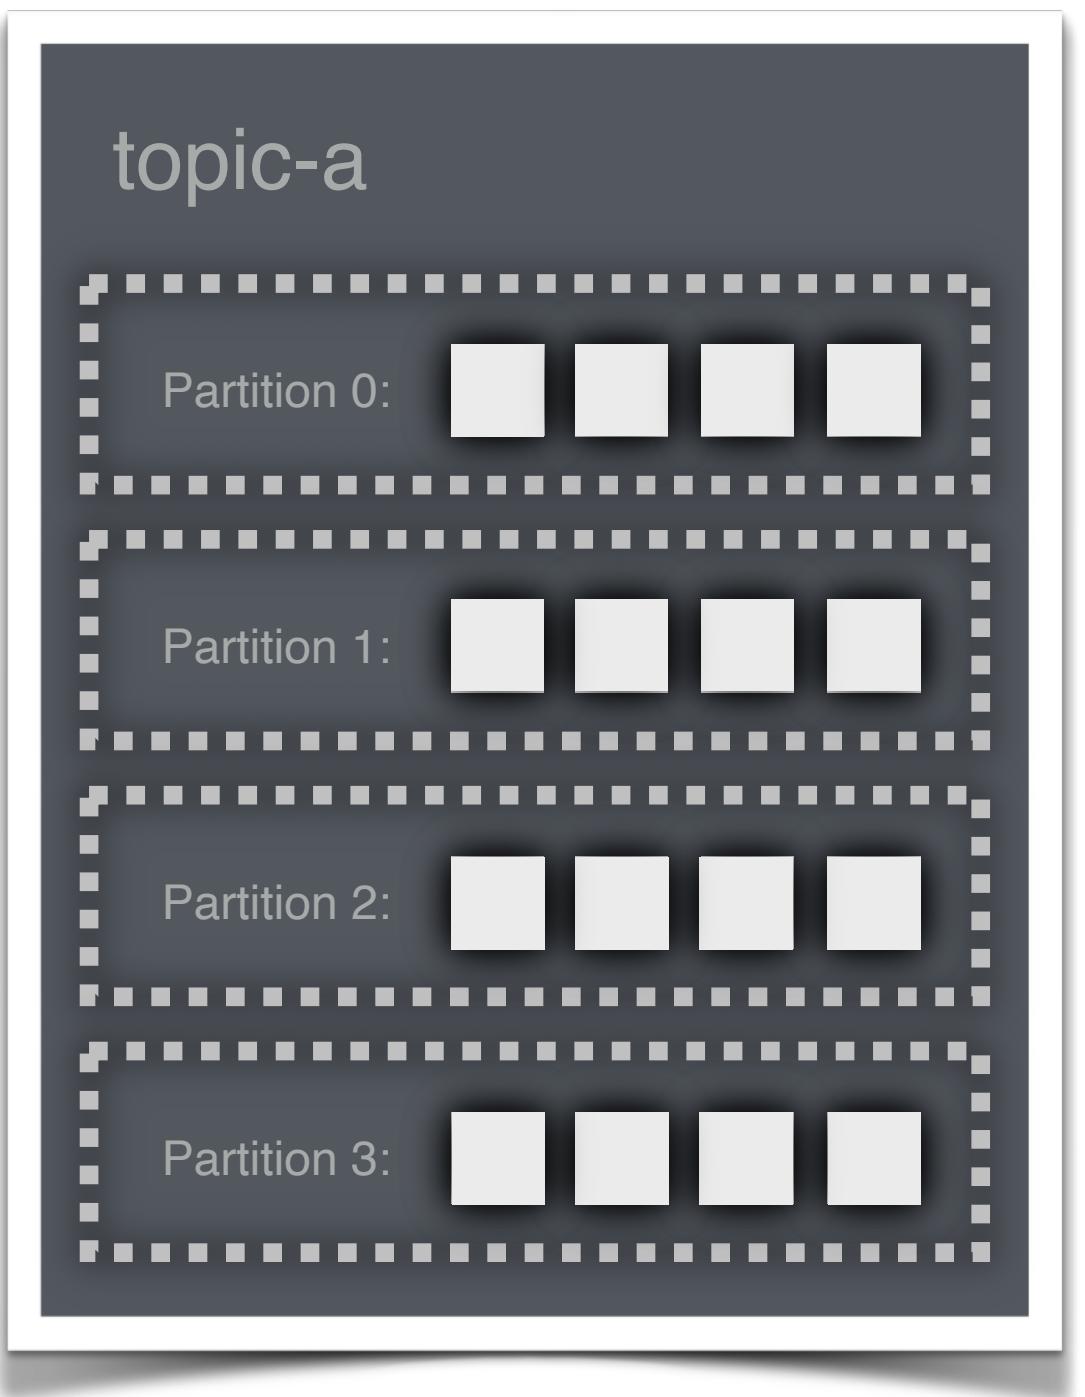
$\text{murmur2(bytes)} \% \text{ number partitions}$

Kafka Programming Consumers & Groups

Kafka Consumer Groups

- Consumers are typically done as a group
- A single consumer will end up inefficient with large amounts of data
- A consumer may never catch up
- **Every consumer should be on it's own machine, instance, pod**





consumer-1

consumer-2

consumer-3

consumer-4

Kafka's Goal

Kafka scales to large amount of different consumers without affecting performance

Kafka Consumer Threading

- There are no multiple consumers that belong to the same group in one thread.
One consumer per one thread.
- There are not multiple threads running one consumer, **One consumer per one thread.**

consumer_offsets

- Topic on Kafka brokers that contain information about consumer and their offsets, stored in Kafka's data directory

Establishing Properties

- Construct a `java.util.Properties` object
- Provide two or more locations where the bootstrap servers are located
- Provide a "Team Name", officially called a `group.id`
- Provide a DeSerializer for the key
- Provide a DeSerializer for the value

```
Properties properties = new Properties();
properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");
properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
    "org.apache.kafka.common.serialization.StringDeserializer");
properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
    "org.apache.kafka.common.serialization.IntegerDeserializer");
```

Create a Consumer Object

- Construct a `org.apache.kafka.clients.consumer.KafkaConsumer` object
- A Kafka Consumer is **not thread-safe**
- Inject the Properties into the Consumer object

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Processing Messages

- Use the consumer that you have constructed, and call poll with pulse time
- The poll is a max sleep time, if a message is ready it will download a batch, ConsumerRecords
- Iterate through each record in the batch with a for-loop, process the message

```
while (!done.get()) {  
    ConsumerRecords<String, String> records =  
        consumer.poll(Duration.of(500, ChronoUnit.MILLISECONDS));  
    for (ConsumerRecord<String, String> record : records) {  
        System.out.format("offset: %d\n", record.offset());  
        System.out.format("partition: %d\n", record.partition());  
        System.out.format("timestamp: %d\n", record.timestamp());  
        System.out.format("timeStampType: %s\n", record.timestampType());  
        System.out.format("topic: %s\n", record.topic());  
        System.out.format("key: %s\n", record.key());  
        System.out.format("value: %s\n", record.value());  
    }  
}
```

Be a good citizen, close your resources

- When you need to terminate, let the *Kafka group coordinator* know you are done
- Close the Consumer

```
consumer.close();
```

Consumer Offset Reset

Consumer Offset Reset

- When a consumer is assigned it must determine what offset for the partition to start
- Position is set according to your consumer's reset policy, either earliest or latest
- This is only applicable if there is no valid offset in the `__consumer_offset__` topic

```
Properties properties = new Properties();
...
properties.put(ConsumerConfig.AUTO_OFFSET_RESET_CONFIG, "earliest");

KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Manual Commits

Manual Commits

- Consumers will commit automatically every 5 seconds
- If your consumer fails within those 5 seconds another consumer may get those messages
- Might be too long before commit, in such case you may want to handle commits yourself
- Turn off your auto-commit by setting `ENABLE_AUTO_COMMIT_CONFIG` to false

```
Properties properties = new Properties();
...
properties.put(ConsumerConfig.ENABLE_AUTO_COMMIT_CONFIG, false);
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);
```

Synchronous Commit

- After turning off auto-commit
- Committing and blocking synchronously yourself
- Synchronous commits will block

Synchronous Commit

- After turning off auto-commit
- Committing and blocking synchronously yourself
- Synchronous commits will block

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(properties);  
...  
consumer.commitSync(); //Block
```

Asynchronous Commit

- After turning off auto-commit
- Committing and letting it commit asynchronously
- Advantage: Speed, Higher throughput
- Disadvantage: You will get error reports asynchronously

```
consumer.commitAsync((offsets, exception) -> {  
    //offsets  
    //exception  
});
```

Using both kinds of commits

- You can commit asynchronously within the loop
- Then commit synchronously outside the loop
- Reason, if you are out of the loop, likely you want to shut down

```
Runtime.getRuntime().addShutdownHook(new Thread(() -> done.set(true)));

while (!done.get()) {
    ConsumerRecords<String, String> records = consumer.poll(Duration.of(500, ChronoUnit.MILLISECONDS));
    for (ConsumerRecord<String, String> record : records) {
        ...
    }
    consumer.commitAsync((offsets, exception) -> {
        //offsets
        //exception
    });
}
consumer.commitSync(); //Block
consumer.close();
```

Consumer Rebalancing

Partition Rebalance

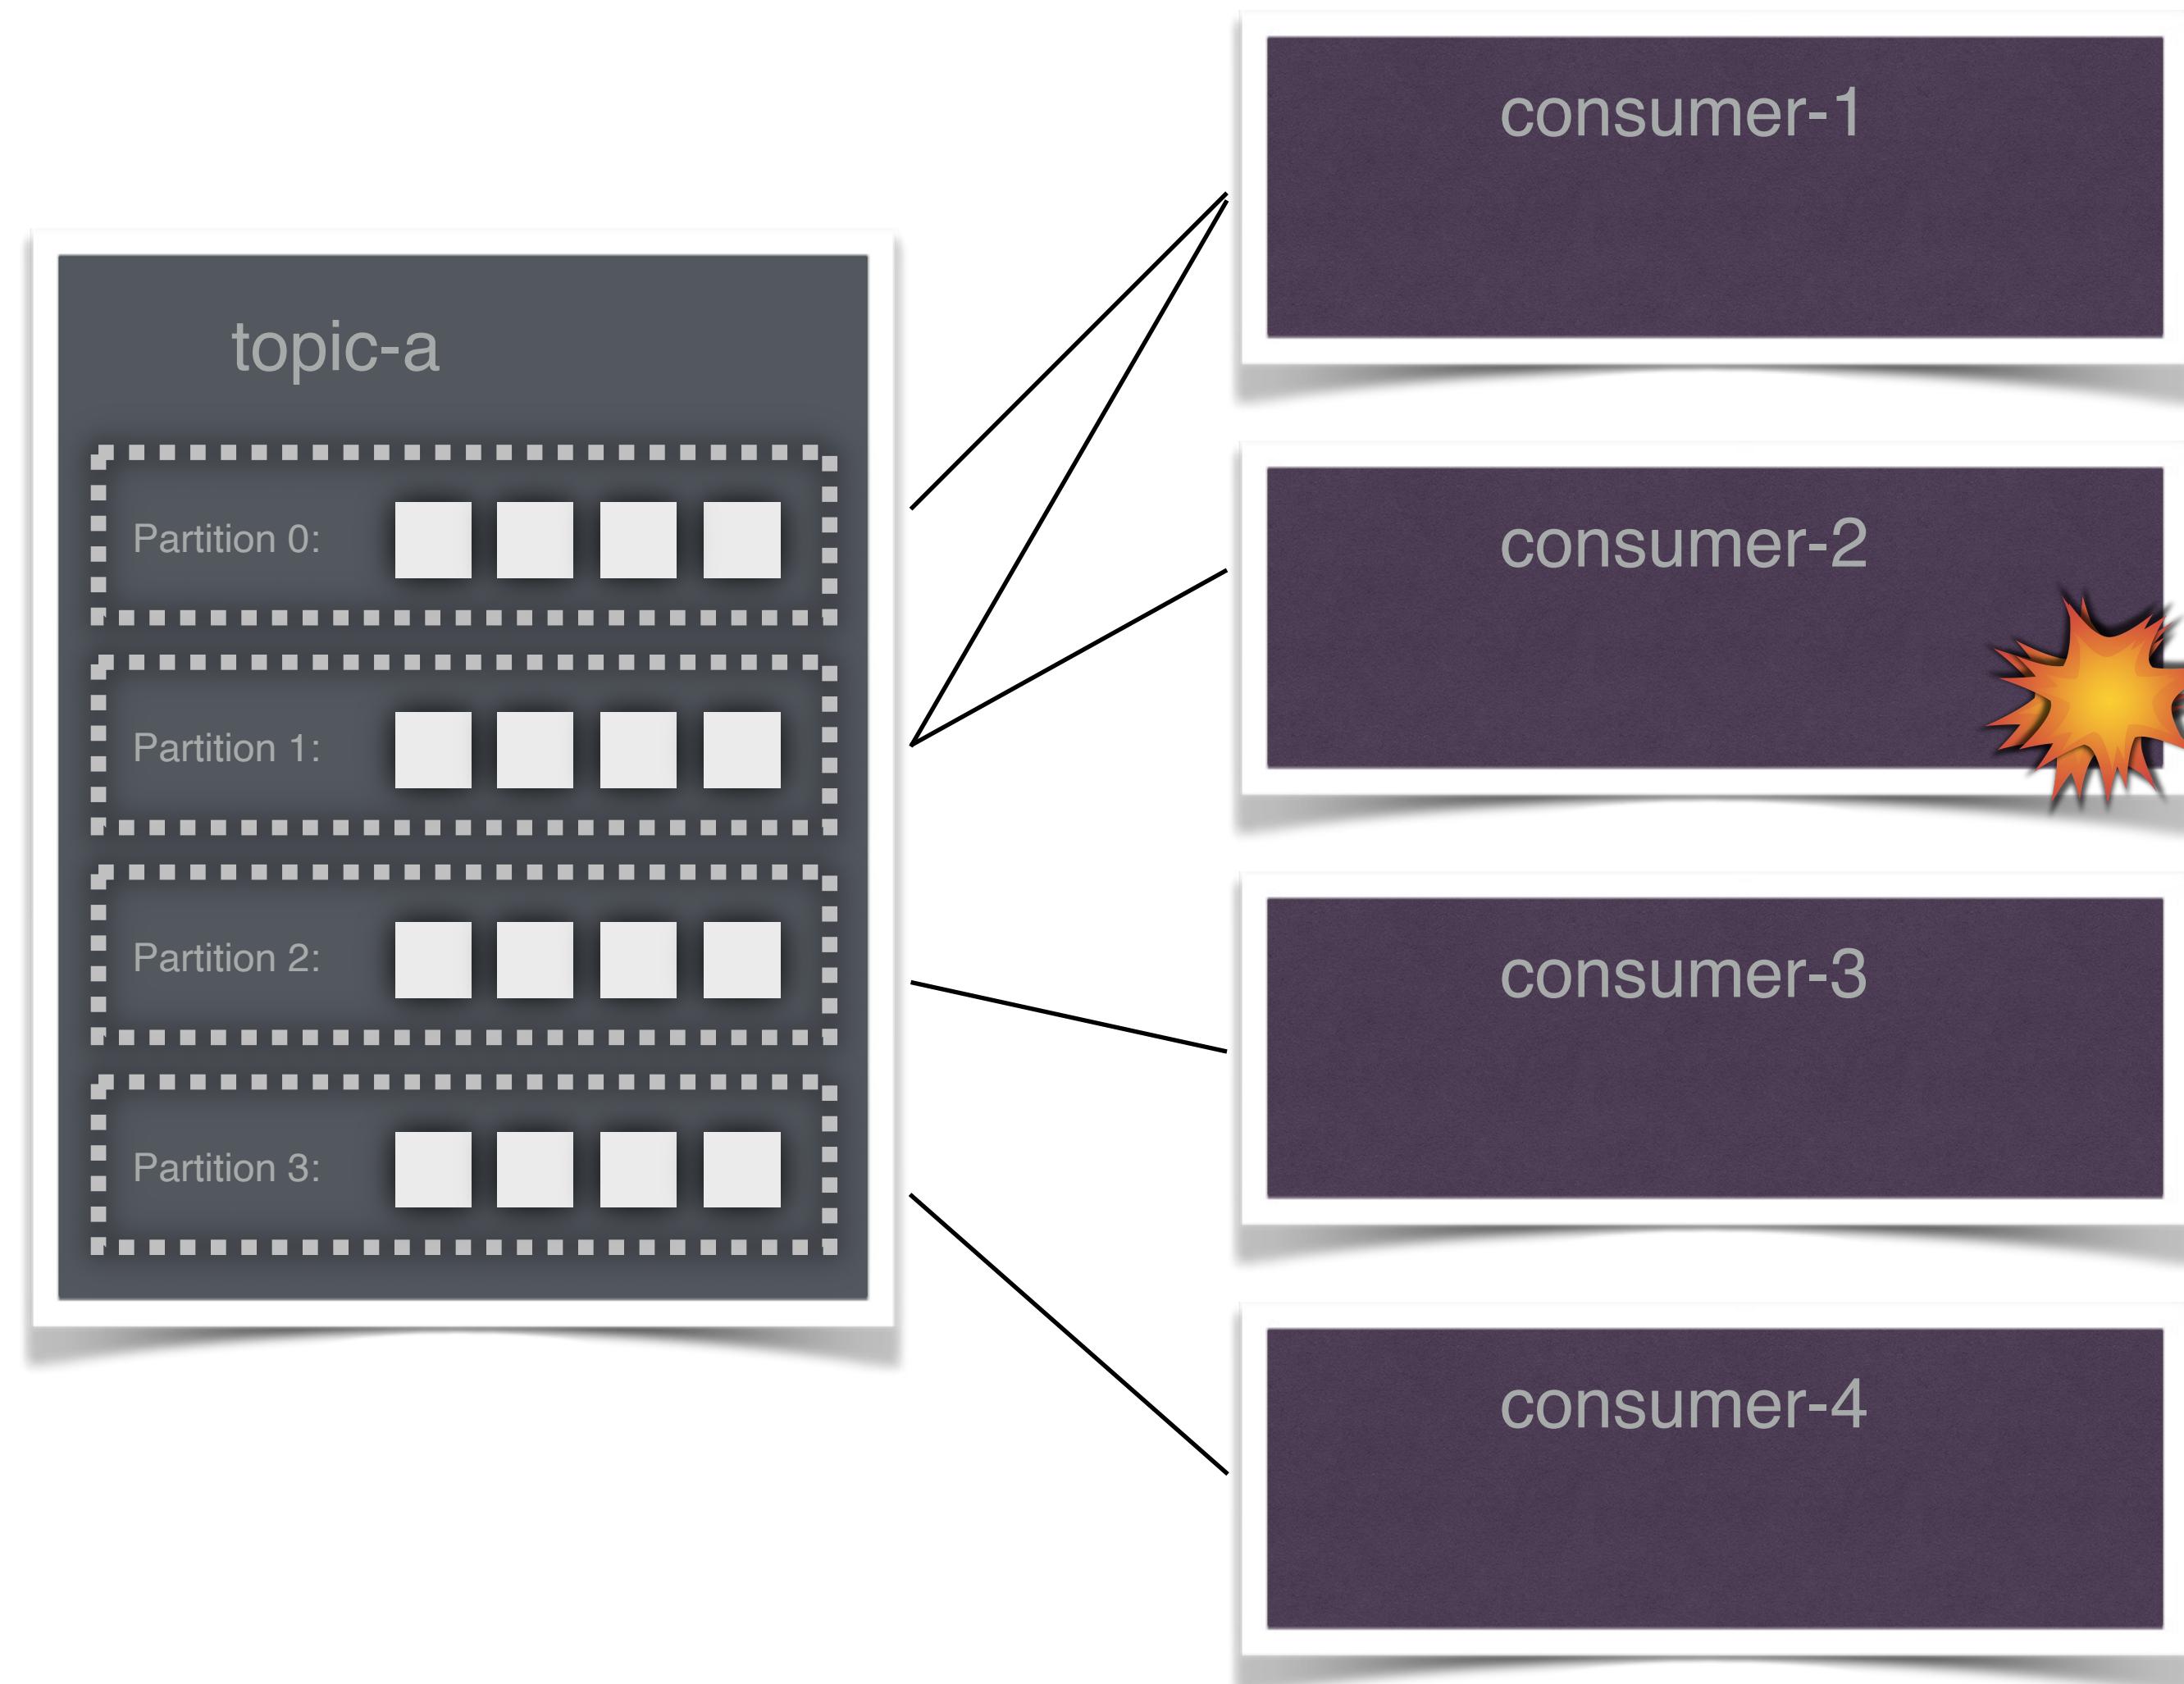
- When one partition is moved from one consumer to another, this is known as a *rebalance*.
- A way to mitigate when either consumers go down, or when consumers are added.
- Although unavoidable, this will cause an unfortunate pause, and it ***will lose state***.

Groups and Heartbeats

- At regular intervals, consumers will send heartbeats to the broker, to let it know it is alive and still reading data
- Heartbeats are sent to a Kafka broker called the *group coordinator*
- They are sent when the consumer polls and consumes a record

Leaving the Group

- When a consumer leaves the group, it lets the *group coordinator* know it is done
- The Kafka broker then triggers a rebalance of the group.



Consumer Rebalance Listener

- You can implement a `ConsumerRebalanceListener` and be alerted when a rebalance has occurred
- The rebalance listener requires that you implement `onPartitionsRevoked`, and `onPartitionsAssigned`

```
consumer.subscribe(Collections.singletonList("my_orders"),
    new ConsumerRebalanceListener() {
        @Override
        public void onPartitionsRevoked(Collection<TopicPartition> collection) {
            System.out.println("Partition revoked:" +
                collectionTopicPartitionToString(collection));
            consumer.commitAsync();
        }

        @Override
        public void onPartitionsAssigned(Collection<TopicPartition> collection) {
            System.out.println("Partition assigned:" +
                collectionTopicPartitionToString(collection));
        }
    });
}
```

Retention Settings

Retention

- Durable Storage over a Period of Time
- Can either be configured in time or size
- Once reached messages are expunged

Retention Formula

Who administers Kafka can update the default retention time or size

```
log.retention.hours = 120  
log.retention.bytes = 3221225472
```

Default Hours are set to (168 hours) or one week

You may also use log.retention.minutes or log.retention.ms, whichever is smallest wins.

You can also configure retention per topic, like the following

```
bin/kafka-configs --zookeeper localhost:2181 --entity-type topics --entity-name my-topic --alter --add-config retention.ms=128000
```

Retention Evaluation

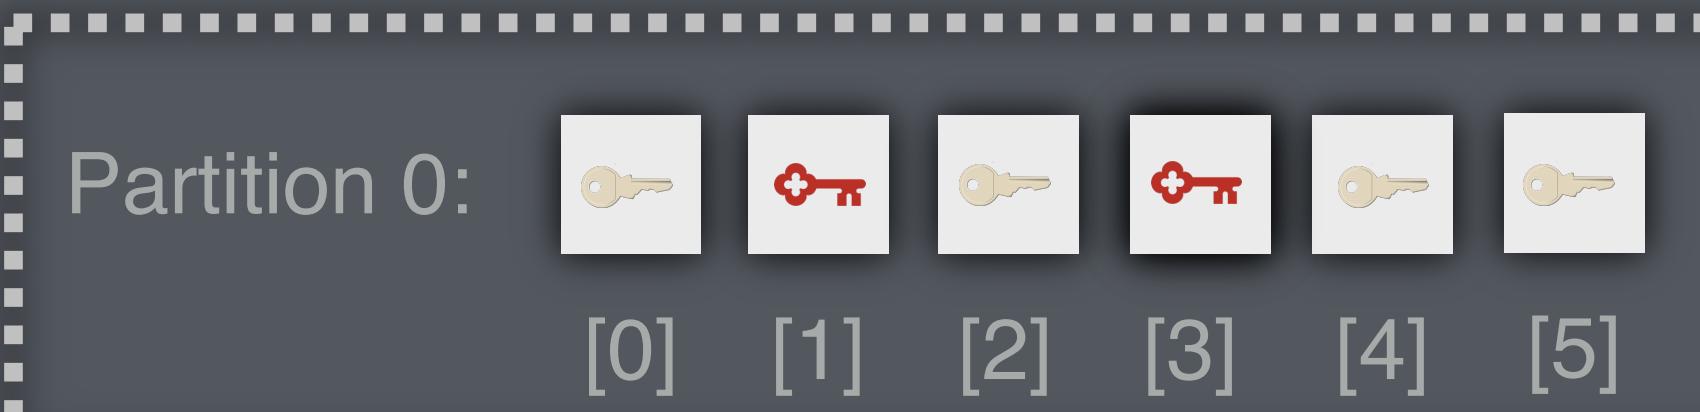
- Durable Storage over a Period of Time
- Can either be configured in time or size
- Once reached messages are expunged

Compaction

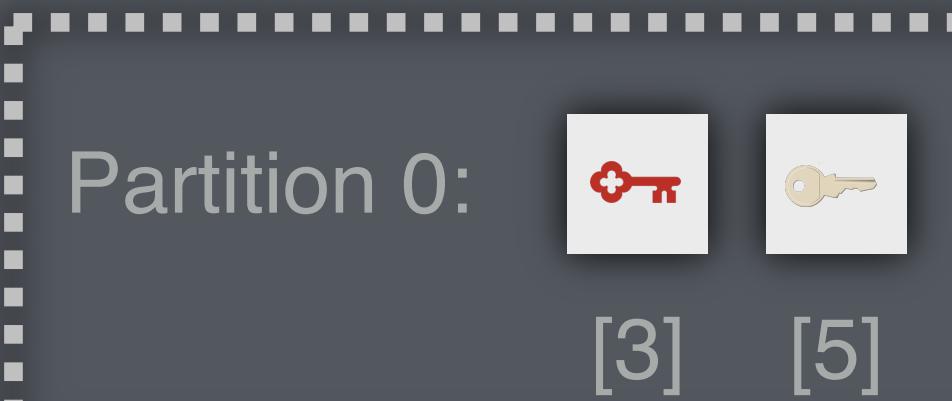
Compaction

- A form of retention where messages of the same key where only the latest message will be retained.
- Compaction is performed by a *cleaner thread*

kafka broker: 0



kafka broker: 0



kafka broker: 0



clean

dirty

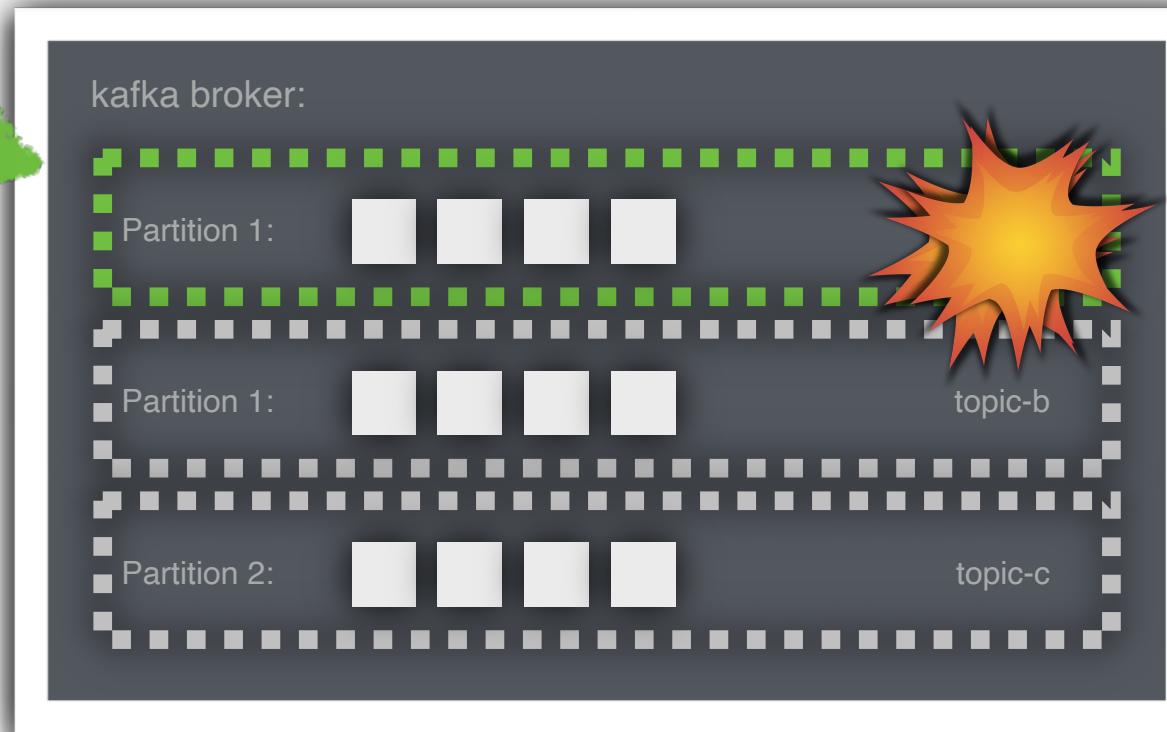
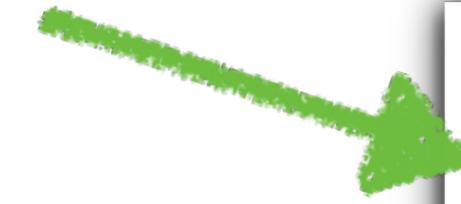
Kafka will start compacting when 50% of the topic contains dirty records

Resiliency

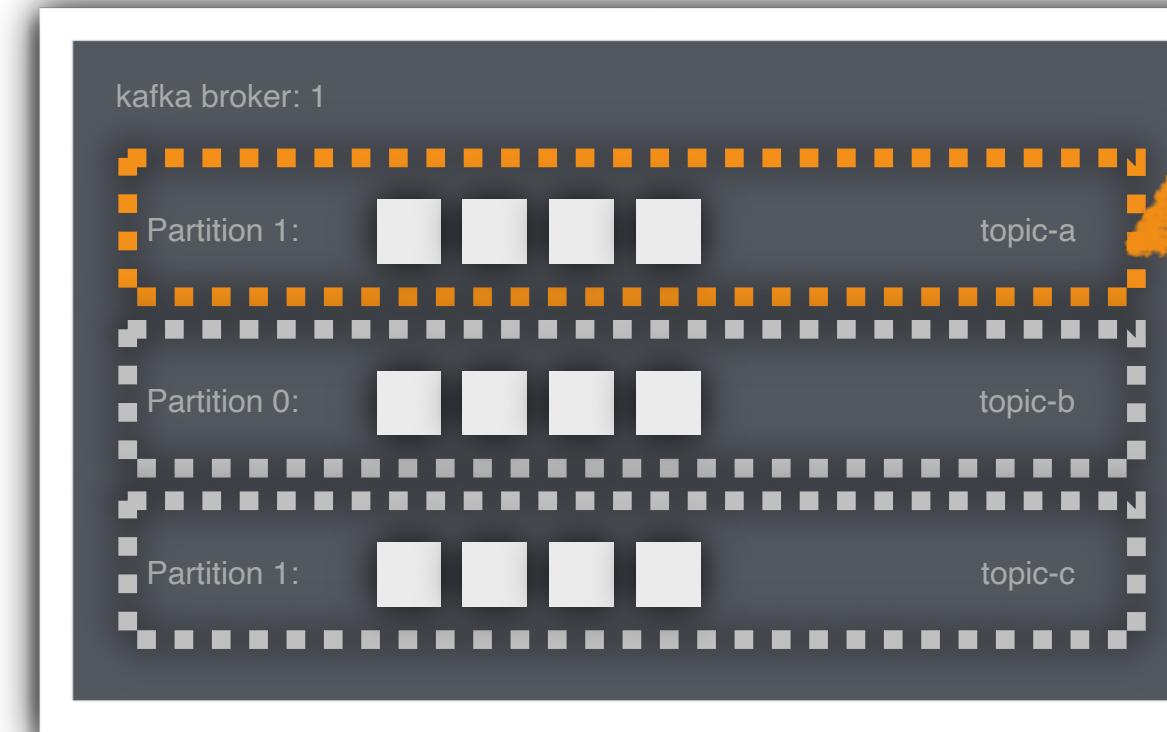
ISR (In sync replicas)

- Given a leader partition, an in sync replica is one that has been kept up to date within the last 10 seconds
- This is configurable
- During a crash, the closest ISR will retain control for failover

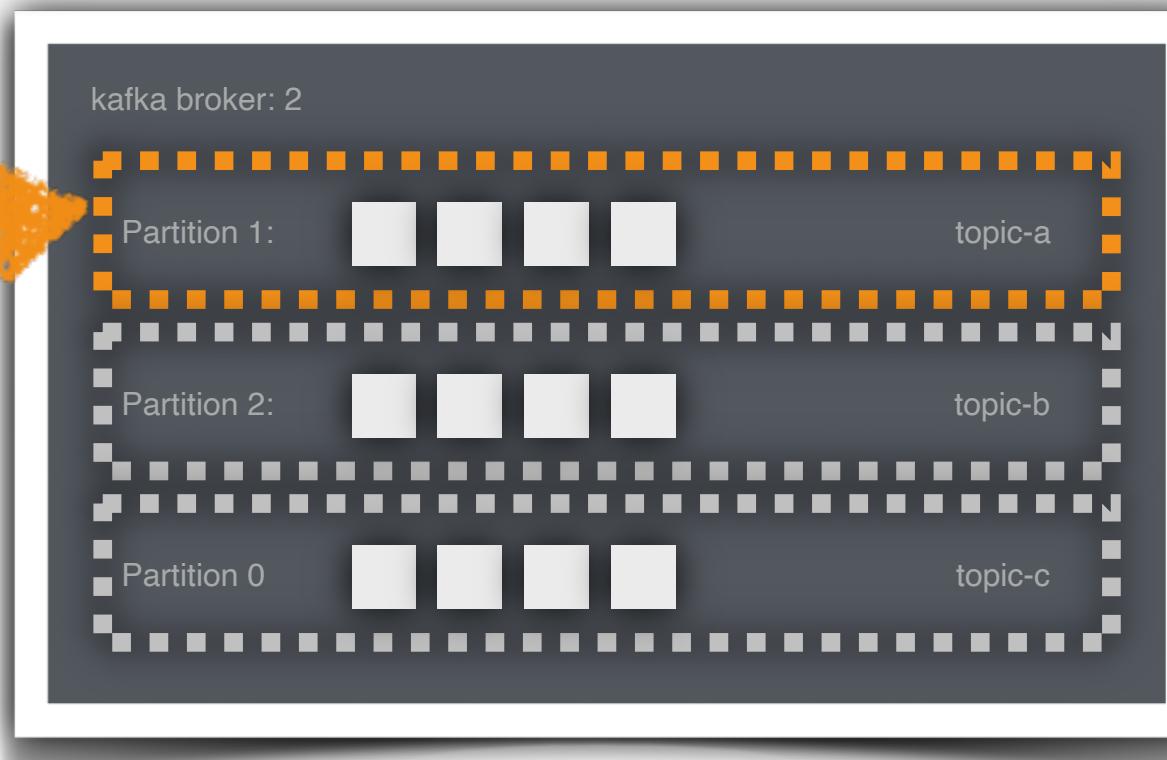
Leader



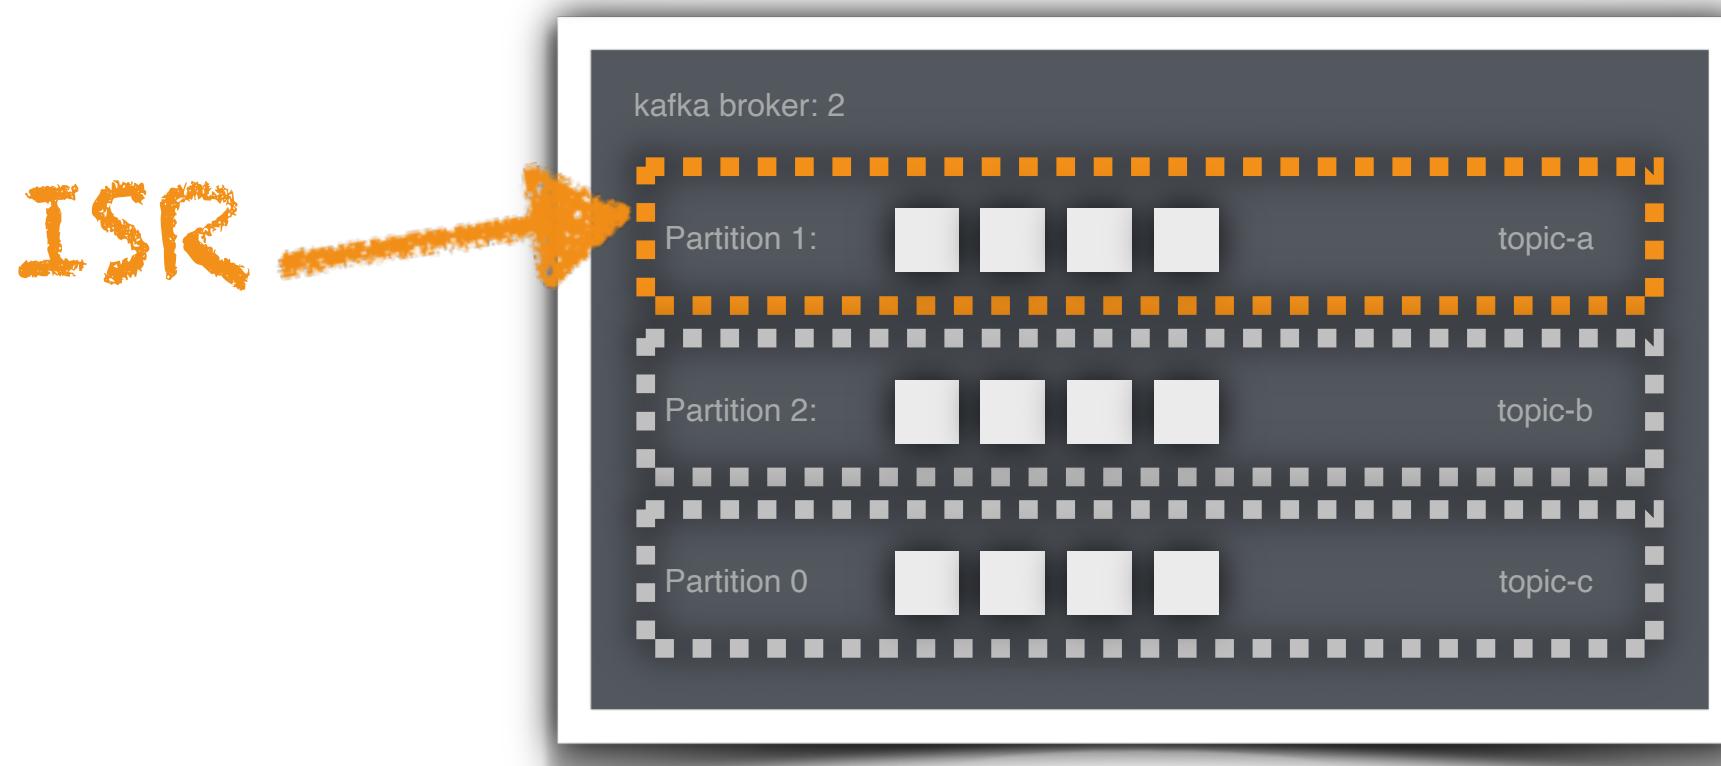
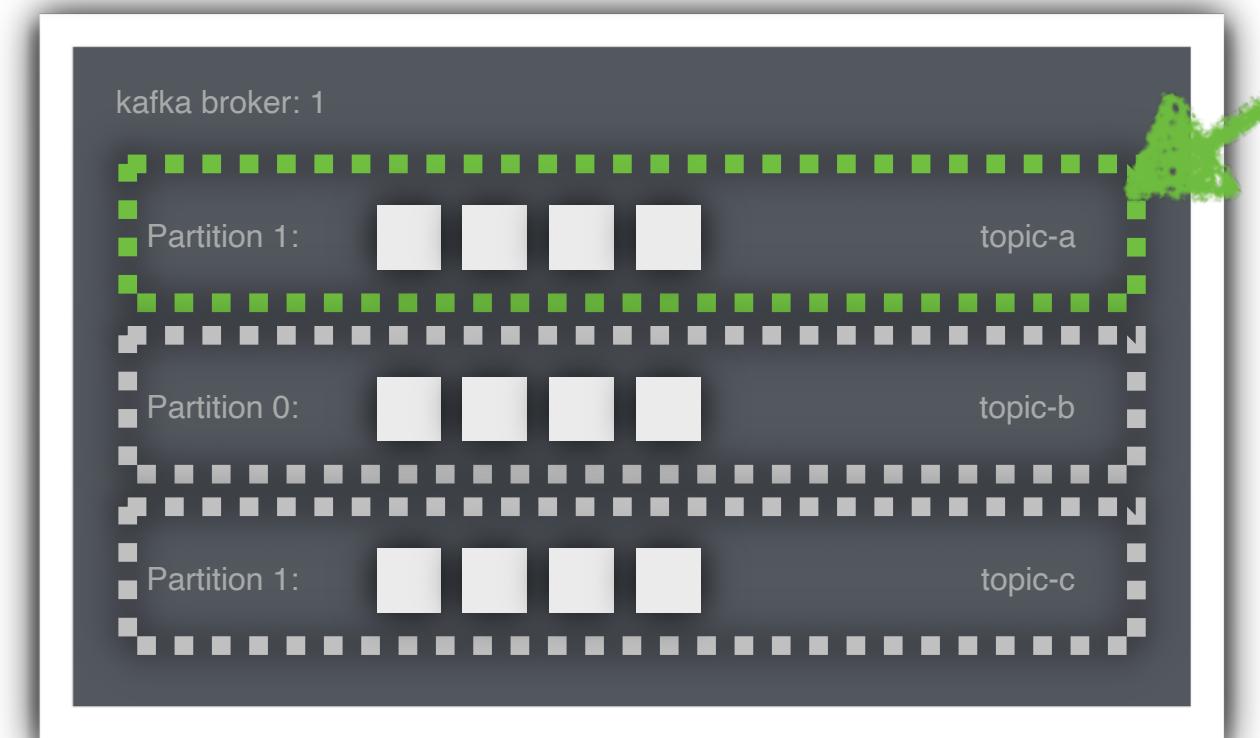
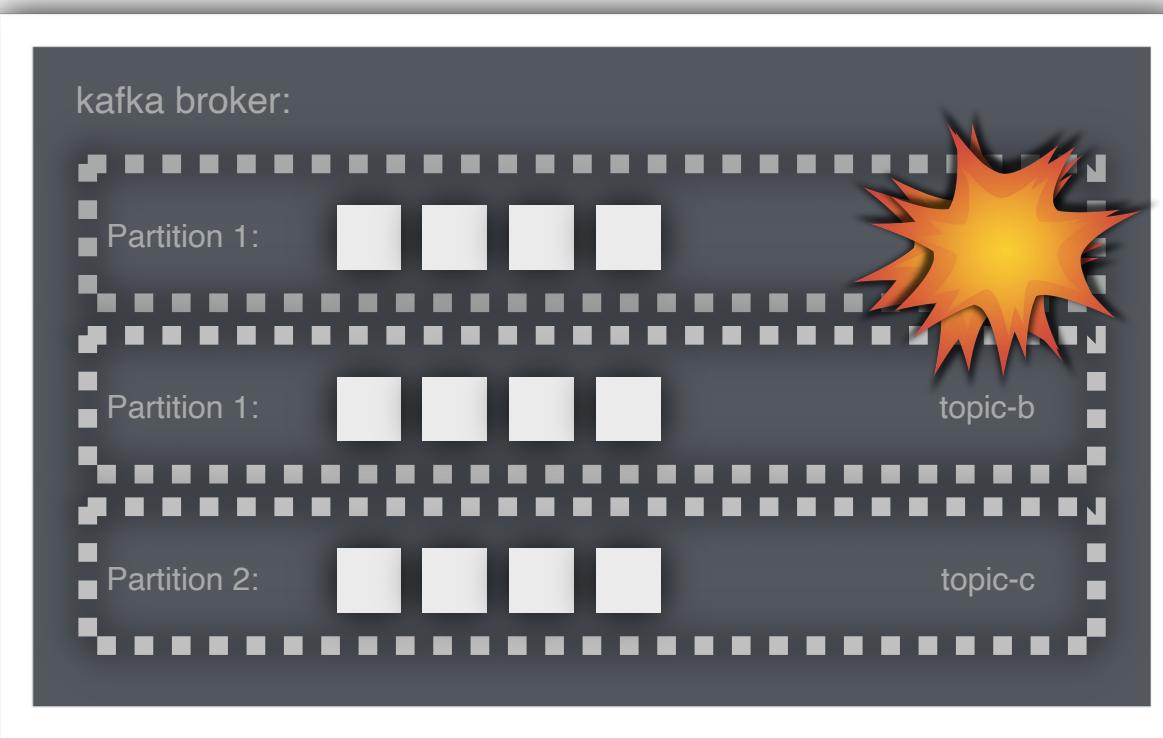
ISR



ISR



Leader



Exactly Once Semantics

**Mathias Verraes**

@mathiasverraes

Follow

There are only two hard problems in distributed systems:
2. Exactly-once delivery
1. Guaranteed order of messages
2. Exactly-once delivery

RETWEETS

6,775

LIKES

4,727

10:40 AM - 14 Aug 2015



69



6.8K

**4.7K**

Exactly Once Semantics

- If a stream processing application consumes message A and produces message B such that $B = F(A)$
- *Exactly once* processing means that A is considered consumed if and only if B is successfully produced, and vice versa.

How Exactly Once Processing is Lost (1)

- The producer.send() could result in duplicate writes of message *B* due to internal retries.
- This is addressed by the idempotent producer
- *Exactly once* processing means that *A* is considered consumed if and only if *B* is successfully produced, and vice versa.

How Exactly Once Processing is Lost (2)

- We may reprocess the input message A , resulting in duplicate B messages being written to the output
- Reprocessing may happen if the stream processing application crashes after writing B but before marking A as consumed.
- Thus when it resumes, it will consume A again and write B again, causing a duplicate.

How Exactly Once Processing is Lost (3)

- Finally, in distributed environments, applications will crash or—worse!—temporarily lose connectivity to the rest of the system.
- Typically, new instances are automatically started (like Kubernetes) to replace the ones which were deemed lost.
- Through this process, we may have multiple instances processing the same input topics and writing to the same output topics, causing duplicate outputs and violating the exactly once processing semantics.
- We call this the problem of “zombie instances.”



Requirements for EOS

- Transaction API
 - All of the messages included in the transaction will be successfully written or none of them will be.
- Zombie Fencing
 - We solve the problem of zombie instances by requiring that each transactional producer be assigned a unique identifier called the `transactional.id`.

transaction.id and epochs

- Producers register the Producer's transaction.id
- Identifies producer with a unique identifier across process restarts
- Registers [transactional.id](#) with broker
- Kafka broker checks for open transactions with the given transactional.id and completes them
- Increments an epoch associated with the [transactional.id](#)
- Any producers with same transactional.id and an older epoch are considered zombies and are fenced off, ie. future transactional writes from those producers are rejected.

Consumers and transactions

- Guarantees with consumers fall short
- Consumers will
 - not know whether these messages were written as part of a transaction
 - not know when transactions start or end
 - not guaranteed to be subscribed to all partitions which are part of a transaction
- **Kafka guarantees that a consumer will eventually deliver only non-transactional messages or committed transactional messages**

Transactions in the Producer

To "turn on" transactions in the Producer

- Add transaction id to identify the producer
- Initialize the transaction

```
Properties properties = new Properties();
properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
    StringSerializer.class);
properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
    IntegerSerializer.class);
properties.put(ProducerConfig.TRANSACTIONAL_ID_CONFIG, "my-transaction-id");

producer.initTransactions();
```

Sending Transactionally

```
producer.initTransactions();
```

```
try {  
    producer.beginTransaction();  
    producer.send(record1);  
    producer.send(record2);  
    producer.commitTransaction();  
} catch(ProducerFencedException e) {  
    producer.close();  
} catch(KafkaException e) {  
    producer.abortTransaction();  
}
```

Using transactions:

- Beginning Transaction with `beginTransaction`
- Send records as normal
- commit the Transaction
- If fenced, close it, nothing more we can do
- Other Exceptions, abort that transaction

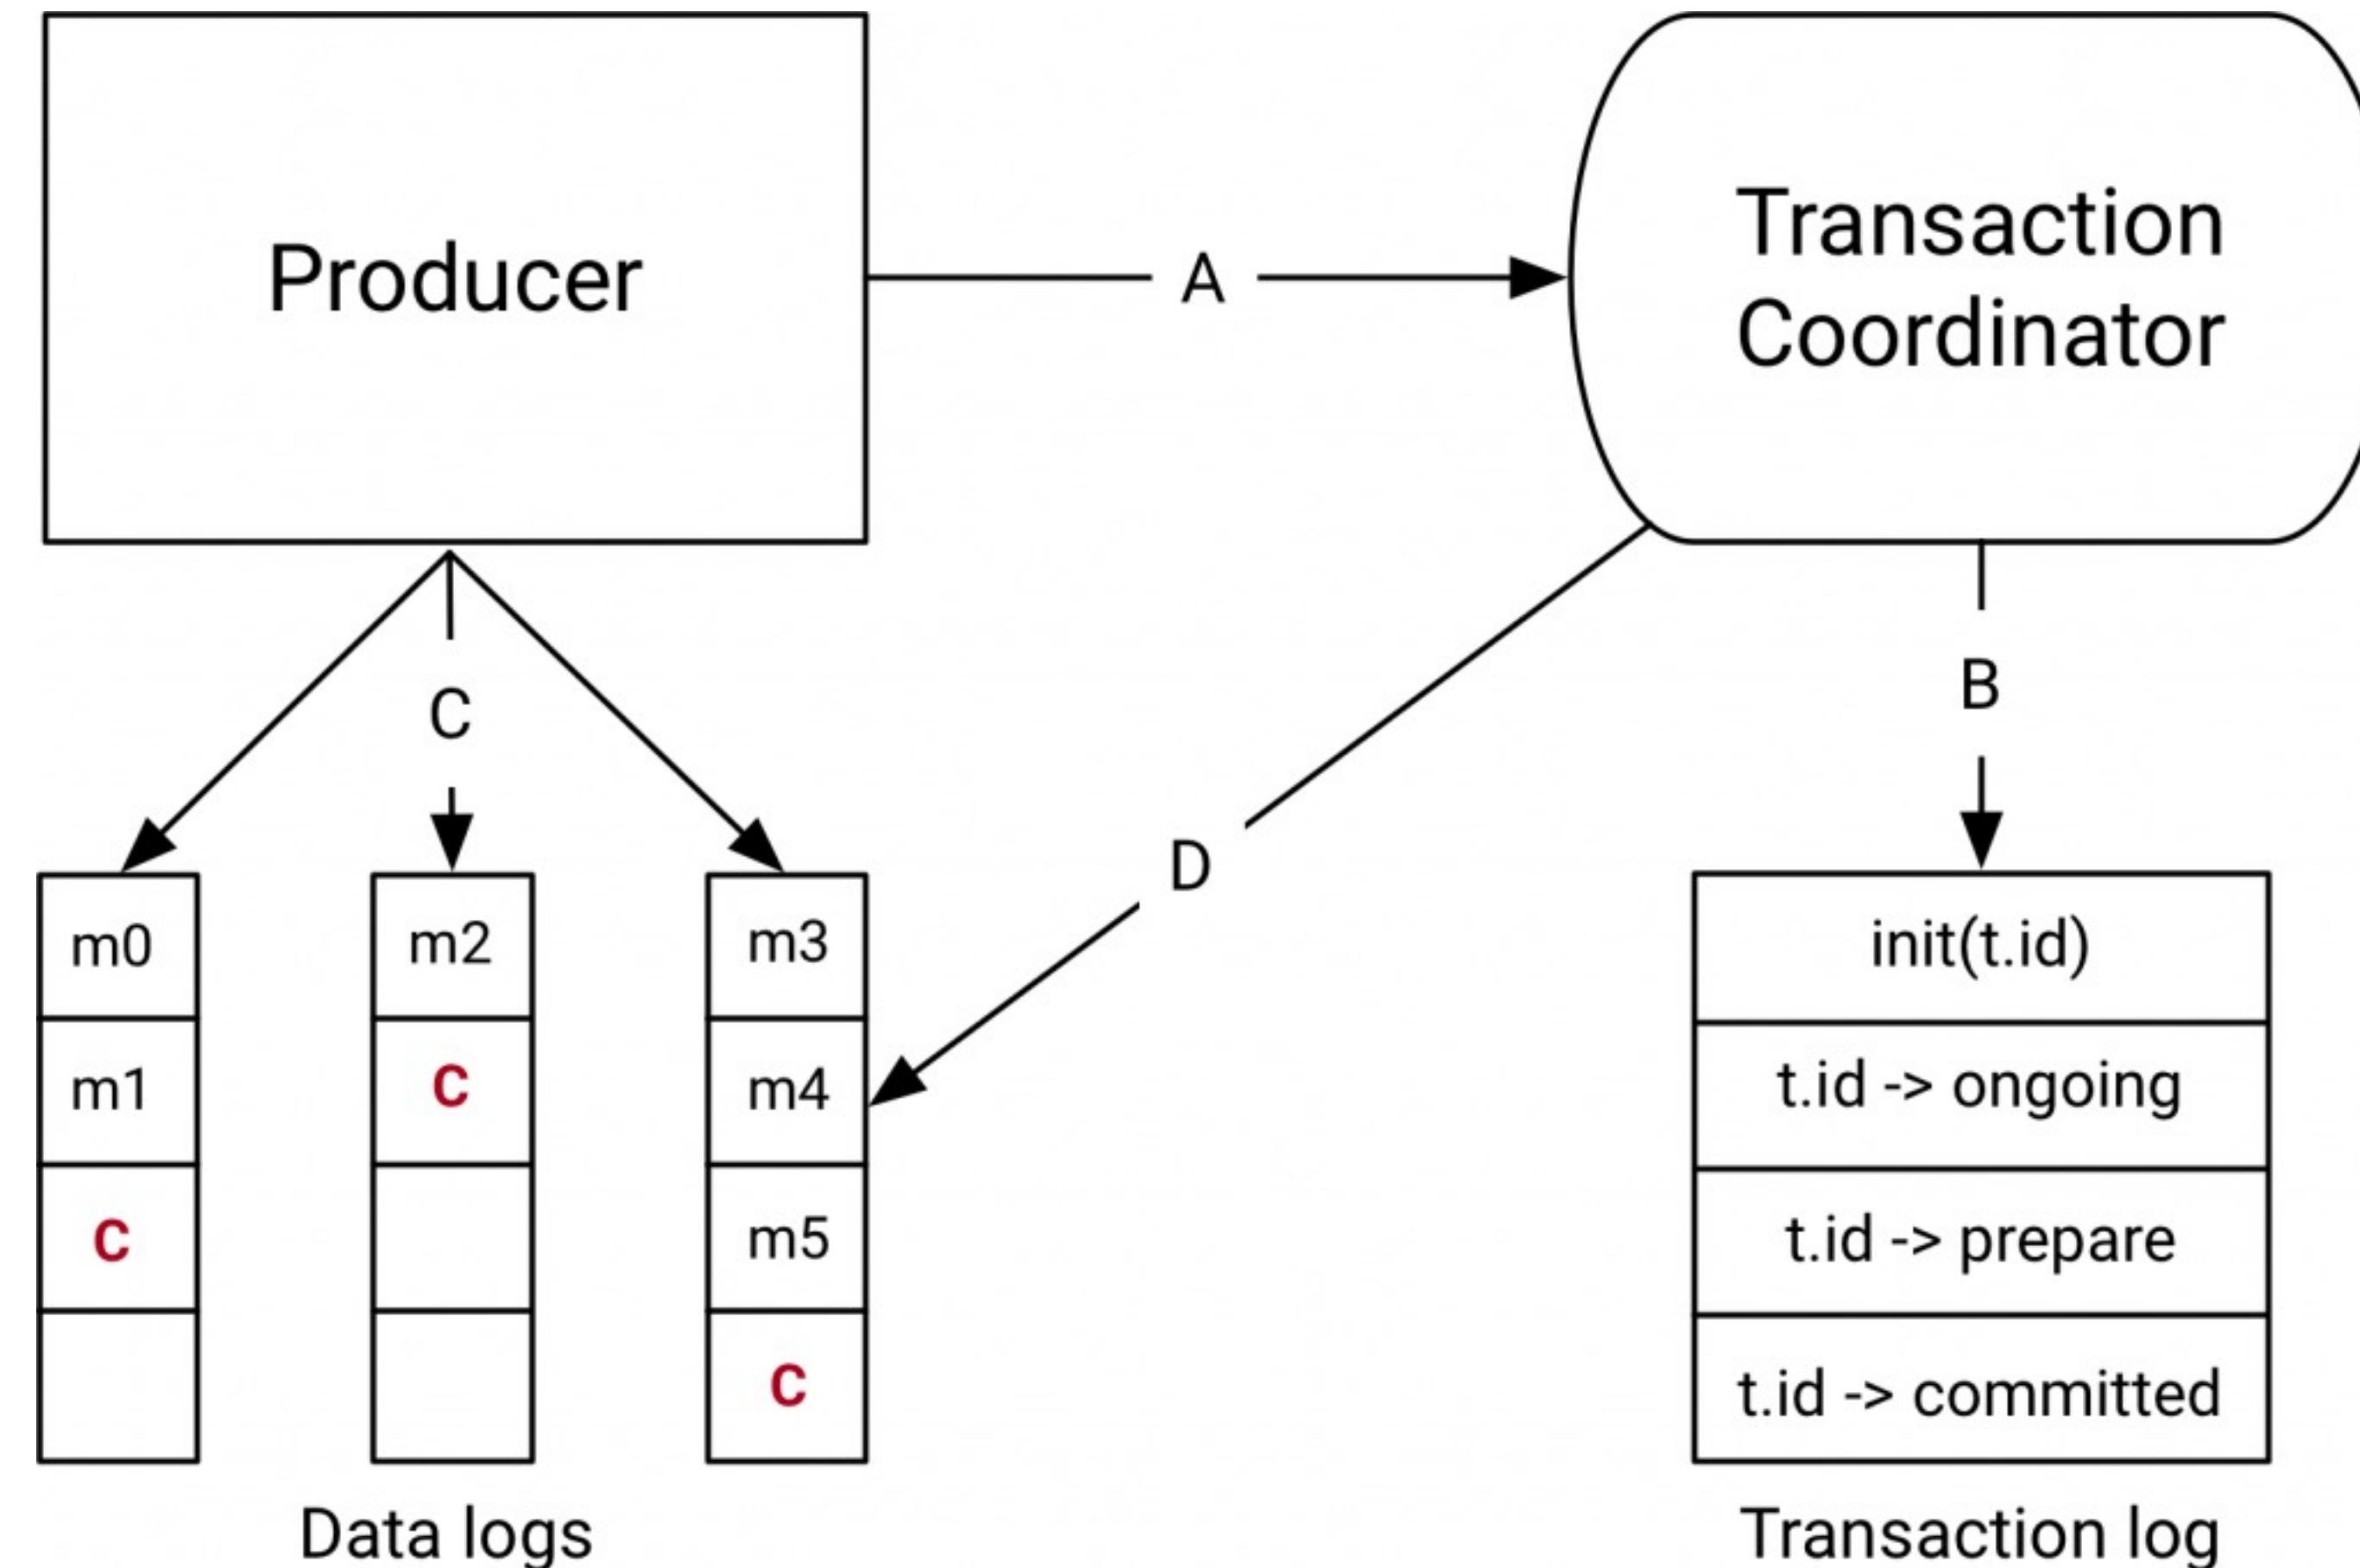
What about the consumer?

For the consumer, we will need to establish:

- `read_committed` - In addition to reading messages that are not part of a transaction, you can also read ones that are, after the transaction is committed
- `read_uncommitted` - Read all messages in offset order without waiting for transactions to be committed. This option is similar to the current semantics of a Kafka consumer.

```
Properties properties = new Properties();
properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ConsumerConfig.GROUP_ID_CONFIG, "sf_giants");
...
properties.put(ConsumerConfig.ISOLATION_LEVEL_CONFIG, "read_committed");
```

How Transactions Work



Enable Idempotence

- We talk about transactions which is a part of ensuring EOS
- We also need to ensure that we handle idempotence
 - *Idempotence* - denoting an element of a set which is unchanged in value when multiplied or otherwise operated on by itself
 - Each batch of messages sent to Kafka will contain a sequence number that the broker will use to dedupe any duplicate send
 - Turn it on in the Producer with enable.idempotence=true

```
properties.put(ProducerConfig.ENABLE_IDEMPOTENCE_CONFIG, true);
```

Avro

Avro

- Created by Doug Cutting; Creator of Hadoop
- Serialization is defined by schema
- Schemas are JSON Based
- Codegen available at command line
- Codegen available by Maven, Gradle, SBT Plugin
- Supports the following languages:
 - C, C++, Java, Perl, Python, Ruby, PHP
- Named after WWI, WWII Aircraft Company, A.V. Roe



Types of Avro Serialization

- Generic
 - Develop code that reflects your schema
- Specific
 - Code generate your class from a schema, an `.avsc` file
 - This is the most common form
- Reflection
 - Auto-create schema from an existing class

Avro Primitive Types

Avro Type	Java Type
null	null
double	double
float	float
int	int
long	long
bool	bool
string	Unicode CharSequence
bytes	Sequence of 8-bit unsigned bytes

Avro Complex Types

```
{  
  "namespace": "com.xyzcorp",  
  "type": "record",  
  "doc" : "An music album",  
  "name": "Album",  
  "fields": [  
    {  
      "name": "name",  
      "type": "string"  
    },  
    {  
      "name": "yearReleased",  
      "type": [  
        "int",  
        "null"  
      ]  
    }  
  ]  
}
```

Avro Field Options

Avro Field	Description
name	Name of the Field
doc	Documentation
type	Type of the Field
default	Default Value
order	What order does this impact record
aliases	Any other names

Avro Array

List of Strings, but can be any type for items

```
{"type": "array", "items": "string"}
```

Avro Map

All keys in an Avro map are string

```
{"type": "map", "values": "long"}
```

Avro Enum

When autogen in Java, RainbowColors will be an enum

```
{ "type" : "enum",
  "name" : "RainbowColors",
  "doc" : "Colors of the Rainbow",
  "symbols" : ["RED", "ORANGE", "YELLOW",
    "GREEN", "BLUE", "INDIGO",
    "VIOLET"]}
```

Full Specification for Avro

- Read more about Avro, <https://avro.apache.org/docs/current/spec.html>
- Includes other types and feature:
 - UUID
 - Timestamp
 - Date/Time
 - Sorting
 - Codecs

Running Avro Tools

Avro Tools is a jar file that can be executed to generate Java from schema

```
java -jar avro-tools-1.8.2.jar compile schema <schema file> <destination>
```

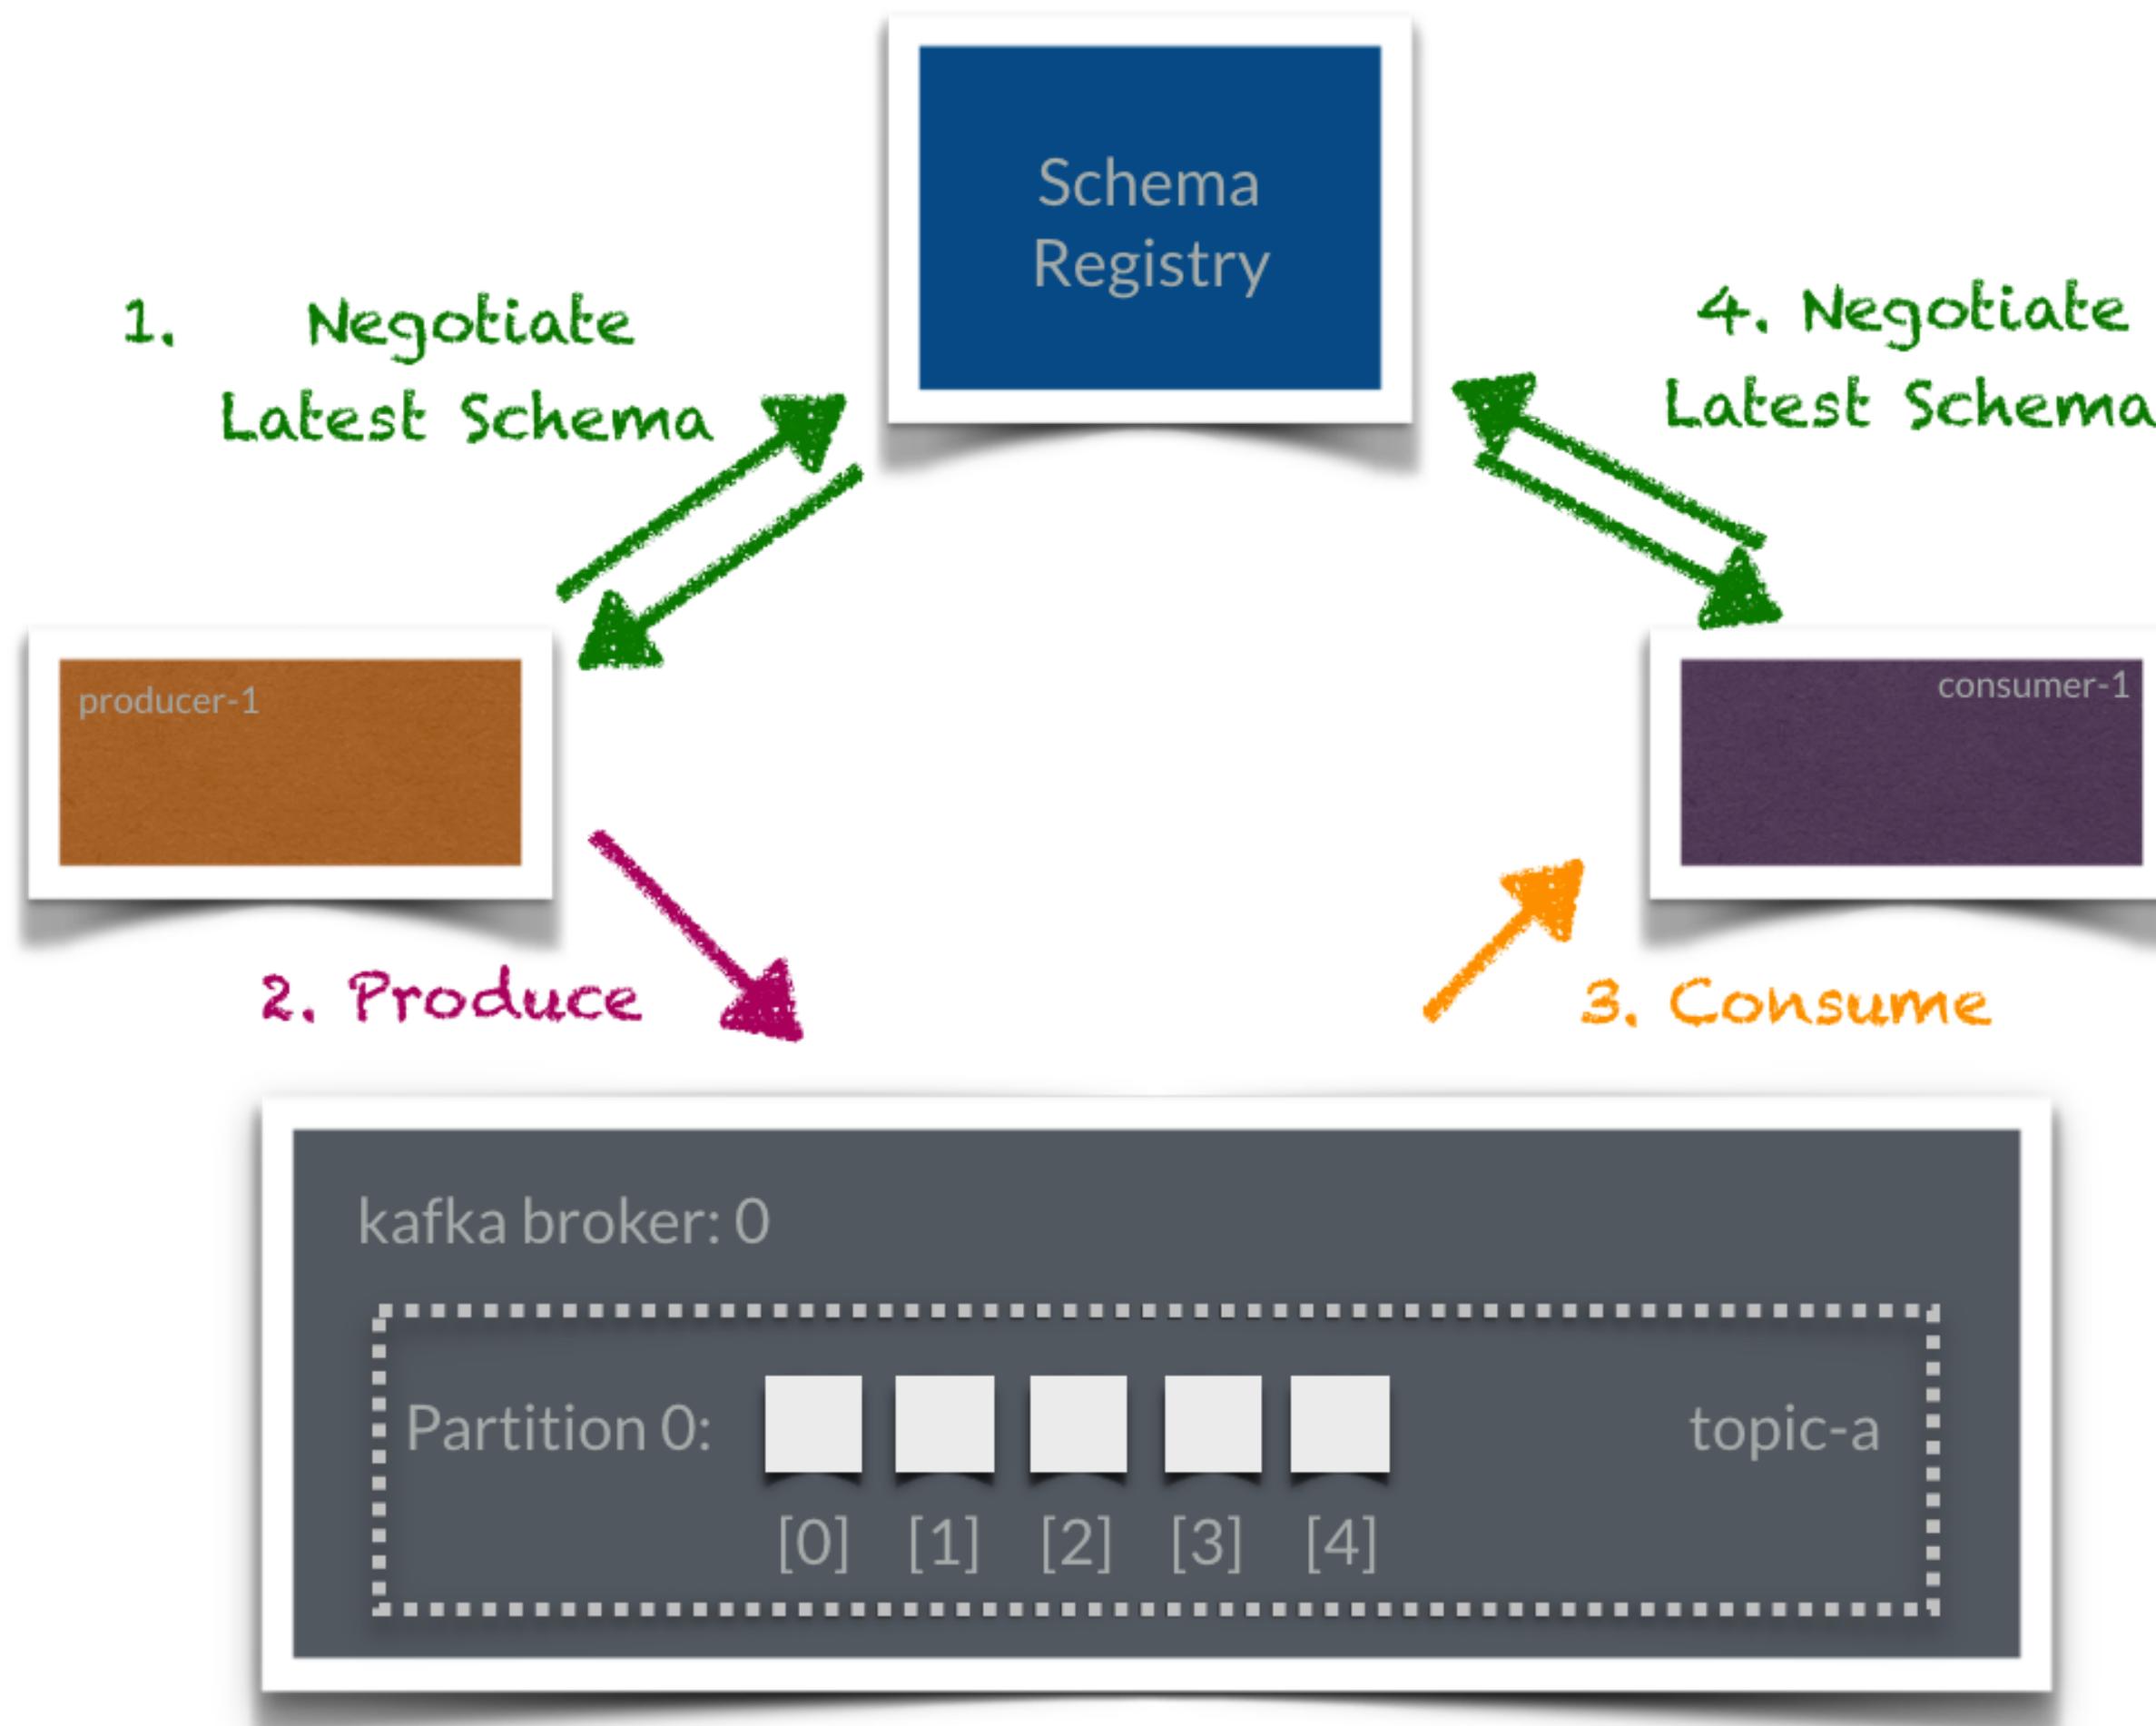
Avro Tools is available for most IDEs like Maven, and Gradle

Schema Registry

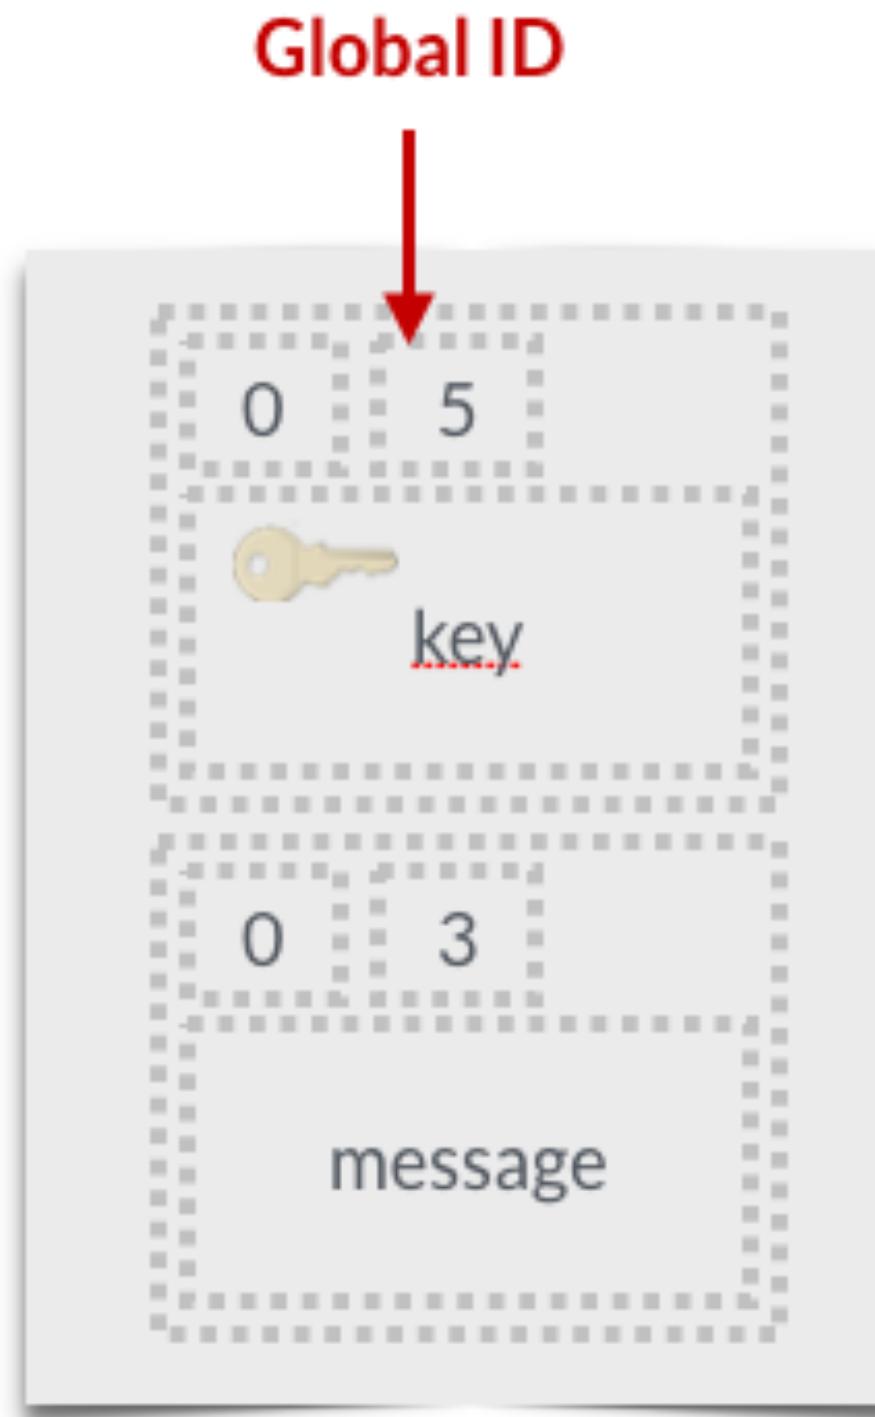
Schema Registry

- Service that constrains and typechecks messages before sending
- Stores schemas in registry for and assigns a global ID
- Ensures *backward* and *forward* compatibility
- Does so through Serializers and Deserializers
- Automatically done through Java Kafka Client API

Schema Registry



Schema Registry Message Format



- The first byte is a zero byte
- The second byte is the global identifier of the Schema that registered in the Schema Registry
- The Global ID is used to as a means of communicating which Schema to use with messages

Registering a New Schema Manually

```
POST /subjects/my_topic-key/versions HTTP/1.1
Host: schemaregistry.example.com
Accept: application/vnd.schemaregistry.v1+json,
         application/vnd.schemaregistry+json,
         application/json
```

```
{
  "schema": {
    "type": "record",
    "name": "test",
    "fields": [
      {
        "type": "string",
        "name": "field1"
      },
      {
        "type": "int",
        "name": "field2"
      }
    ]
  }
}
```

- Posts to a subject
- A subject is combination of
 - Topic
 - -key or -value
- Example of a subject
 - stocks-key
 - sales-value



Note: We aren't going manually do this we just wanted to show what an avro producer will do

Getting Response from the Registry

HTTP/1.1 200 OK

Content-Type: application/vnd.schemaregistry.v1+json

{"id":1}



The {"id":1} is the global id, this is used as reference to this schema

Schema Registry Java Producer

```
Properties properties = new Properties();
properties.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,
    "localhost:9092");
properties.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
    StringSerializer.class);
properties.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
    io.confluent.kafka.serializers.KafkaAvroSerializer.class);
properties.setProperty("schema.registry.url", "http://localhost:8081");
```

- What makes schema registry work is the Serializer
- You must add `schema.registry.url` and specify the location of the registry

Example of Sending Custom Avro Object

```
Album album = new Album("Purple Rain", "Prince", 1984,  
    Arrays.asList("Purple Rain", "Let's go crazy));
```

```
ProducerRecord<String, Album> producerRecord =  
    new ProducerRecord<>("music_albums", "Prince", album);  
  
producer.send(producerRecord);
```

- The ProducerRecord's value is accepting an actual Java Object
- We can just send the object just like any other type

Schema Registry Consumer

```
Properties properties = new Properties();
properties.put(ConsumerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
properties.put(ConsumerConfig.GROUP_ID_CONFIG, "my_group");
properties.put(ConsumerConfig.KEY_DESERIALIZER_CLASS_CONFIG,
    "org.apache.kafka.common.serialization.StringDeserializer");
properties.put(ConsumerConfig.VALUE_DESERIALIZER_CLASS_CONFIG,
    io.confluent.kafka.serializers.KafkaAvroDeserializer.class);
properties.setProperty("schema.registry.url", "http://localhost:8081");
properties.setProperty("specific.avro.reader", "true");
```

- What makes schema registry work is the Deserializer
- You must add `schema.registry.url` and specify the location of the registry
- You must add `specific.avro.reader` and specify that you used specific avro mode

Receiving a Custom Object from Kafka

```
while (true) {
    ConsumerRecords<String, Album> records =
        consumer.poll(Duration.of(500, ChronoUnit.MILLIS));
    for (ConsumerRecord<String, Album> record : records) {
        System.out.format("offset: %d\n", record.offset());
        System.out.format("partition: %d\n", record.partition());
        System.out.format("timestamp: %d\n", record.timestamp());
        System.out.format("timeStampType: %s\n",
                          record.timestampType());
        System.out.format("topic: %s\n", record.topic());
        System.out.format("key: %s\n", record.key());
        Album a = record.value();
        System.out.format("value: %s\n", a.getTitle());
        System.out.format("value: %s\n", a.getArtist());
    }
}
```

Streaming Concepts



TK45
115m →
ÜNIKC.

TK45
120m
ÜNIKC.

Stream Processing

Endless supply of data

“Replayable”

Real Time Processing for Fraud Detection, High End Sales Detection, Internet of Things, and More.

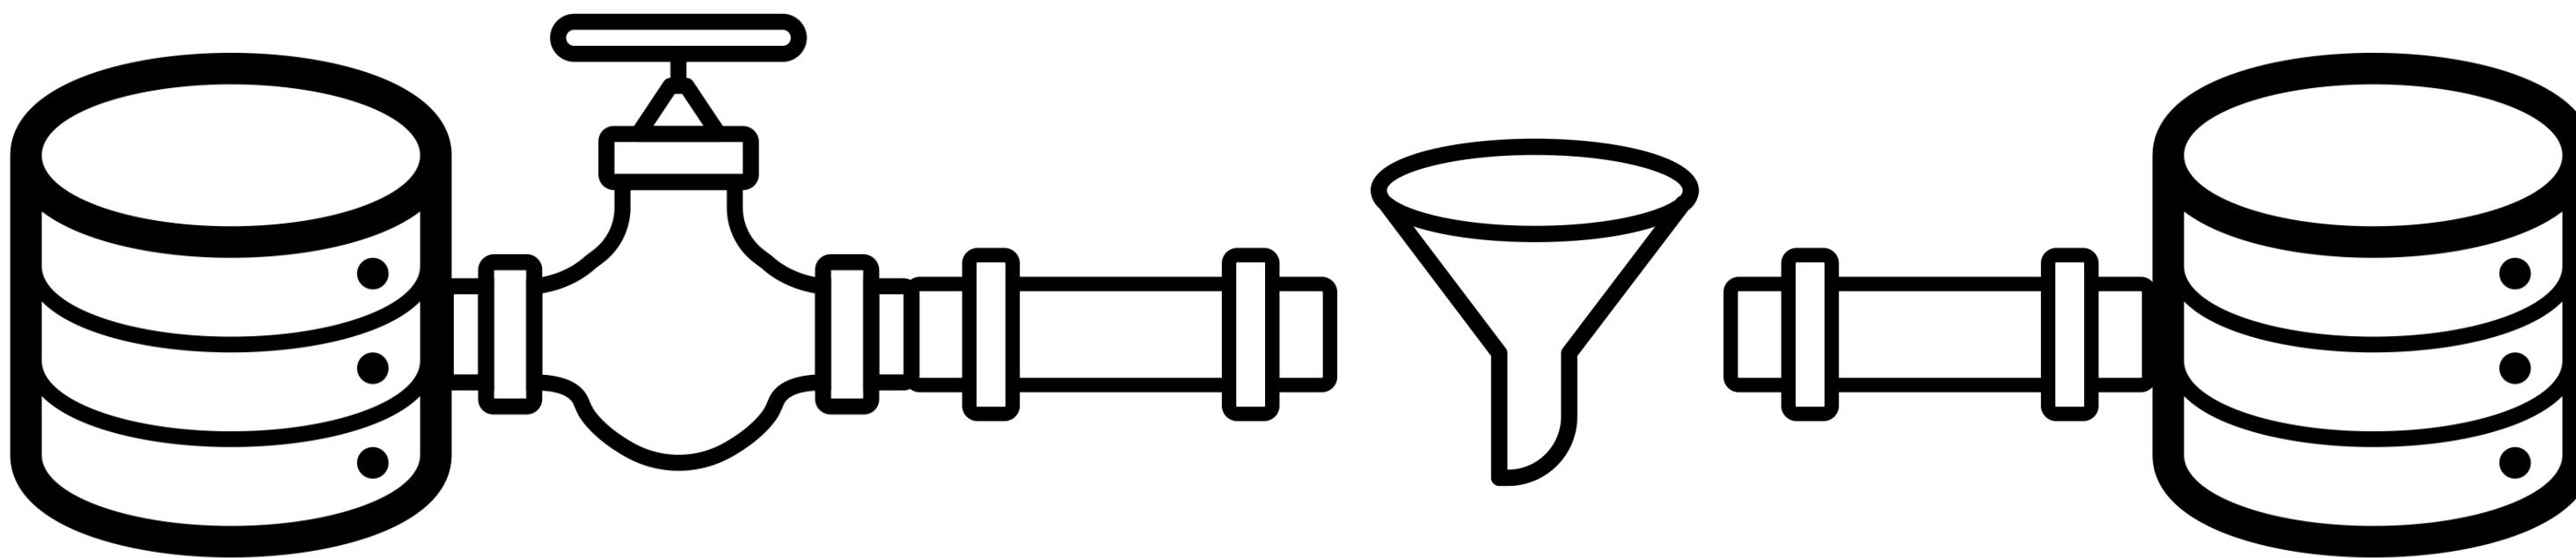
Will require built-in accumulator table to perform real time data processing, like counting, and grouping

NY

100.00

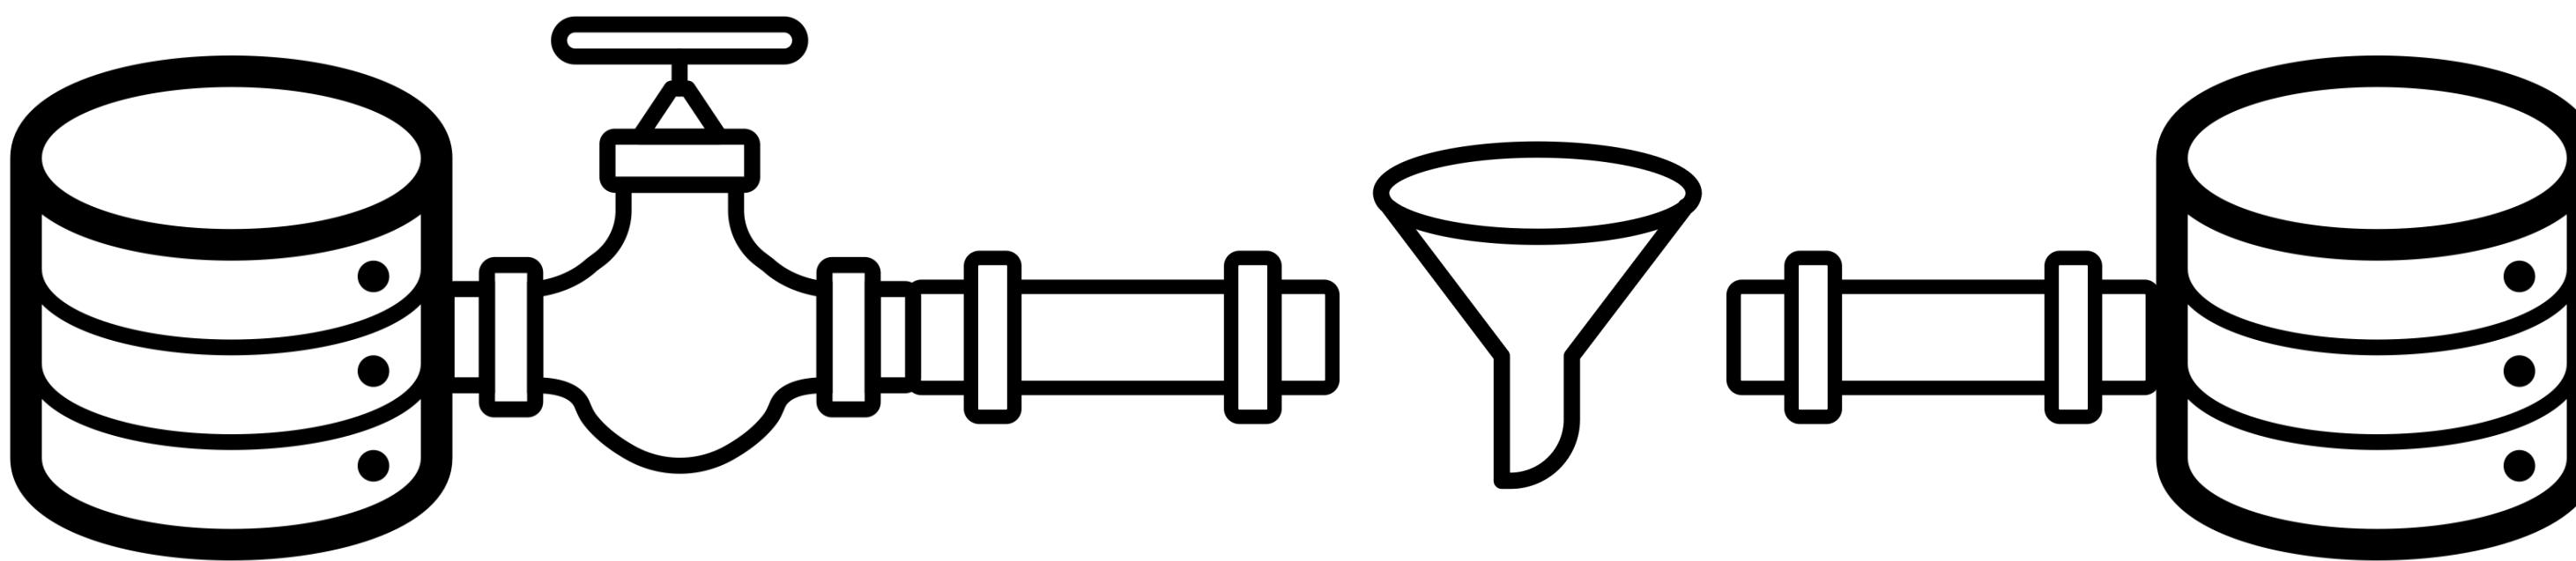
Filtering

NY
.....
100.00



value > 10000

OH
.....
20000

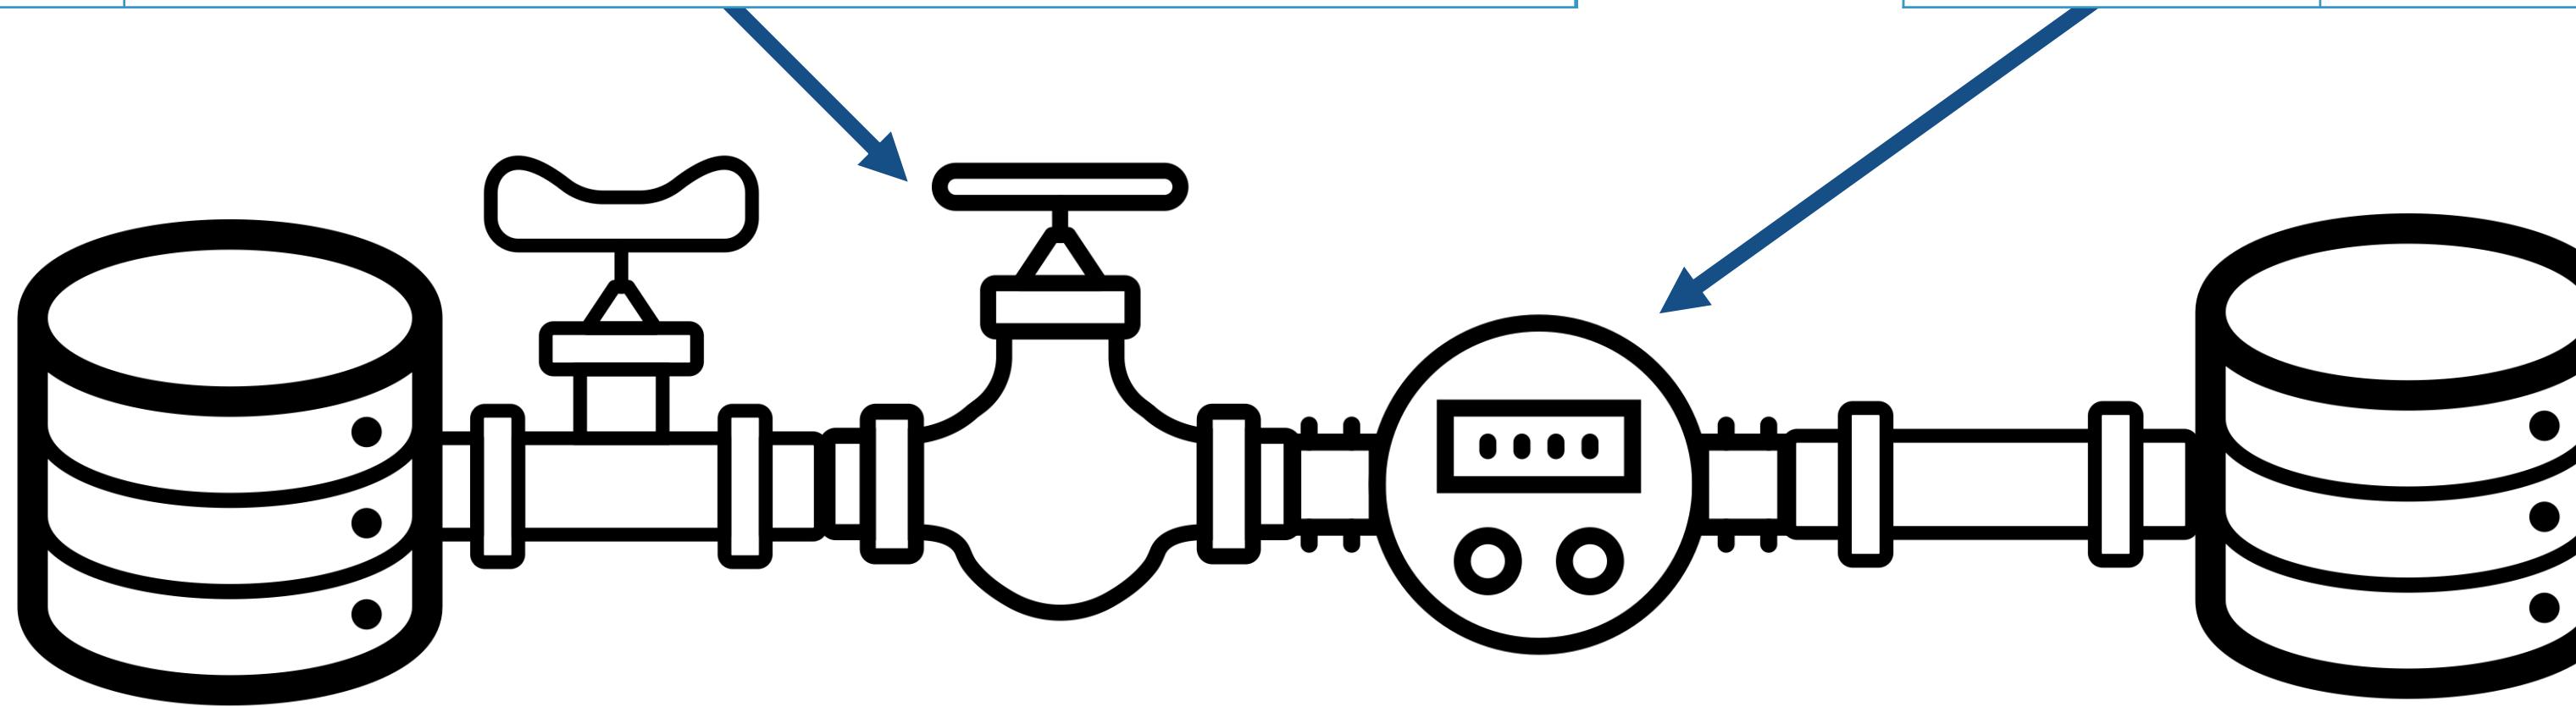


value > 10000

Aggregating

Key	Value
OH	100.0

Key	Value
OH	100.0

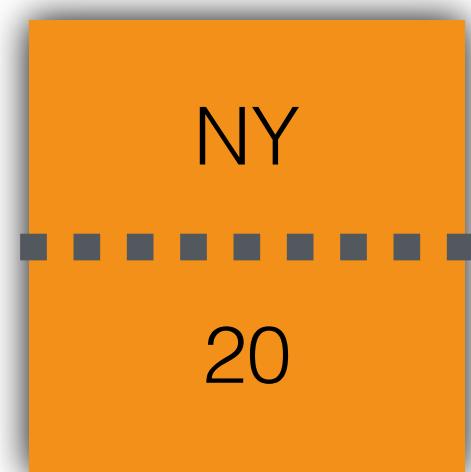
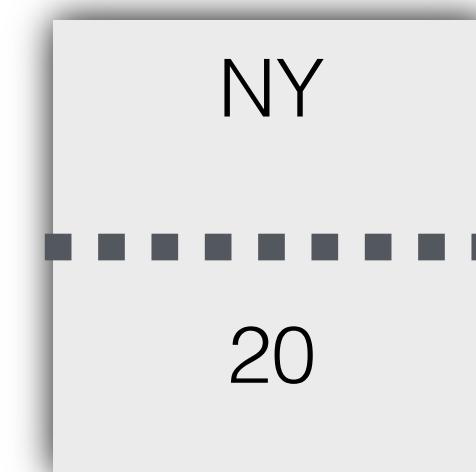
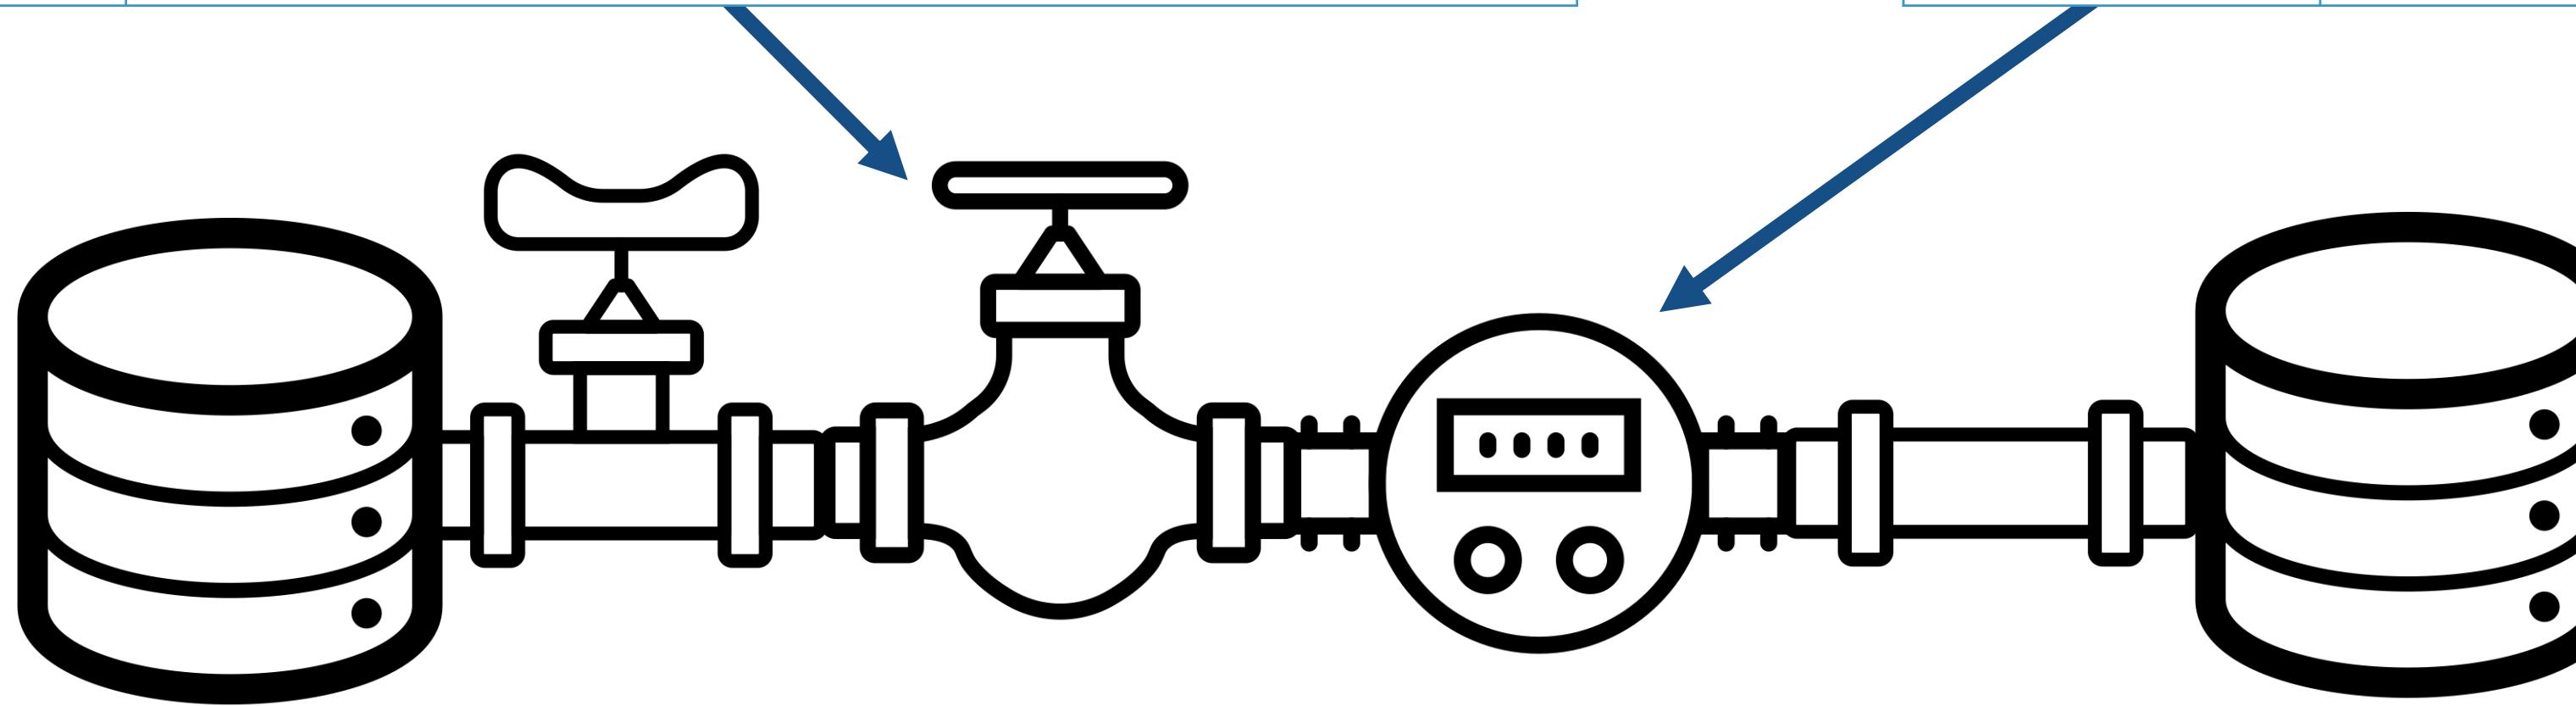


OH
100.00

OH
100.00

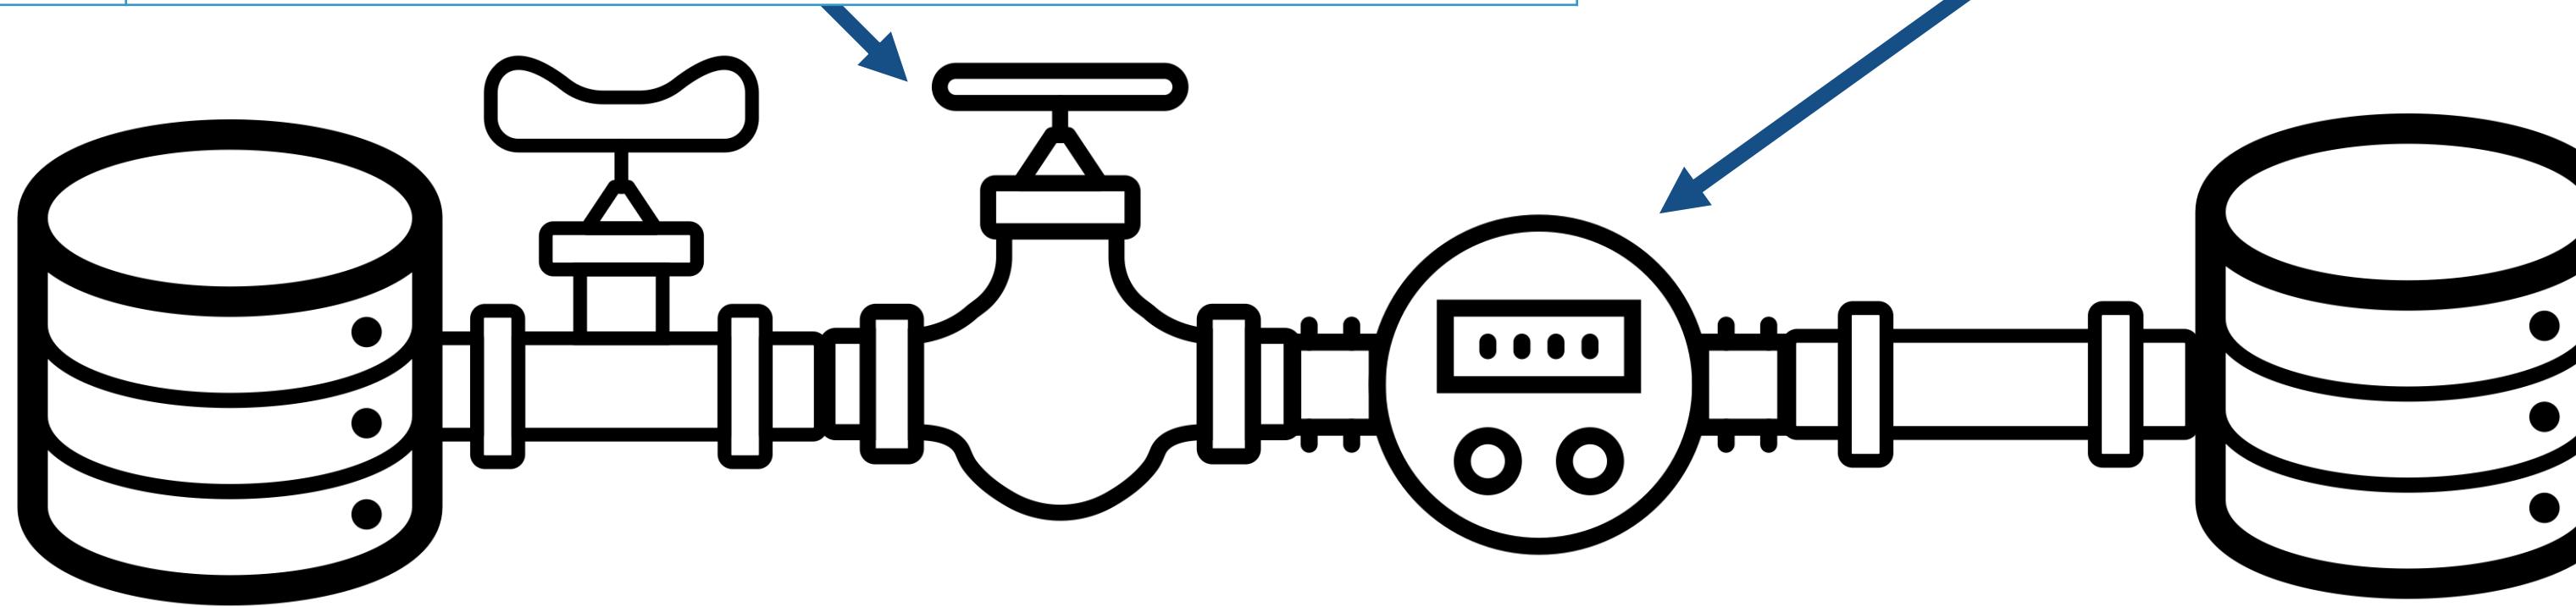
Key	Value
OH	100.0
NY	20.0

Key	Value
OH	100.0
NY	20.0



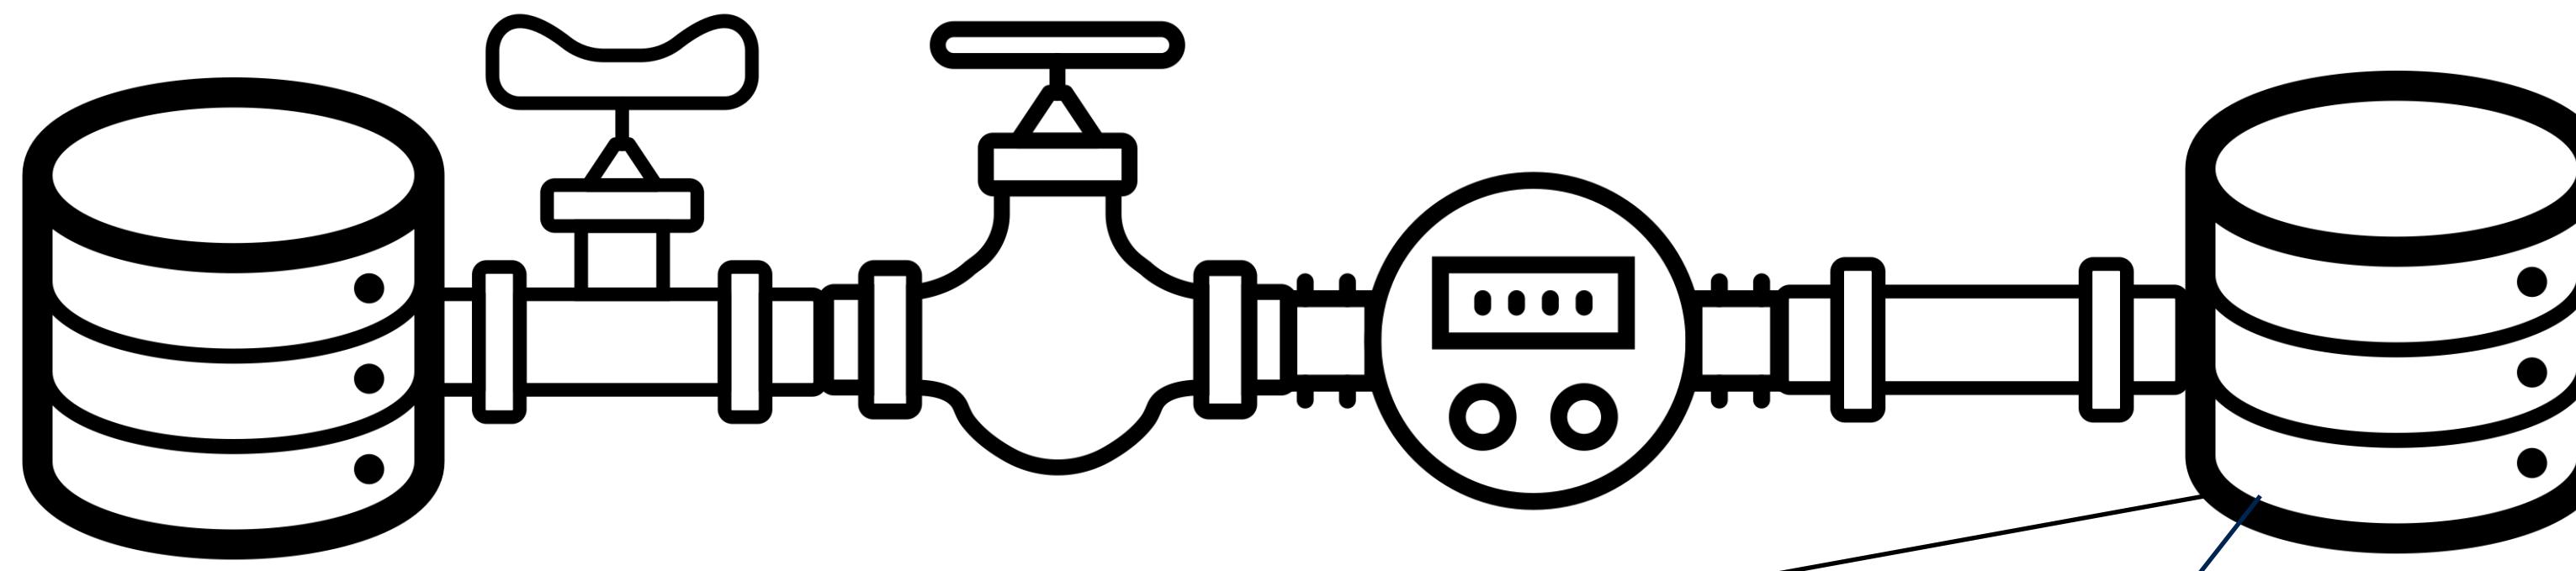
Key	Value
OH	100.0
NY	20.0, 60.0

Key	Value
OH	100.0
NY	80.0



NY
.....
60

NY
.....
80.0

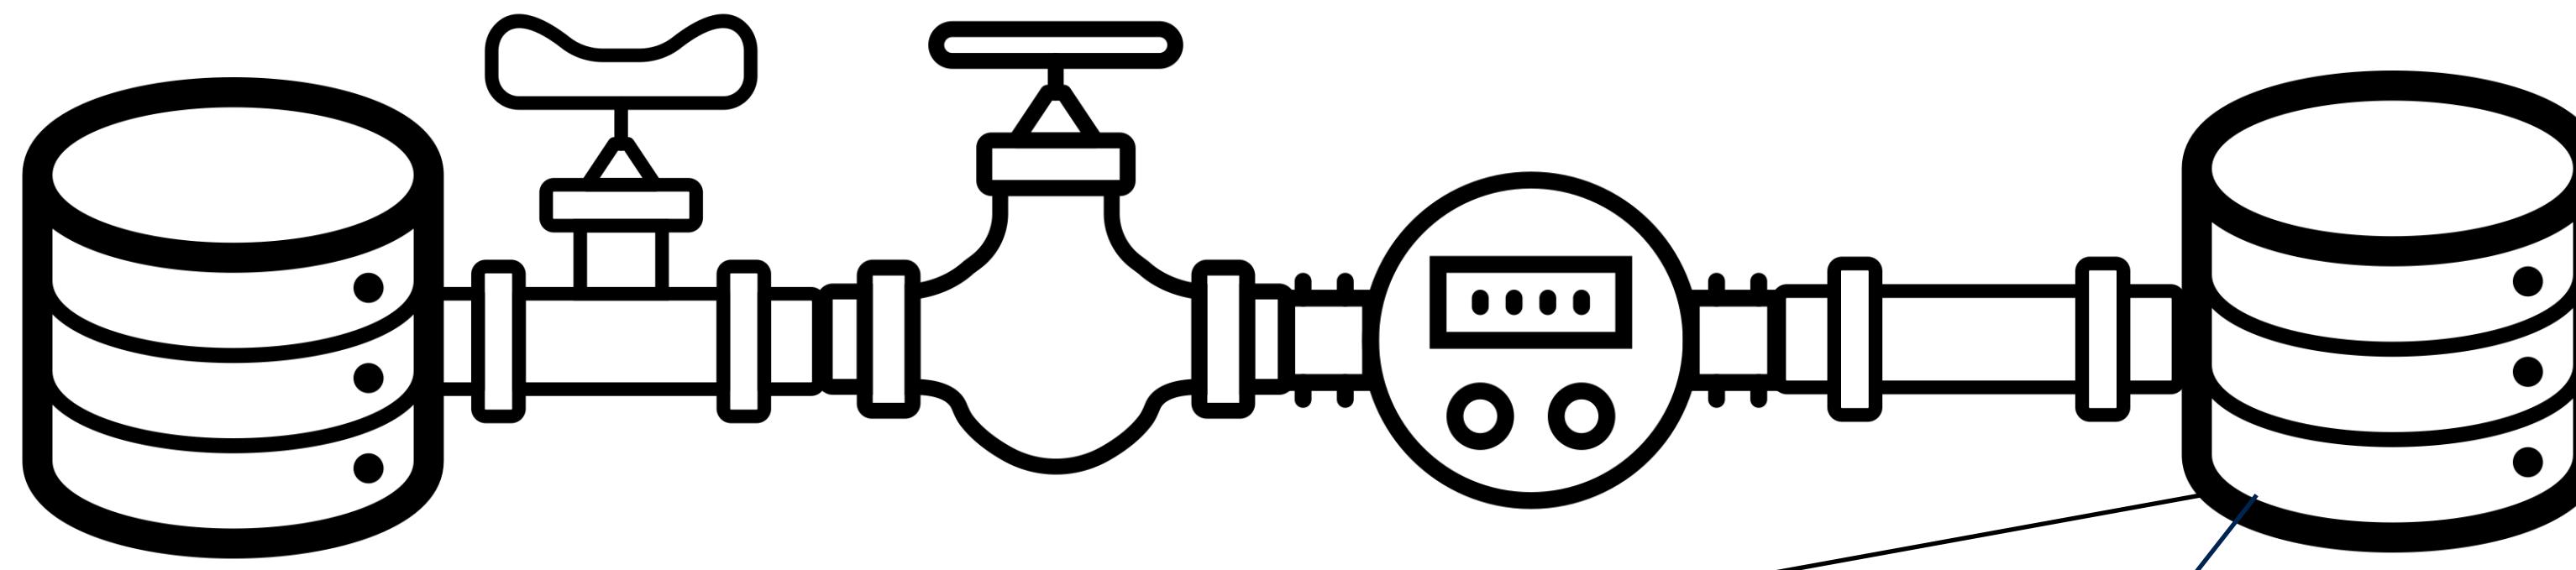


Offset	Key	Value
0	OH	100.0

Partition 0

Offset	Key	Value
0	NY	20.0
1	NY	80.0

Partition 1



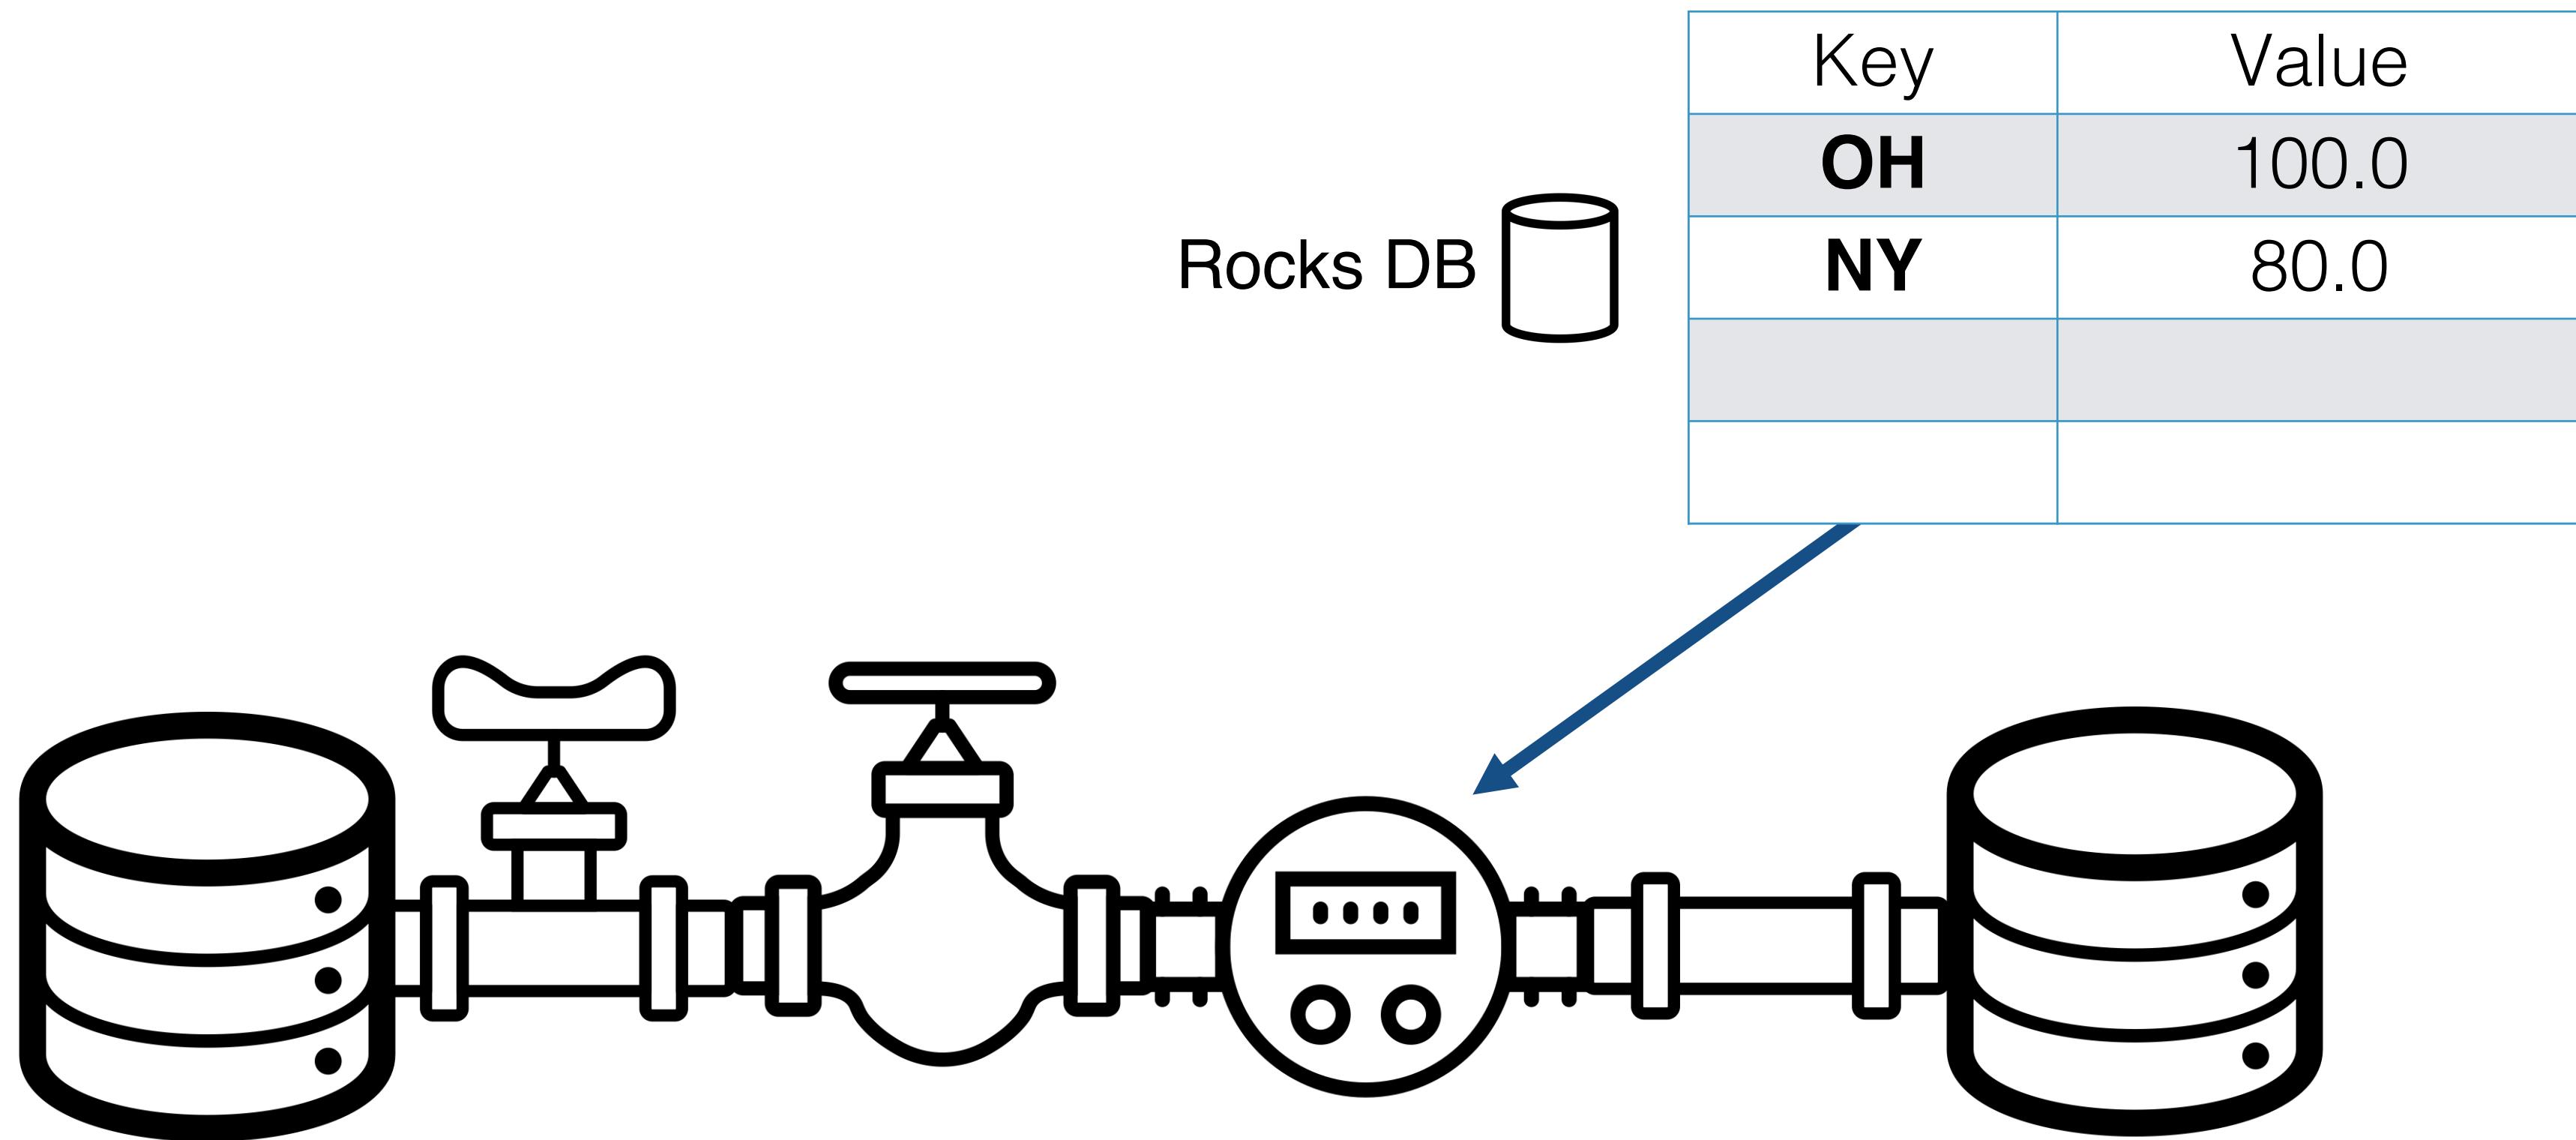
Offset	Key	Value
0	OH	100.0

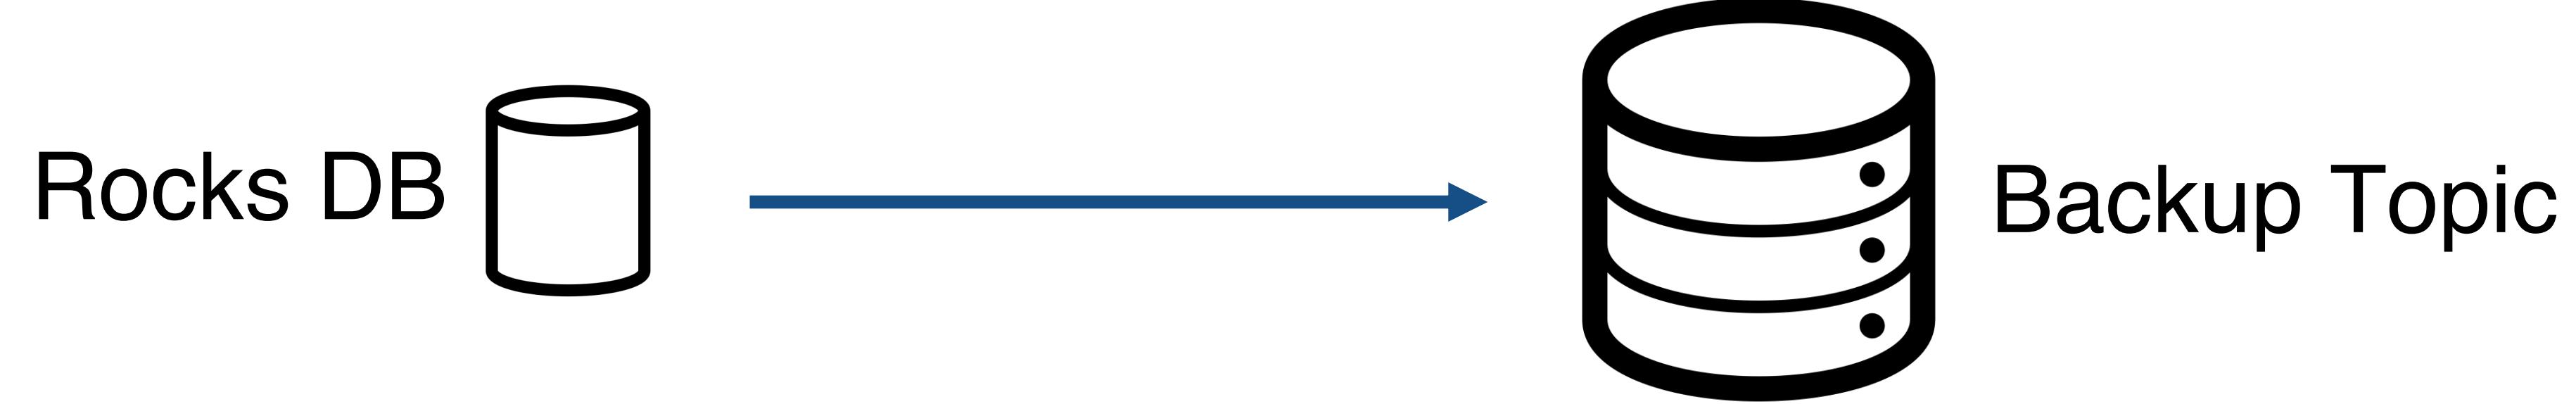
Partition 0

Offset	Key	Value
1	NY	80.0

Partition 1

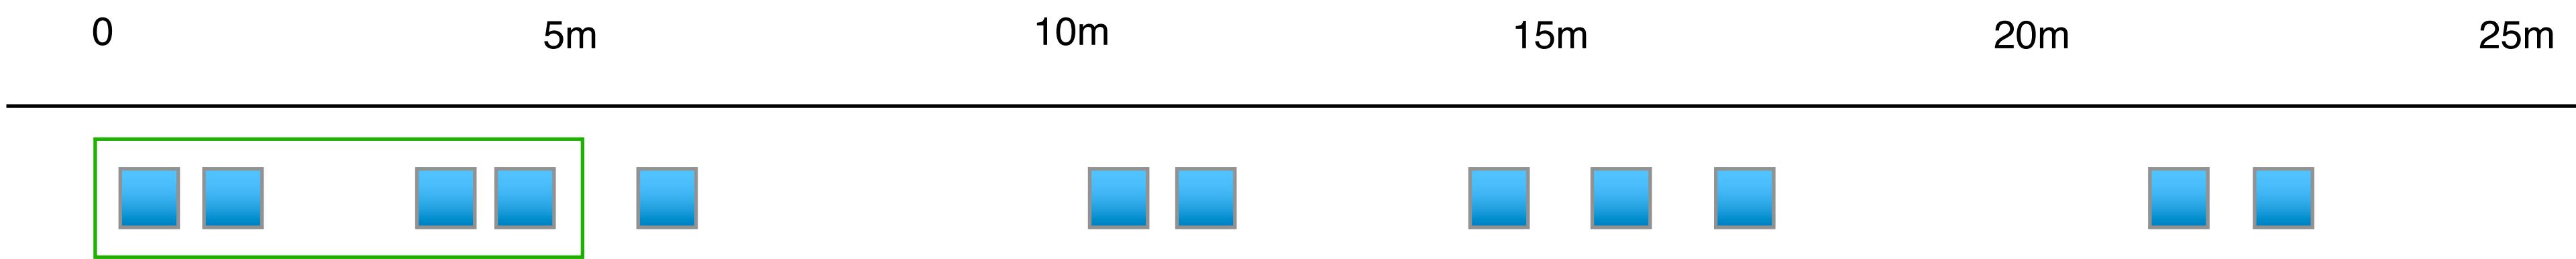
After Compaction



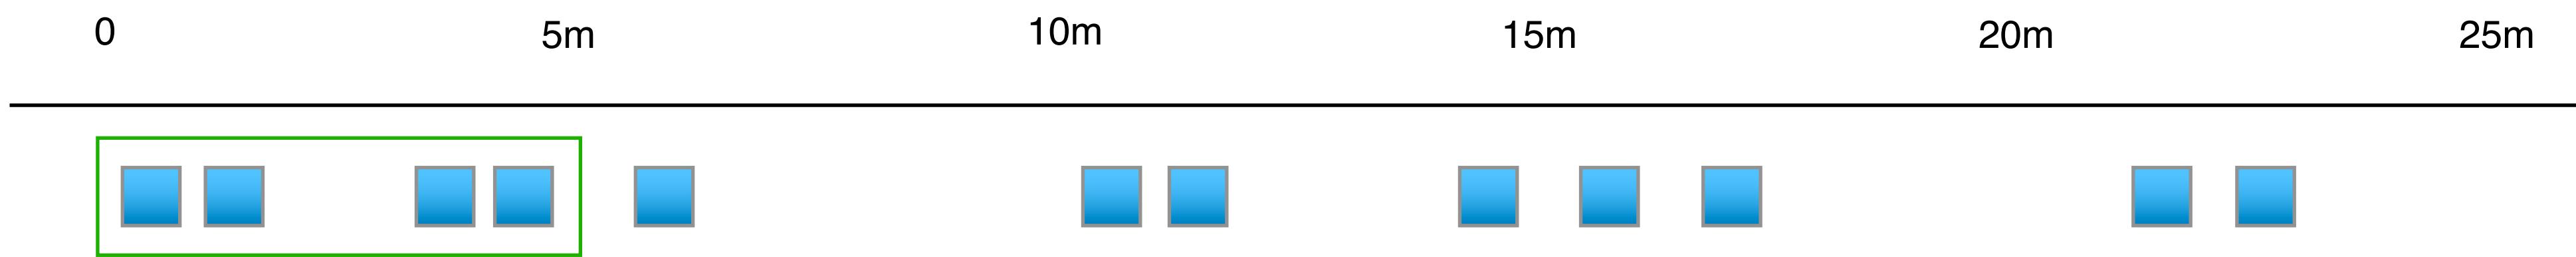


Windowing

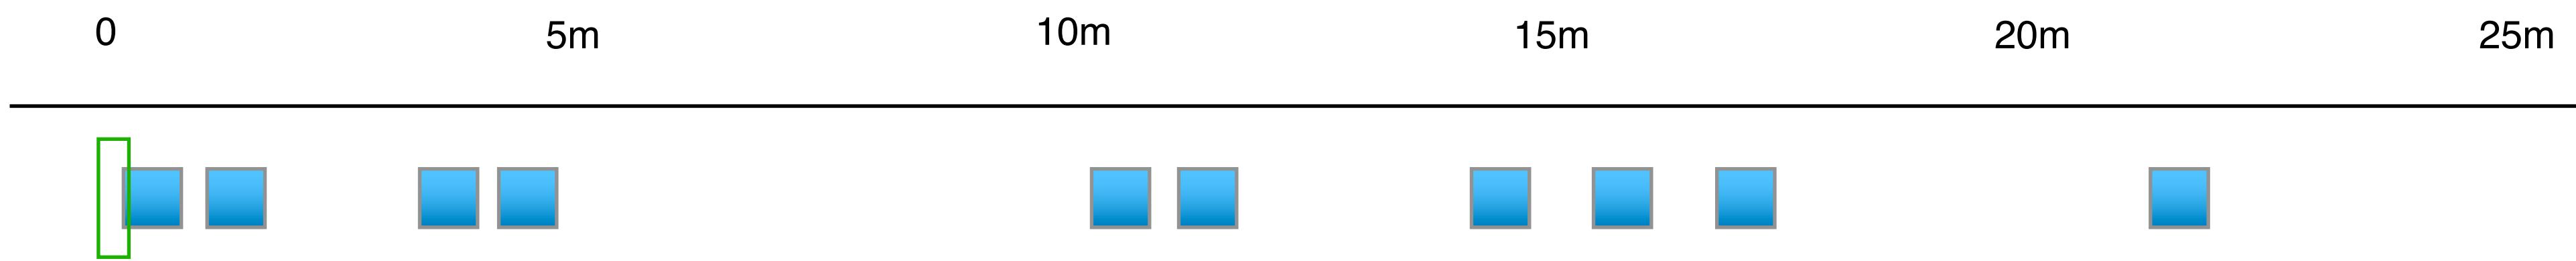
Tumbling Window



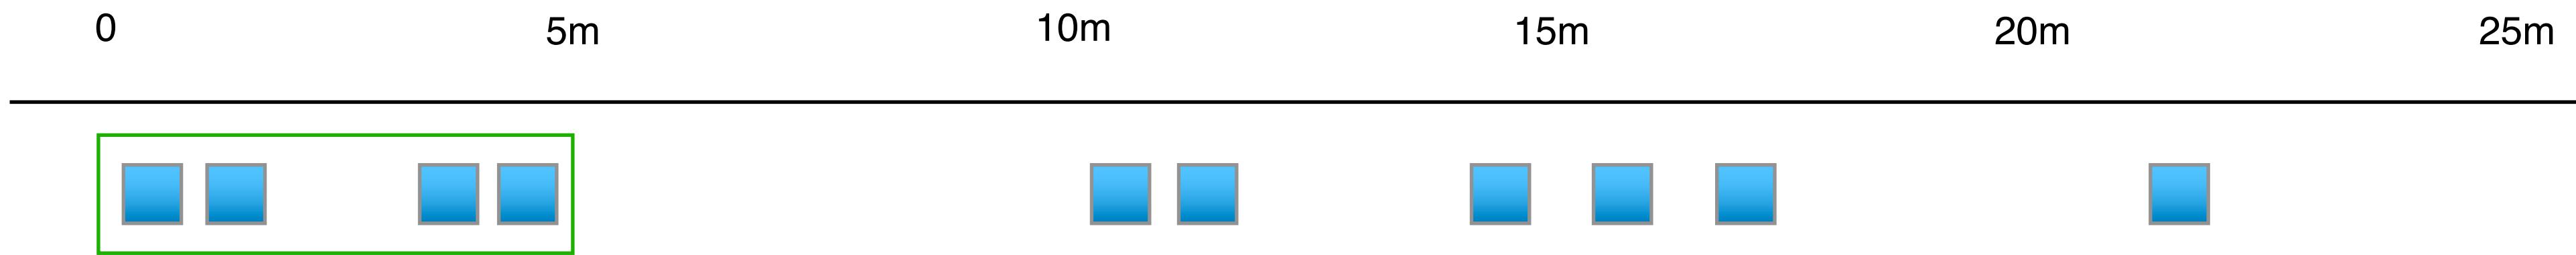
Hopping Window



Session Window

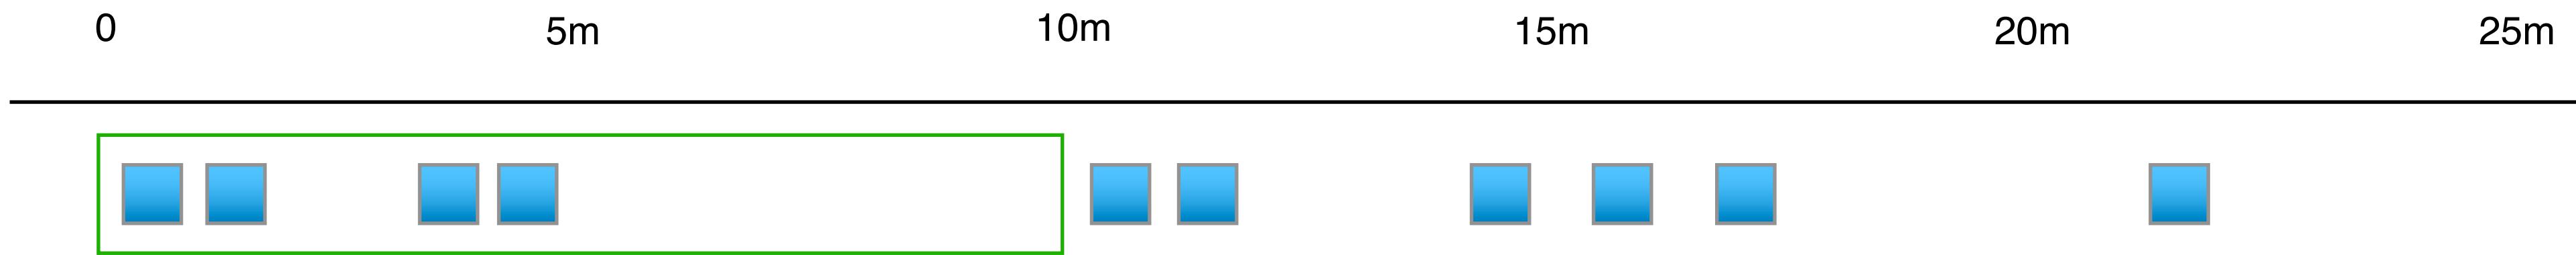


Session Window



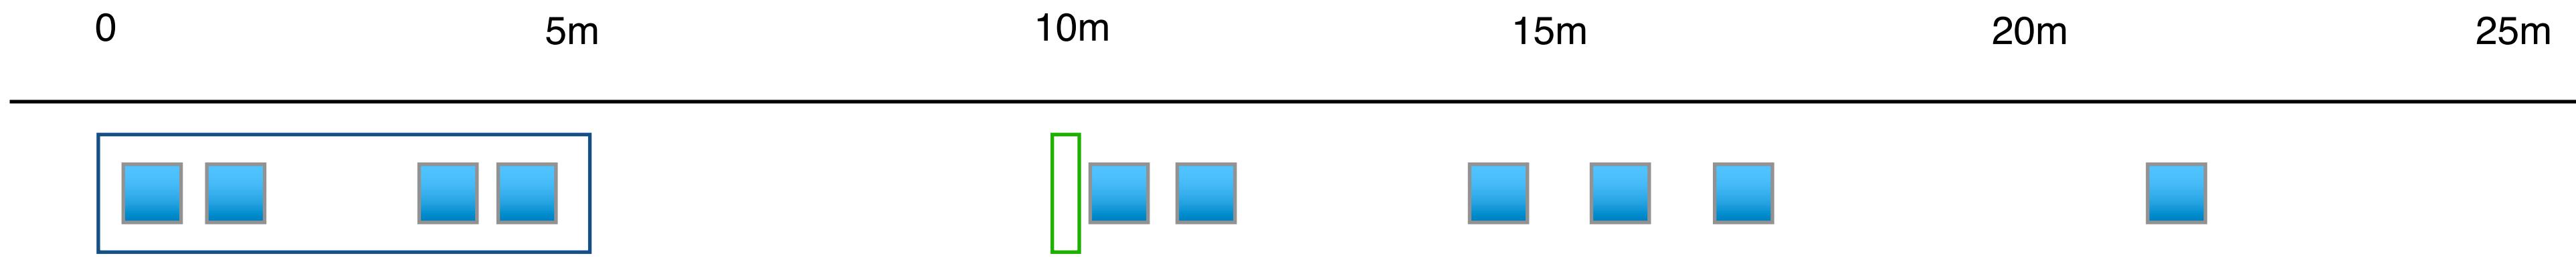
the last element has occurred; now we wait

Session Window



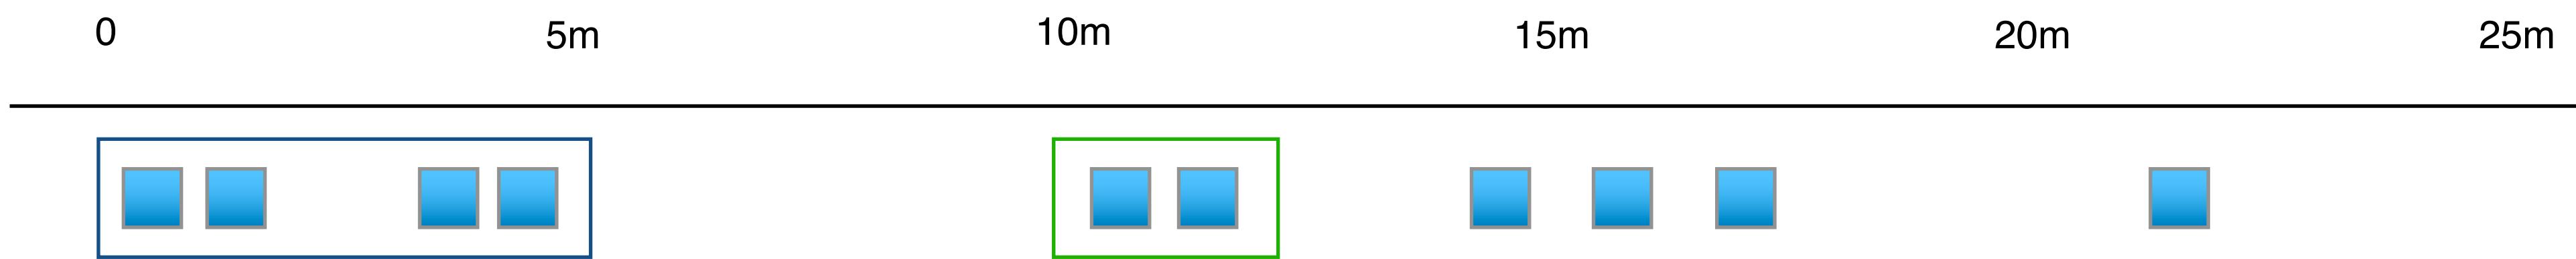
waited 5 minutes = make it a window

Session Window



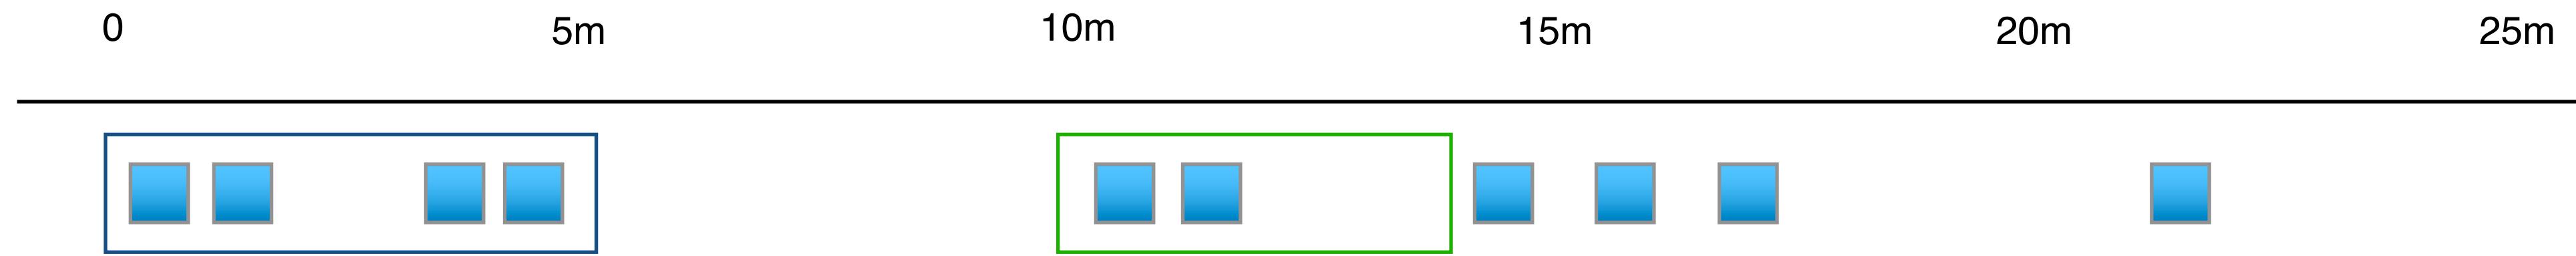
we have a new element; we start a new window

Session Window



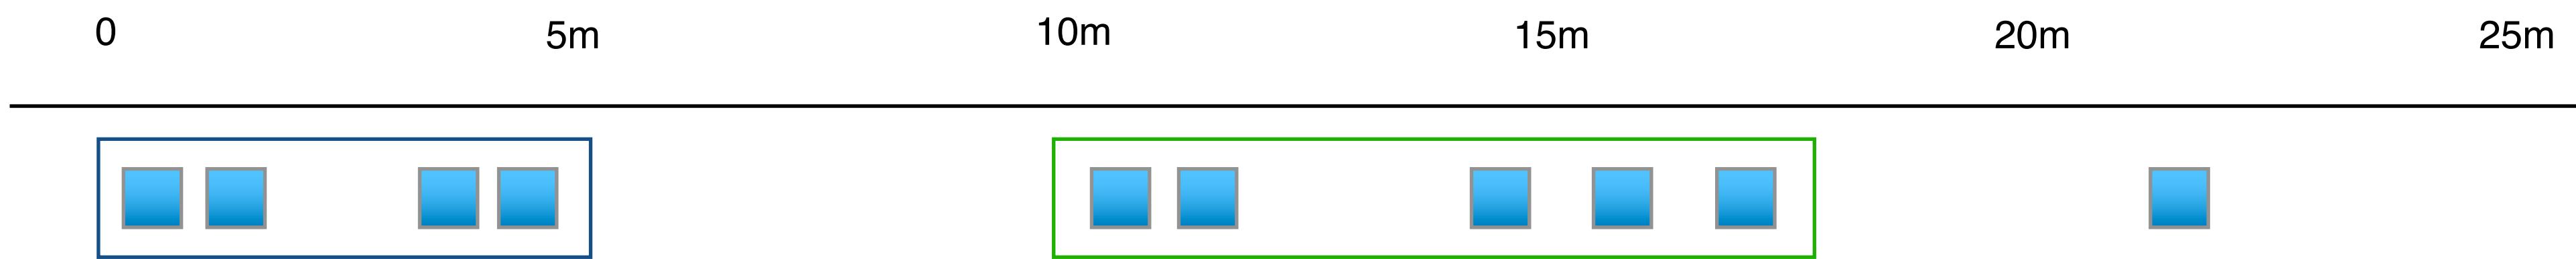
we wait for 5 minutes

Session Window



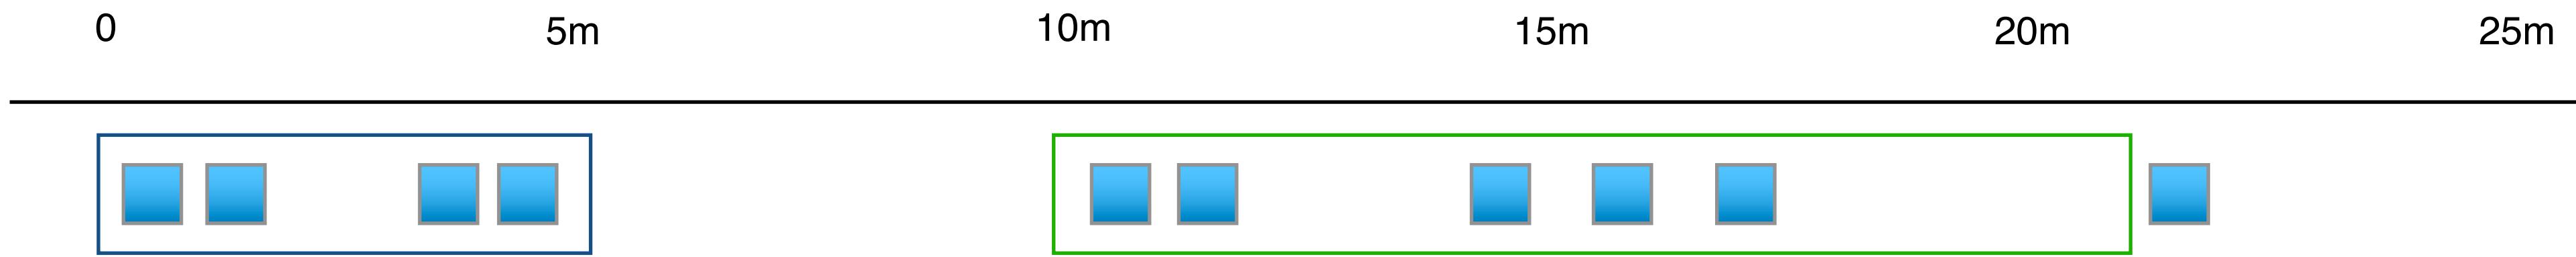
2 minutes elapsed; no windowing yet

Session Window



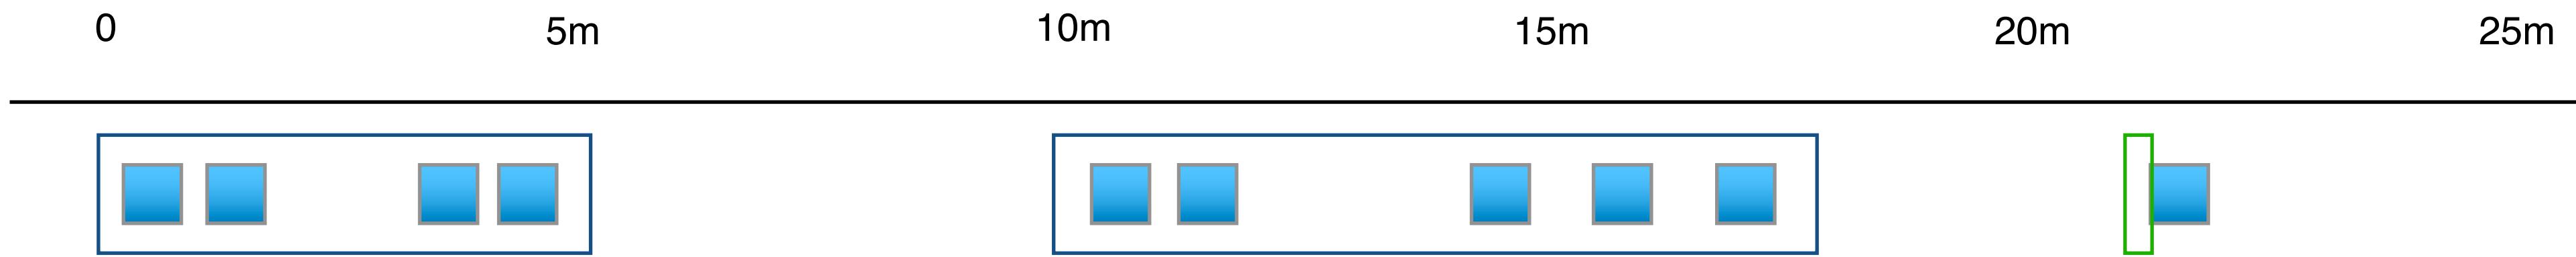
we wait five minutes

Session Window



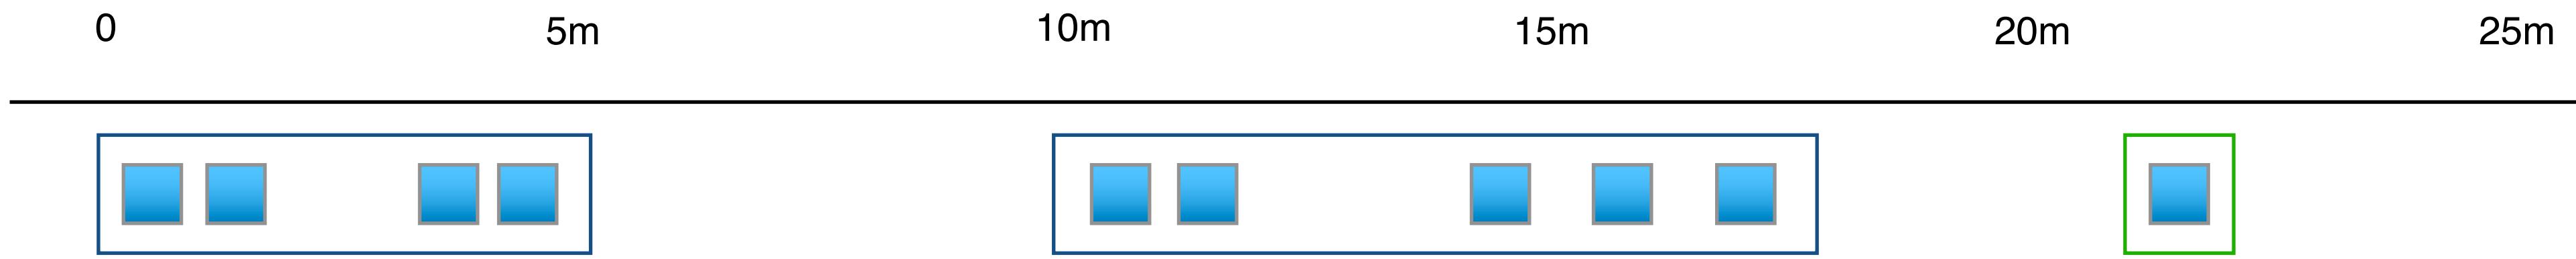
no new elements; we make it a window

Session Window



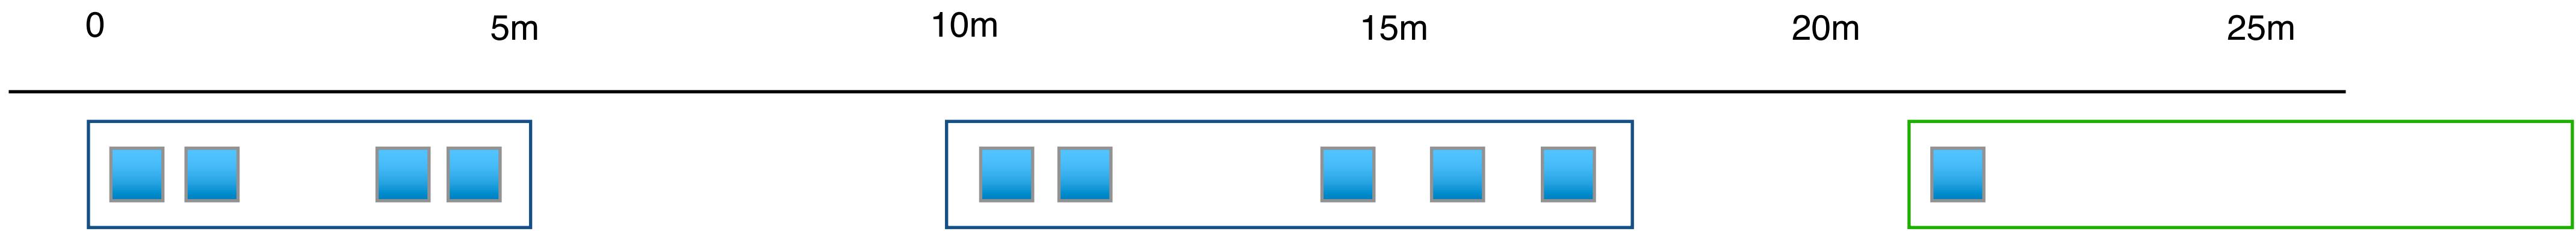
we see something new

Session Window



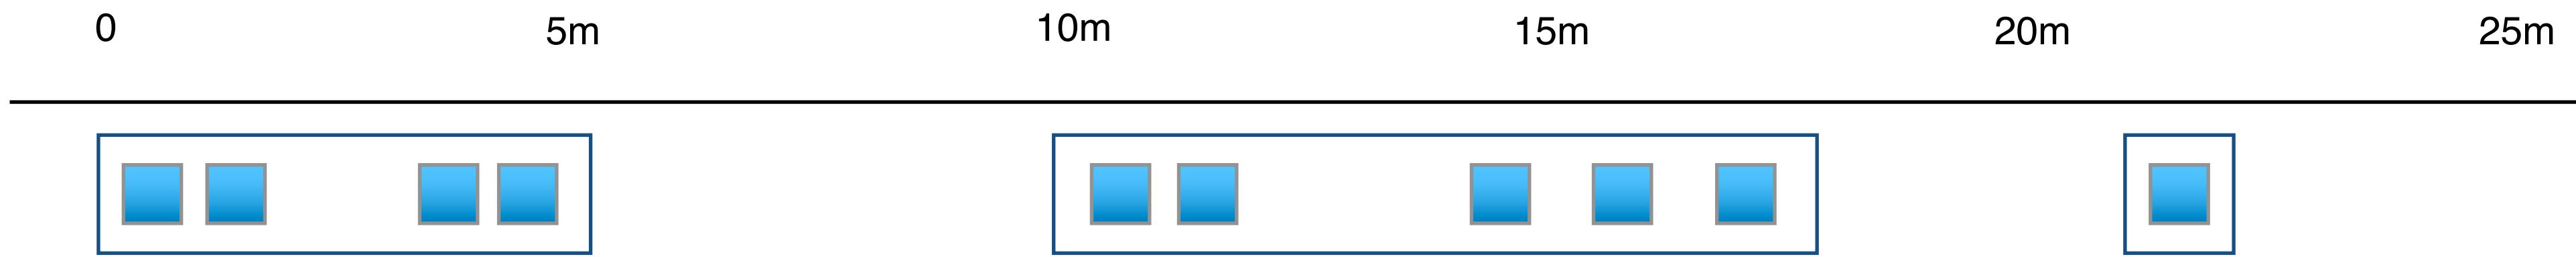
last element; we wait five minutes

Session Window



nothing else; we wait 5 minutes

Session Window



that's a new window

Stream API

Establishing Properties

- Create an `application.id` that represents your application group or "team"
- Serde is combination of Serializer/Deserializer
- Every stream application and KSQL app (later) is a consumer-producer

```
Properties props = new Properties();

props.put(StreamsConfig.APPLICATION_ID_CONFIG, "my_stream_app");
props.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
props.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass());
props.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.Integer().getClass());
```

Create a Stream Builder

- Always start with a `StreamBuilder` object
- This is the GoF builder pattern, where we will create a `Topology` object that represents our data pipeline

```
StreamsBuilder builder = new StreamsBuilder();
```

Read from a Topic

- Recall there are two representations of how to ingest data from a topic
 - Stream - Flow of data
 - Table - An aggregation; Or a database
- How do you read the topic is up to you
- Result is either a stream or table object that can take functional programming DSL.

```
KStream<String, Integer> stream = builder.stream("my_topic");
```

```
KTable<String, Integer> table = builder.table("my_topic");
```

Standard Functional Programming

- map
- filter
- flatMap
- groupBy
- reduce
- window
- join
- leftJoin
- outerJoin

Mapping

- Given a Stream

```
(1, "Hello"), (2, "Zoom"), (3, "Fold")
```

- Applying map

```
stream.map((key, value) -> new KeyValue<>(key + 1, value + "!"));
```

- Resulting in

```
(2, "Hello!"), (3, "Zoom!"), (4, "Fold!")
```

Filtering

- Given a Stream

```
(1, "Hello"), (2, "Zoom"), (3, "Fold"), (4, "Past")
```

- Applying filter

```
stream.filter((key, value) -> key % 2 == 0);
```

- Resulting in

```
(2, "Zoom"), (4, "Past")
```

FlatMap

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying flatMap

```
stream.flatMap((key, value) -> Arrays.asList(  
    new KeyValue(key, value),  
    new KeyValue(key * 100, value + " Hundred")))
```

- Resulting in

```
(1, "One"), (100, "One Hundred"), (2, "Two"), (200, "Two Hundred"), (3, "Three"), (300, "Three Hundred"),  
(4, "Four"), (400, "Four Hundred")
```

Peek

- Peers into the Stream and view what goes through

```
stream1.peek((key, value) ->  
    System.out.printf("key: %s, value: %d", key, value));
```

GroupBy

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy

```
stream.groupBy((key, value) -> key % 2 == 0 ? "Even", "Odd")
```

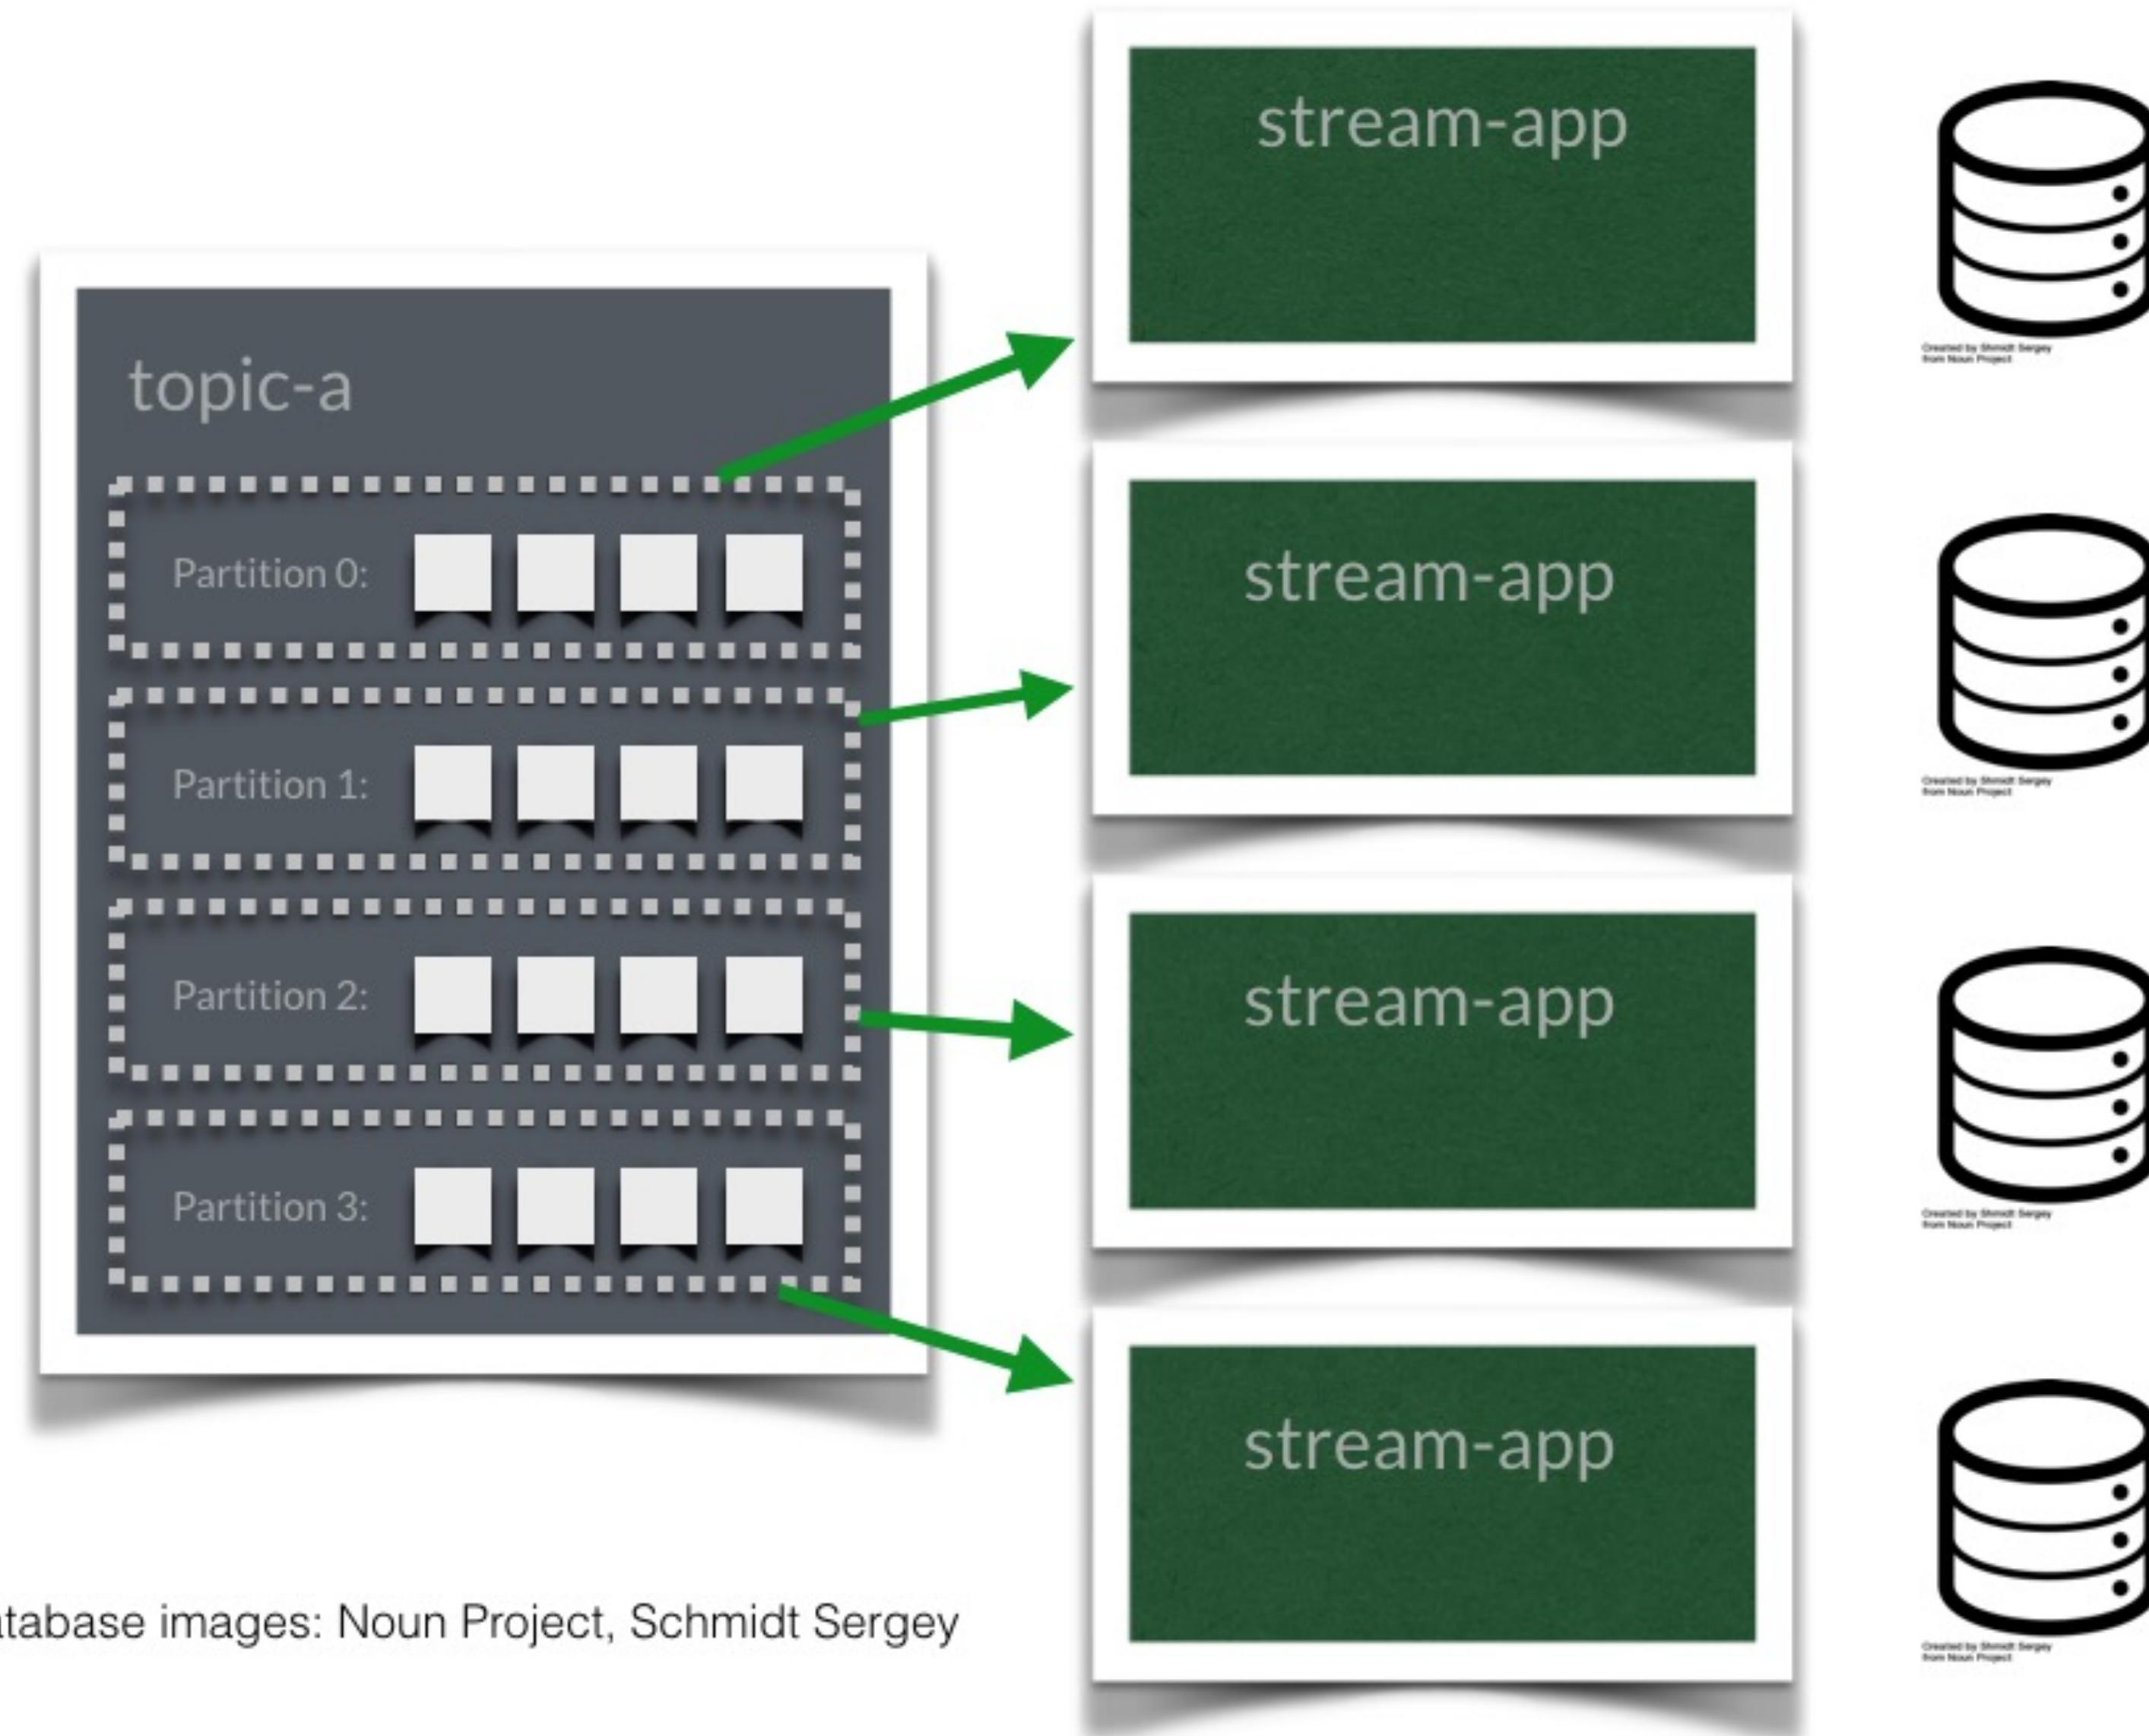
- Resulting in KGroupedStream

Once we have KGroupedStream

- `groupBy` creates a `KGroupedStream` object
- You have to decide what to do with this group
 - `count`
 - `aggregate`
 - `reduce`
 - `windowBy`
- Once you call any of the above, you have a `KTable`

KTable

- Represents a changelog, where the oldest records are not important
- Ideal for counters, and state
- KTable values are stored at each application's ephemeral storage
- Most ephemeral storage is backed by Rocks DB
- All streaming is then backed up as a topic for resilience



GroupByKey

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupByKey

```
stream.groupByKey()
```

- Resulting in KGroupedStream
- Is Synonymous with the Following

```
stream.groupBy((key, value) -> key)
```

Count

- Aggregates the count of what has been grouped
- Records with null key or value are ignored
- Returns KTable that represents the latest rolling count
- Will maintain the count in ephemeral storage, (like RocksDB)

Count

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy with count

```
stream.groupBy((key, value) -> key % 2 == 0 ? "Even" : "Odd").count();
```

- Will result in the RocksDB ephemeral storage

```
("Even", 2)
```

```
("Odd", 2)
```

Reduce

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy with reduce

```
stream.groupBy((key, value) -> key % 2 == 0 ? "Even" : "Odd")  
    .reduce((value1, value2) -> value1 + "," + value2);
```

- Will result in the RocksDB ephemeral storage

```
("Even", "Two,Four")  
("Odd", "One,Three")
```

Aggregate

- Given a Stream

```
(1, "One"), (2, "Two"), (3, "Three"), (4, "Four")
```

- Applying groupBy with count

```
groupedStream.aggregate(() -> "0:Zero,",  
(nextKey, nextValue, aggString) -> aggString + nextKey + ":" + nextValue ",");
```

- Will result in the RocksDB ephemeral storage

```
("Even", "0:Zero,2:Two,4:Four")
```

```
("Odd", "0:Zero,1:One,3:Three")
```

Hopping Windows

- Creating a Time Window

```
import java.util.concurrent.TimeUnit;  
  
import org.apache.kafka.streams.kstream.TimeWindows;  
  
  
  
  
long windowHeightMs = TimeUnit.MINUTES.toMillis(5); // 5 * 60 * 1000L  
long advanceMs =     TimeUnit.MINUTES.toMillis(1); // 1 * 60 * 1000L  
TimeWindows.of(windowHeightMs).advanceBy(advanceMs);
```

- Applying the Window

```
streams.groupBy(...).windowedBy(TimeWindows.of(windowHeightMs).advanceBy(advanceMs))  
;
```

Session Windows

- Creating a Time Window

```
import java.util.concurrent.TimeUnit;  
import org.apache.kafka.streams.kstream.SessionWindows;
```

```
SessionWindows.with(TimeUnit.MINUTES.toMillis(5));
```

- Applying the Window

```
streams.groupBy(...).windowedBy(SessionWindows.with(TimeUnit.MINUTES.toMillis(5));
```

Dump results to a topic

- Dump the results to a topic using to

```
KStream<String, Integer> stream = builder.stream("my_topic");
stream.filter(...).to("new_topic");
```

- Dump the results to a topic using through to post to topic and continue:

```
KStream<String, Integer> stream = builder.stream("my_topic");
stream.filter(...).through("new_topic").flatMap(...).to("other_topic")
```

Build the Topology and Stream

- Once you create your stream, you can build a Topology object:

```
Topology topology = builder.build();
```

- Create a Kafka Streams object with both Topology and Properties you created earlier

```
KafkaStreams streams = new KafkaStreams(topology, props);
```

- And Run...

```
streams.start();
```

Adding a Shutdown Hook

- As always, be a good citizen, properly shutdown resources

```
Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
```

KSQL

KSQL

Open source streaming SQL engine for Apache Kafka

Provides SQL interface for stream processing on Kafka

Scalable, elastic, fault-tolerant, and real-time.

KSQL

Supports a wide range of streaming operations, including:

- Data filtering
- Transformations
- Aggregations
- Joins
- Windowing
- Sessionization

KSQL Server



KSQL CLI

KSQL Architecture

KSQL Server communicates with the broker

KSQL servers are run separately from the KSQL CLI client and Kafka brokers.

You can deploy servers on remote machines, VMs, or containers

KSQL CLI interacts with KSQL Servers via a REST API

KSQL Server runs in two modes: CLI, Headless

KSQL

Kafka Streams

Producer/Consumer API

Starting in CLI Mode

```
./bin/ksql http://localhost:8088
```

Starting in Headless Mode

```
/bin/ksql-start-server \  
server_path/etc/ksql/ksql-server.properties \  
--queries-file /path/to/queries.sql
```



Tapping a Stream Based Topic

```
CREATE STREAM pageviews \
(viewtime BIGINT, \
userid VARCHAR, \
pageid VARCHAR) \
WITH (KAFKA_TOPIC='pageviews', \
VALUE_FORMAT='DELIMITED');
```

Tapping a Table Based Topic

```
CREATE TABLE users \
  (registertime BIGINT, \
  userid VARCHAR, \
  gender VARCHAR, \
  regionid VARCHAR) \
WITH (KAFKA_TOPIC = 'users', \
      VALUE_FORMAT='JSON', \
      KEY = 'userid');
```

Showing the Streams

SHOW STREAMS;

Stream Name	Kafka Topic	Format
-------------	-------------	--------

PAGEVIEWS	pageviews	DELIMITED
-----------	-----------	-----------

Showing the Tables

```
SHOW TABLES;
```

Table Name	Kafka Topic	Format
USERS	users	JSON

Supported SQL Types

- BOOLEAN
- INTEGER
- BIGINT
- DOUBLE
- VARCHAR (or STRING)
- ARRAY<ArrayType> (JSON and AVRO only. Index starts from 0)
- MAP<VARCHAR, ValueType> (JSON and AVRO only)
- STRUCT<FieldName FieldType, ... > (JSON and AVRO only)

Various Operators

Scalar Operators:

ABS, ARRAYCONTAINS, CEIL, CONCAT,

EXTRACTJSONFIELD, FLOOR, GEO_DISTANCE, IFNULL, LCASE, LEN, MASK, RANDOM, ROUND,

STRINGTOTIMESTAMP, SUBSTRING, TIMESTAMPTOSTRING,

TRIM, UCASE

Aggregator Operators:

COUNT, MAX, MIN, SUM, TOPK, TOPKDISTINCT,

WINDOWSTART, WINDOWEND

Displays Information

SHOW TOPICS	Show all topics
SHOW <TOPIC>	Display Topic
SHOW STREAMS	Show Streams
SHOW TABLES	Show Tables
SHOW FUNCTIONS	Show Available Functions
SHOW QUERIES	Show Persistent Queries

Describing

DESCRIBE

List columns in stream or table

DESCRIBE EXTENDED

Describe in detail stream or table;
runtime statistics, and queries that
populate the stream or table

Persistent Query

```
CREATE TABLE users_female AS \  
SELECT userid, gender, regionid FROM users \  
WHERE gender='FEMALE';
```

What makes this a persistent query is the form

```
CREATE (TABLE|STREAM) AS SELECT
```

Once executed, it will continuously run, until terminated

with TERMINATE command

Non-Persistent Query

```
SELECT * FROM pageviews  
WHERE ROWTIME >= 1510923225000  
AND ROWTIME <= 1510923228000;
```

A `LIMIT` can be used to limit the number of rows returned.

Once the limit is reached the query will terminate.

Persistent v. Non Persistent

	Persistent	Non Persistent
Where does data go?	Will store into a topic	Will display on screen
How do I stop it?	Find query id using SHOW QUERIES and TERMINATE <id>	CTRL + C
How to create?	CREATE STREAM AS SELECT ...	SELECT

Selecting Customized Fields

```
CREATE STREAM pageviews_transformed \
    WITH (TIMESTAMP='viewtime', \
          PARTITIONS=5, \
          VALUE_FORMAT='JSON') AS \
    SELECT viewtime, \
           userid, \
           pageid, \
           TIMESTAMPTOSTRING \
               (viewtime, 'yyyy-MM-dd HH:mm:ss.SSS') \
               AS timestamp \
    FROM pageviews \
    PARTITION BY userid;
```

Start from the Beginning

```
SET 'auto.offset.reset' = 'earliest';
```

Tumbling Window

```
SELECT item_id, SUM(quantity)  
FROM orders  
WINDOW TUMBLING (SIZE 20 SECONDS)  
GROUP BY item_id;
```

Hopping Window

```
SELECT item_id, SUM(quantity)  
FROM orders  
WINDOW HOPPING (SIZE 20 SECONDS,  
                  ADVANCE BY 5 SECONDS)  
GROUP BY item_id;
```

Session Window

```
SELECT item_id, SUM(quantity)  
FROM orders  
WINDOW SESSION (20 SECONDS)  
GROUP BY item_id;
```

Conclusion